

## ✓ Part I : Global includes

```
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")

import os
import pathlib
from matplotlib import pyplot as plt
import torch
import numpy as np
import cv2
import tensorflow as tf

# Uncomment to disable GPU usage.
# This is required for some models like Pridnet which has too many trainable parameters
tf.config.set_visible_devices([], 'GPU')

from tqdm.notebook import tqdm
import random

import data_importer

from tensorflow.compat.v1 import ConfigProto
from tensorflow.compat.v1 import InteractiveSession

config = ConfigProto()
config.gpu_options.allow_growth = True
session = InteractiveSession(config=config)
```

## ✓ Part II : Loading test images

```
from dataloader_lodopab_ct import get_validation_dataloader
noisy_dataset = get_validation_dataloader("../../../Dataset/LoDoPaB-CT/ground_truth_validation/")

    number of image paths : 28
    validation data image len : , 28

def reconstruct_image_from_patches(patches, num_patches_per_row):
    patch_size = patches.shape[1] # Assuming square patches
    num_patches = patches.shape[0]

    # Calculate the number of rows
    num_patches_per_col = num_patches // num_patches_per_row

    # Initialize an empty image to store the reconstructed result
    reconstructed_image = np.zeros((num_patches_per_col * patch_size, num_patches_per_row * patch_size))
```

```
# Reshape the patches into a 2D array
patches_2d = patches.reshape((num_patches_per_col, num_patches_per_row, patch_size, patch_size))

# Reconstruct the image by placing each patch in its corresponding position
for i in range(num_patches_per_col):
    for j in range(num_patches_per_row):
        reconstructed_image[i * patch_size:(i + 1) * patch_size, j * patch_size:(j + 1) * patch_size] = patches_2d[i, j]

return np.expand_dims(reconstructed_image, axis=-1)

noisy_array = [None] * 28
print(len(noisy_dataset))
for i, data in enumerate(noisy_dataset):
    noisy_array[i] = reconstruct_image_from_patches(torch.squeeze(data[i], axis=0), 8)
    if i == 28:
        break
noisy_array = np.array(noisy_array)
print(noisy_array)

28
[[[ [0.0000000e+00]
  [0.0000000e+00]
  [0.0000000e+00]
  ...
  [2.01545864e-01]
  [2.27233559e-01]
  [2.47729927e-01]]]

[[ [0.0000000e+00]
  [0.0000000e+00]
  [0.0000000e+00]
  ...
  [1.97884500e-01]
  [2.25247324e-01]
  [2.57035911e-01]]]

[[ [0.0000000e+00]
  [1.35118389e-04]
  [4.09101369e-04]
  ...
  [1.93417832e-01]
  [2.24700466e-01]
  [2.67172039e-01]]]

...
[[ [1.91558048e-03]
  [8.41807865e-04]
  [1.29352743e-03]
  ...
  [2.50156038e-02]
  [1.58237405e-02]
  [7.00143981e-04]]]

[[ [2.05797283e-03]
  [2.28885561e-03]
```

```
[2.79895985e-03]
...
[2.37569660e-02]
[1.74138136e-02]
[9.26138554e-03]]

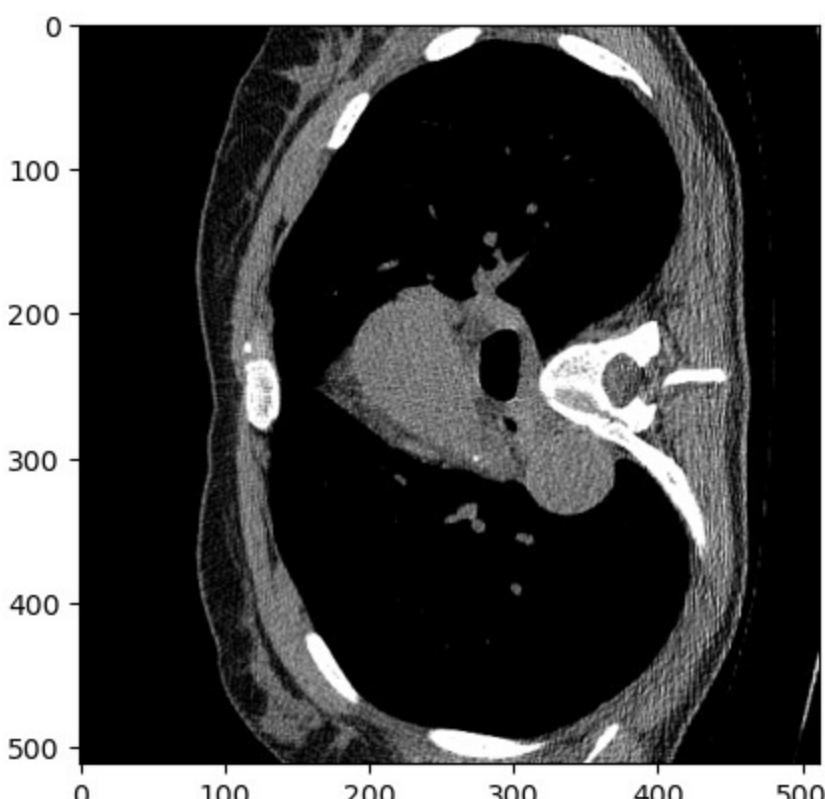
[[1.34693610e-03]
[3.74236028e-03]
[4.11643507e-03]
...
[1.76892225e-02]
[1.77762713e-02]
[2.03140154e-02]]]

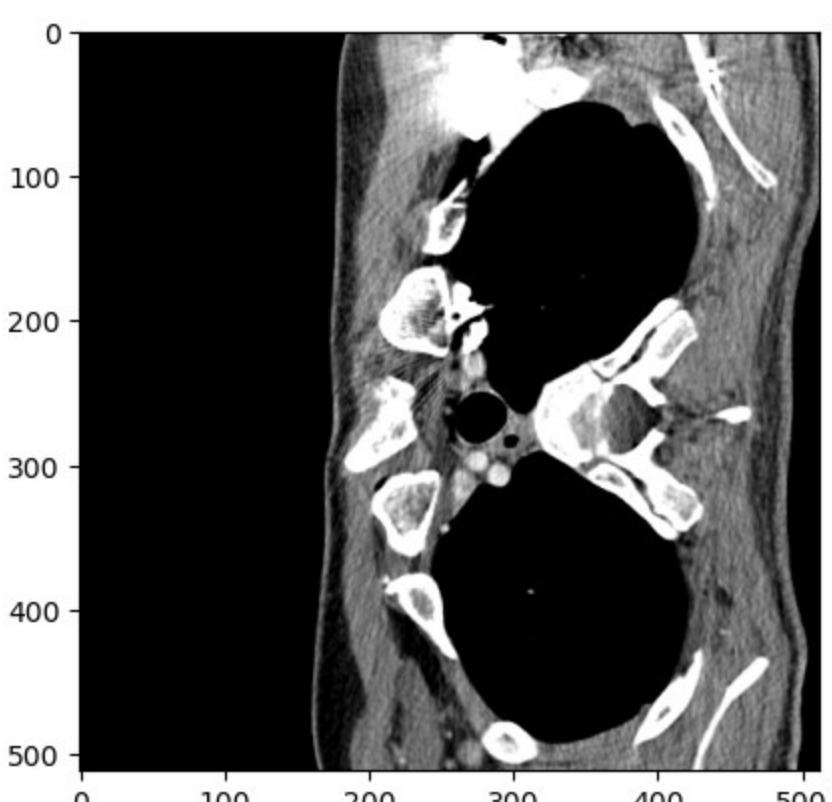
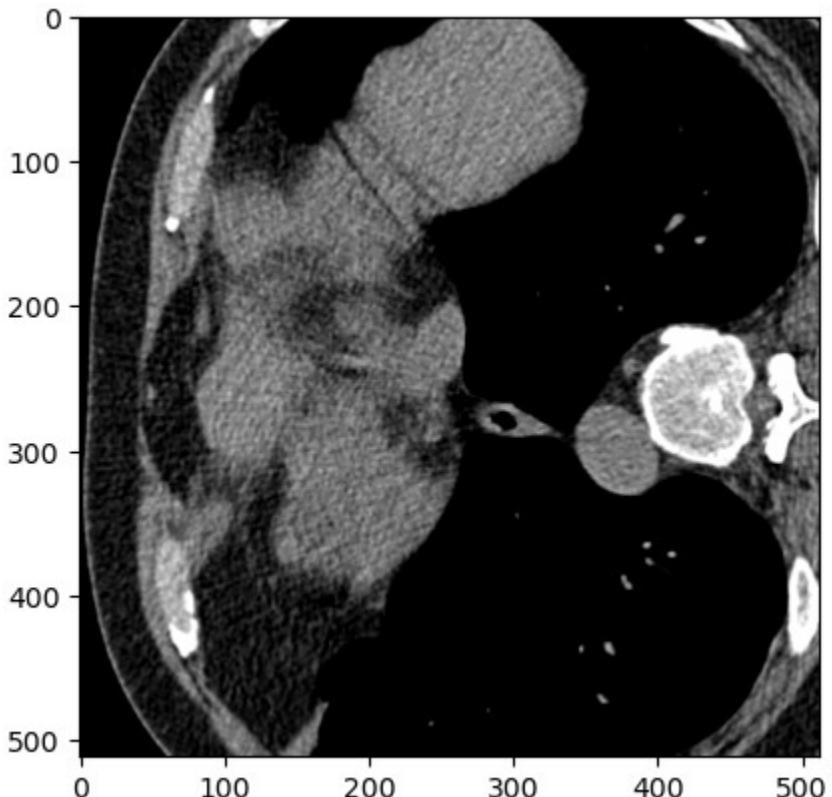
[[[6.74816035e-03]
[3.39281606e-03]
[6.21013518e-04]
...
[1.21287532e-01]
[7.91603923e-02]]]
```

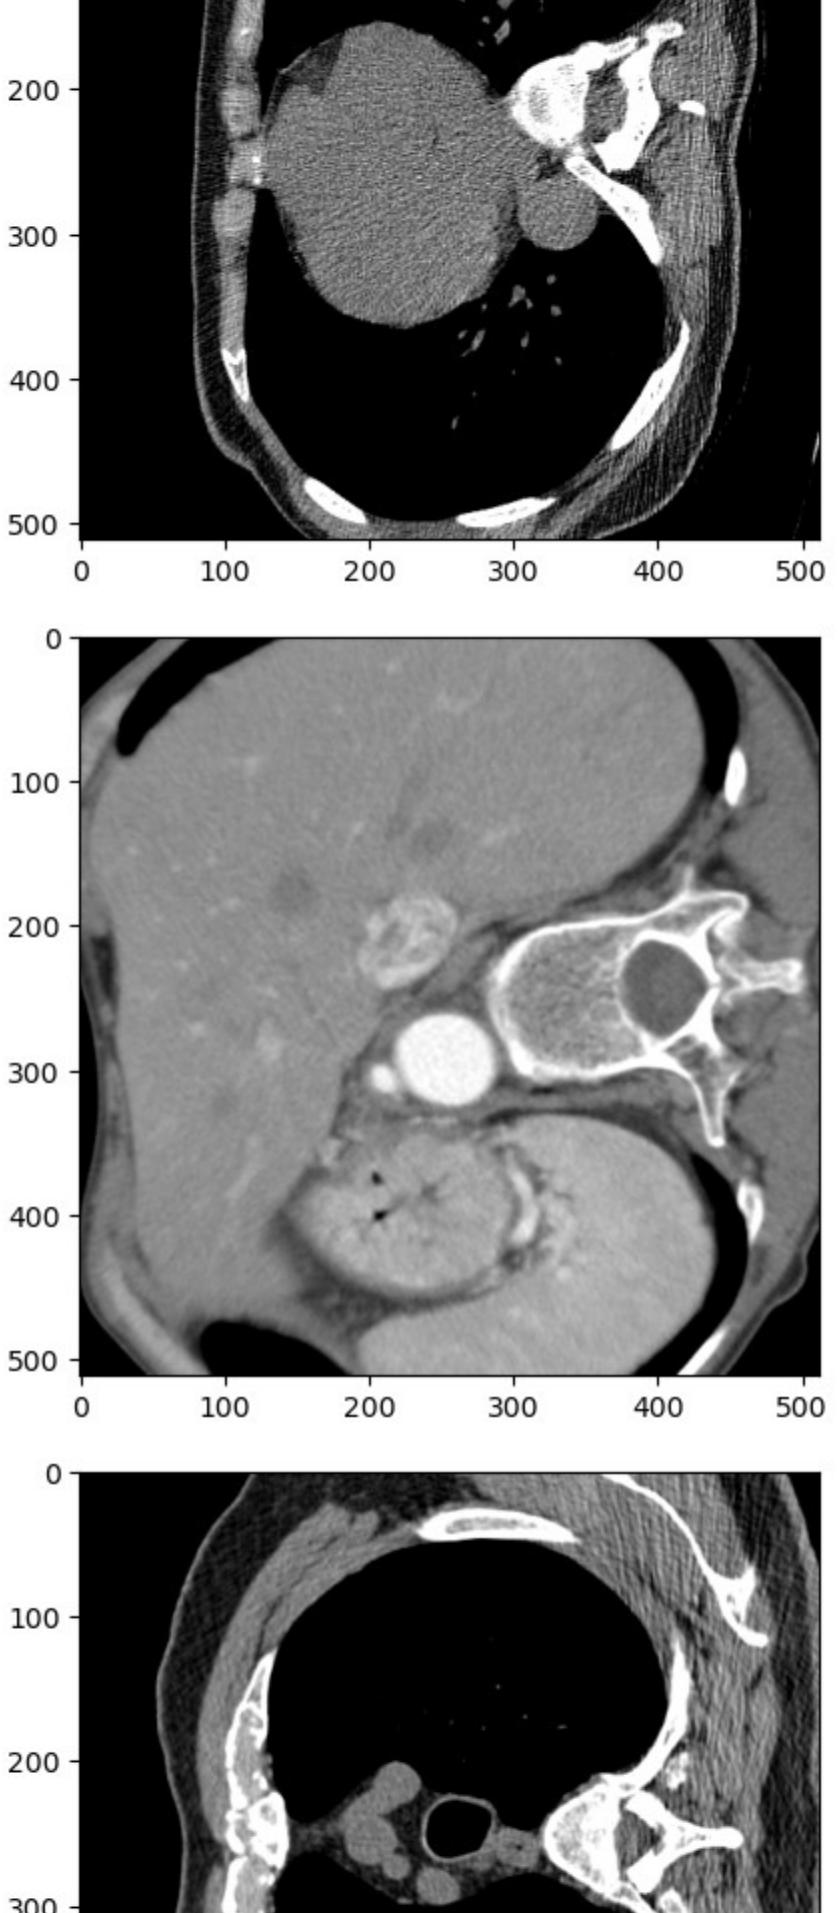
### Visualization of the noisy / ground truth image pair

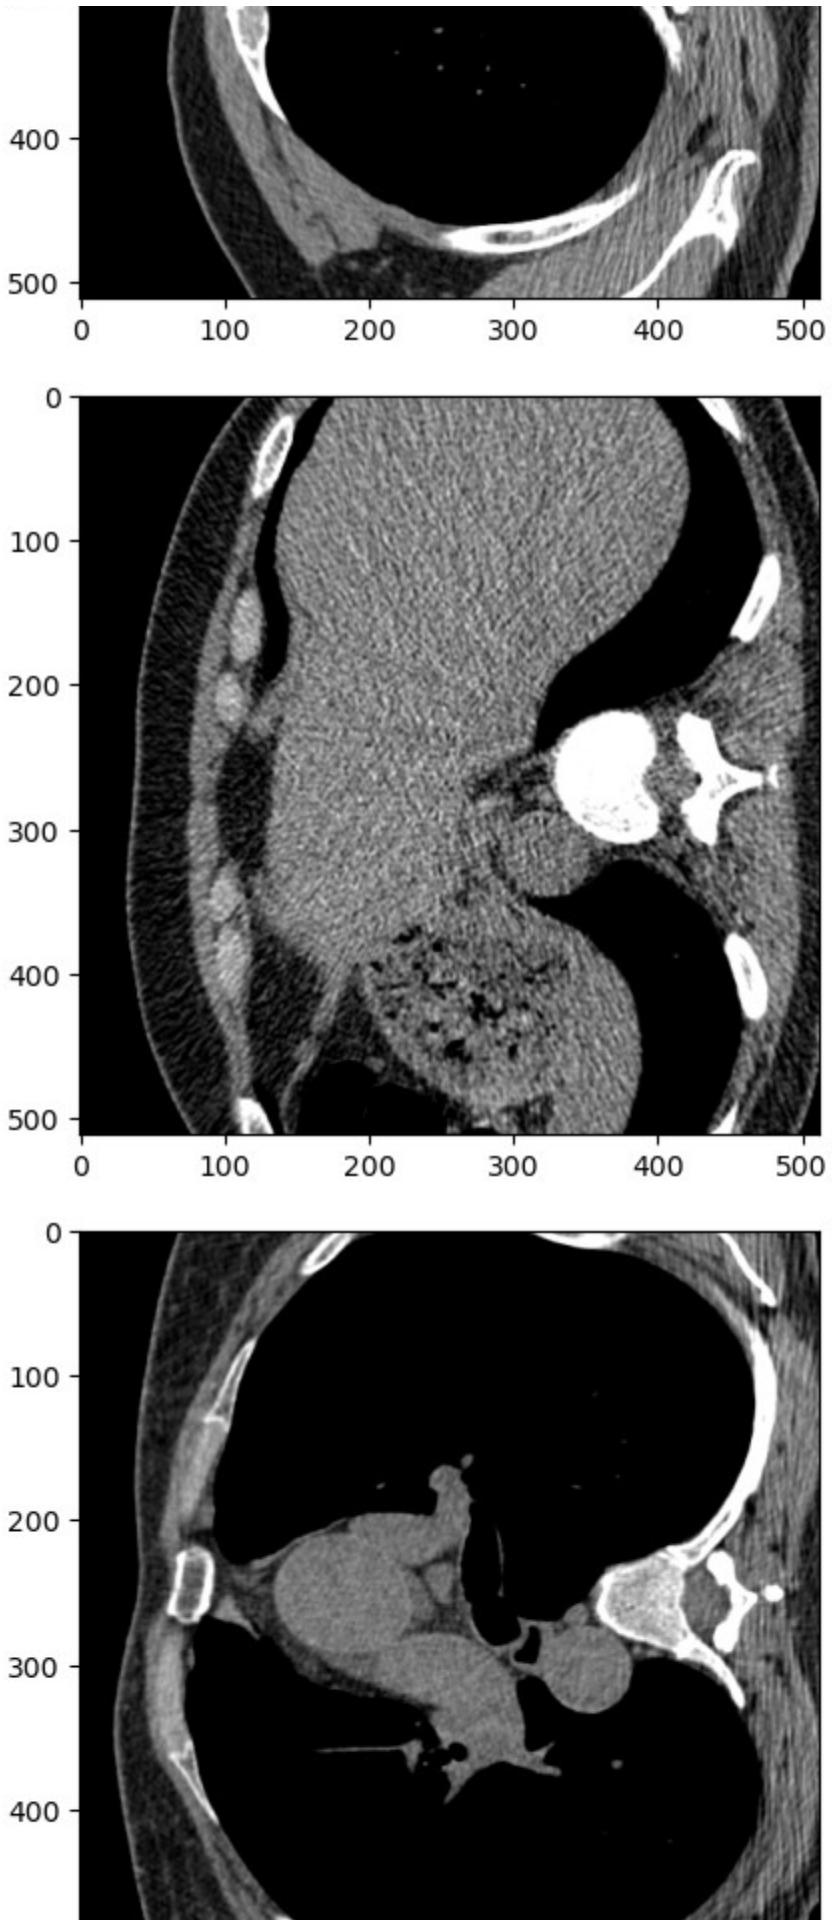
```
from data_importer import denormalize, trunc

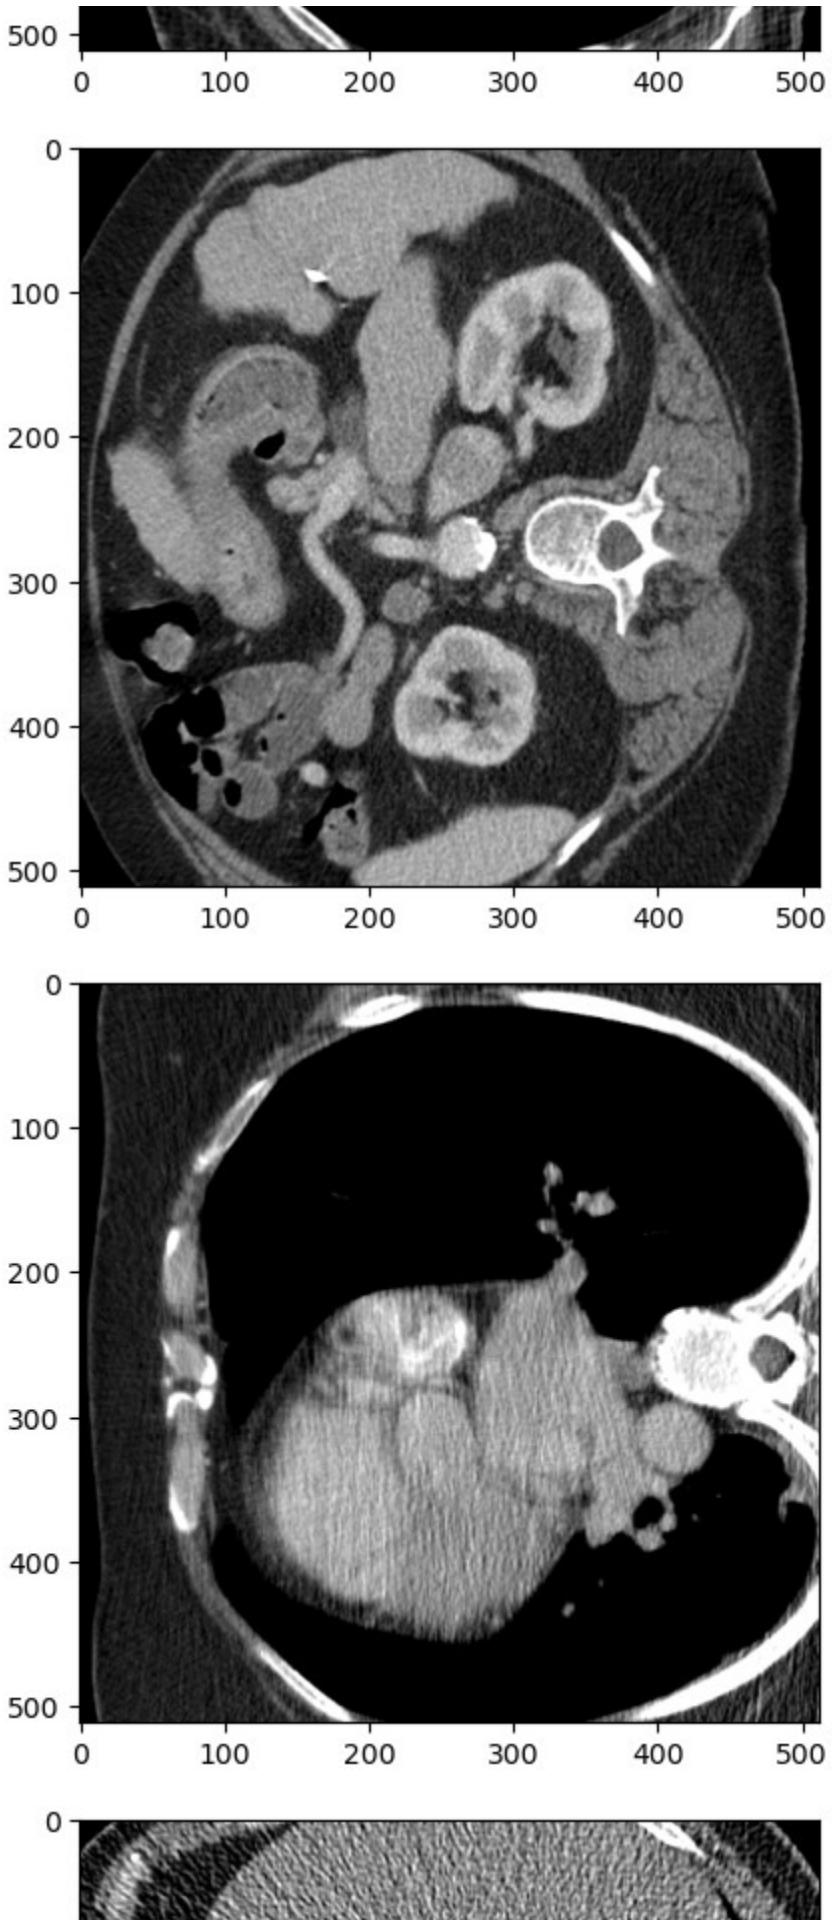
with torch.no_grad():
    for i, data in enumerate(noisy_array):
        plt.imshow(trunc(denormalize(data)), vmin=-160.0, vmax=240.0, cmap='gray')
        plt.show()
```

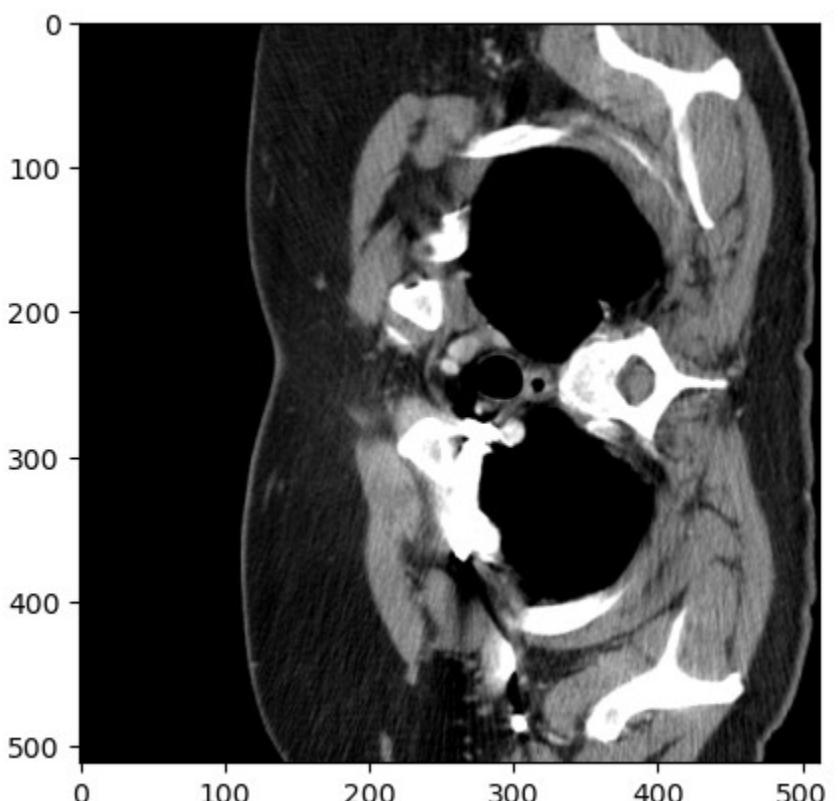
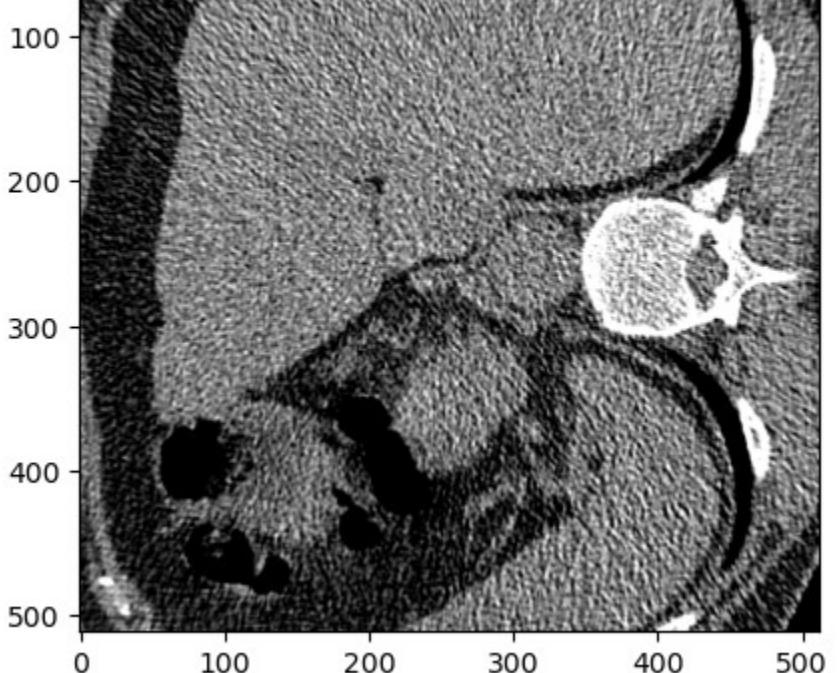


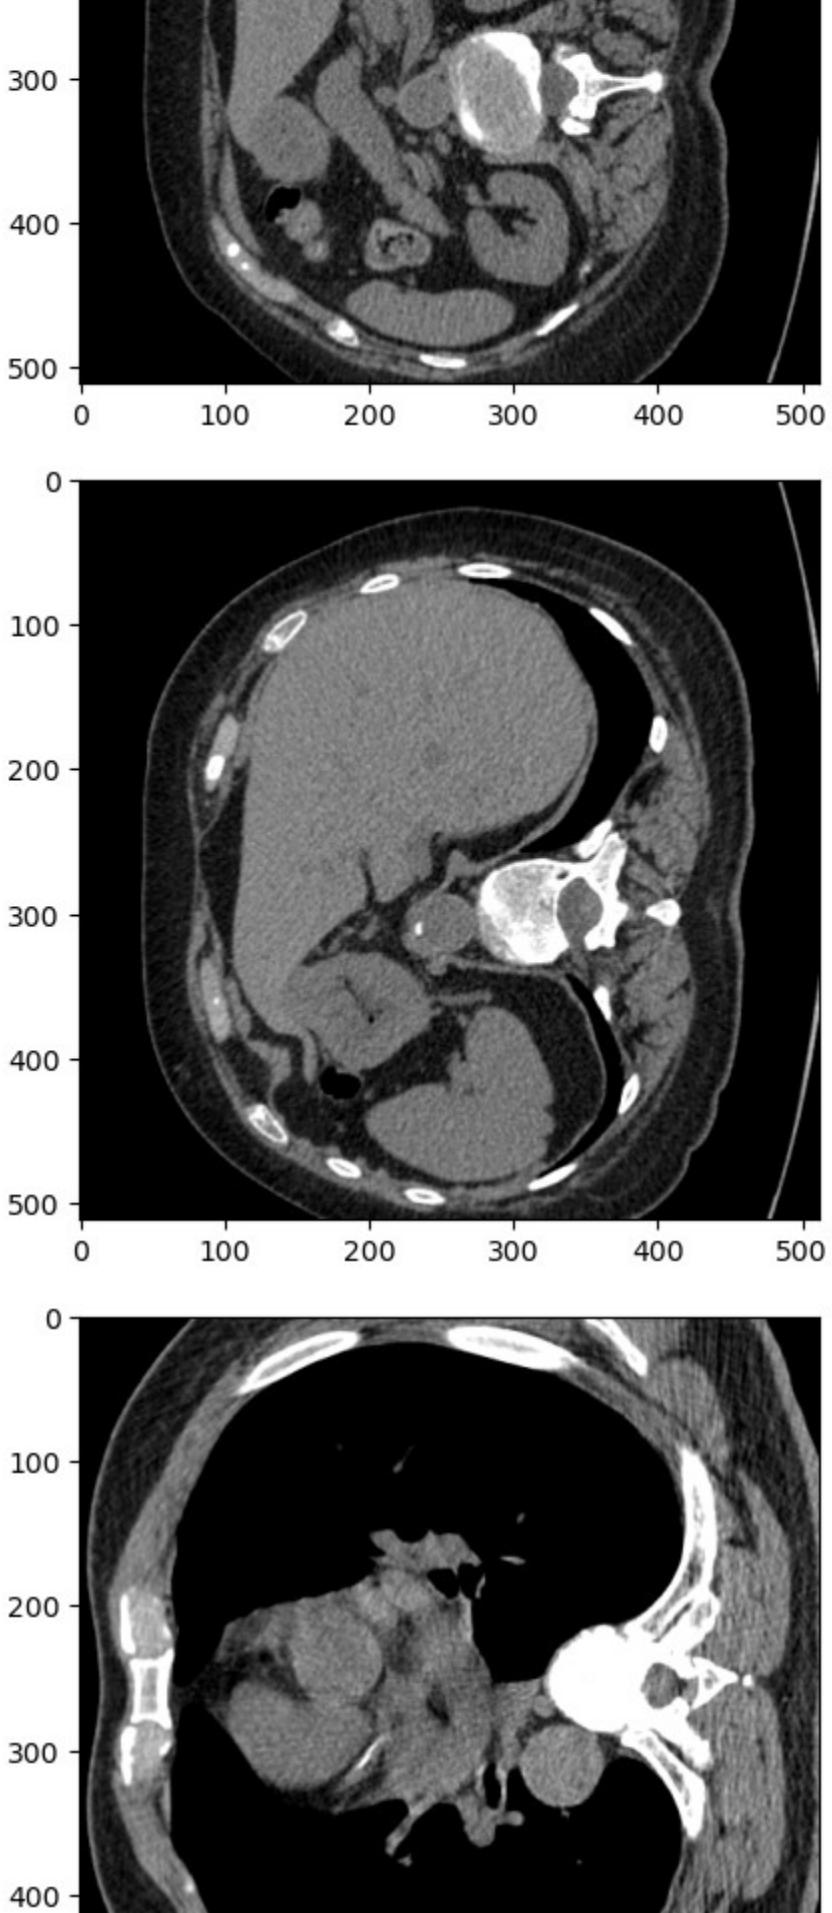


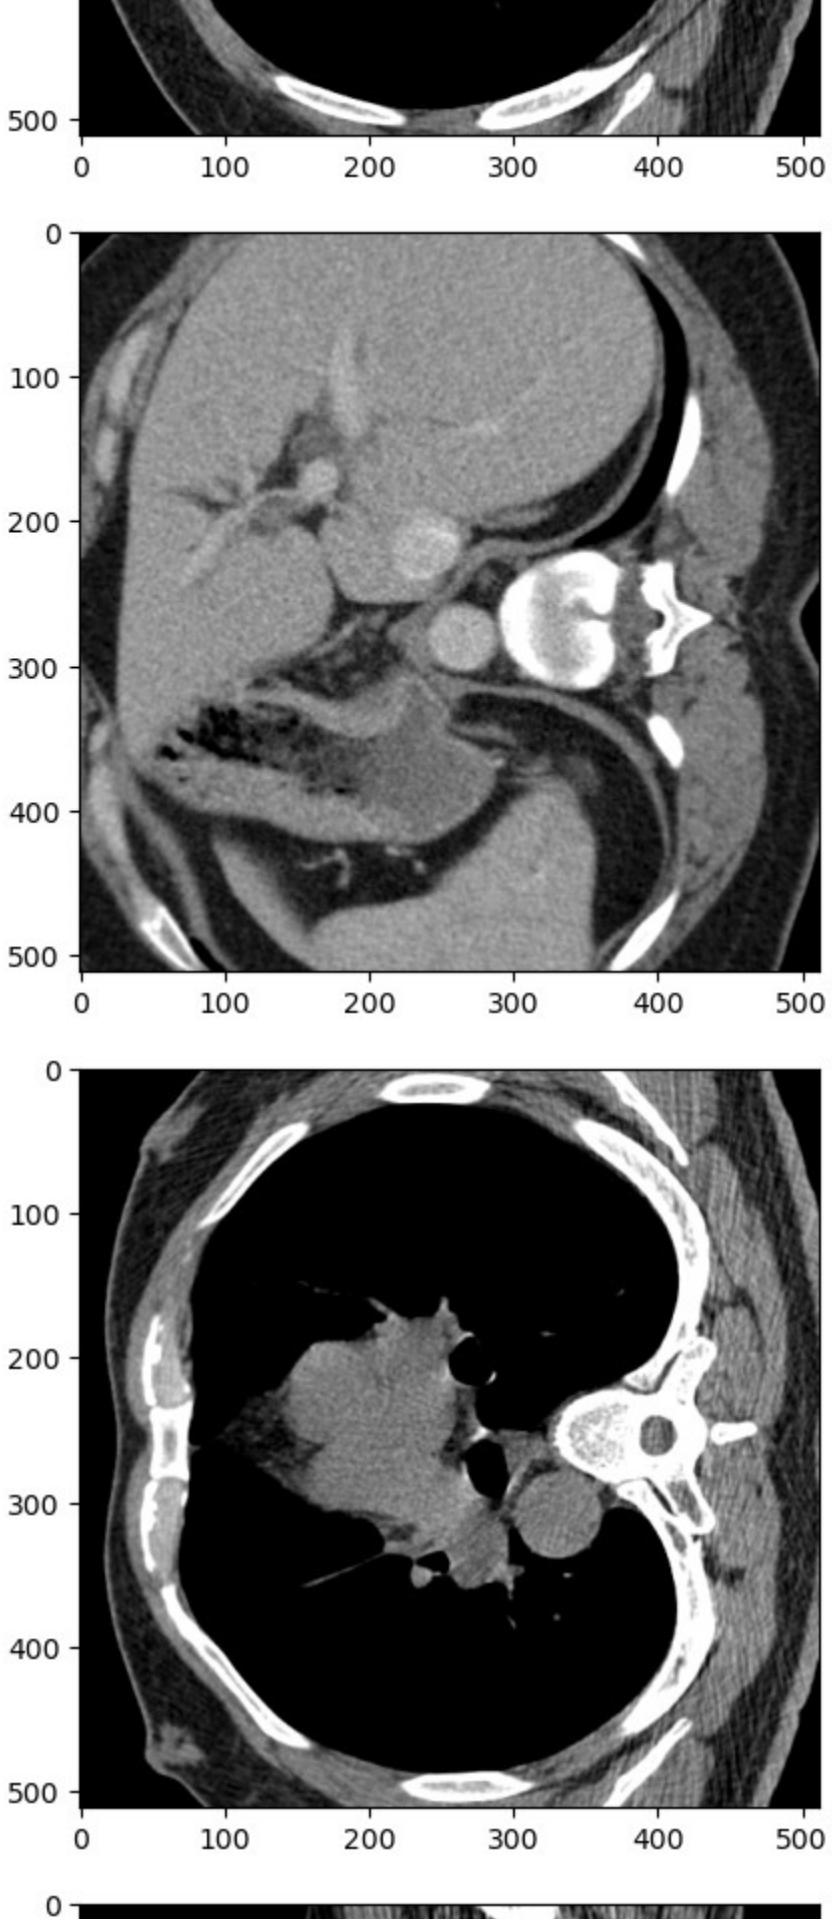


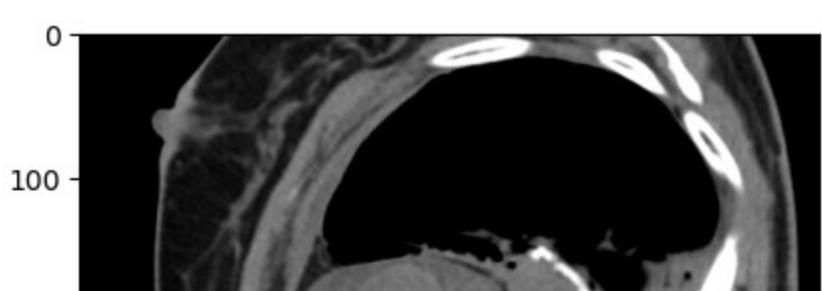
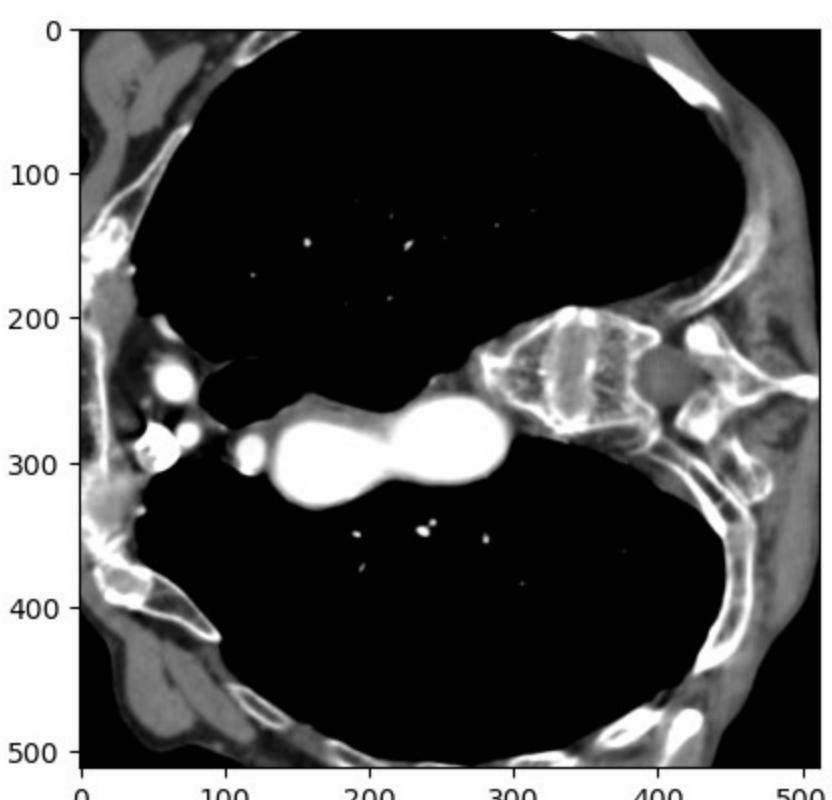
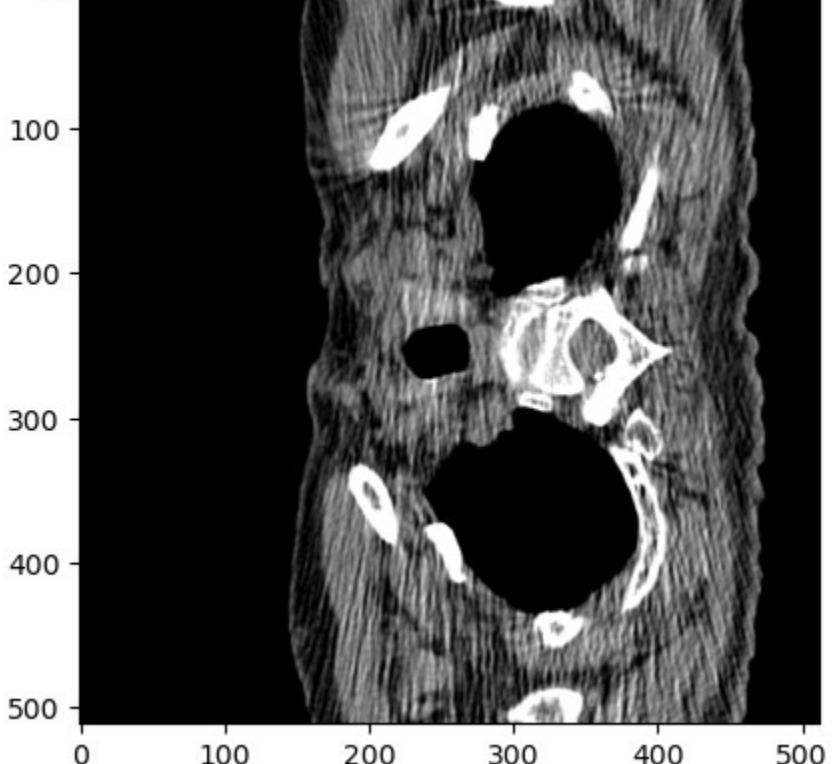


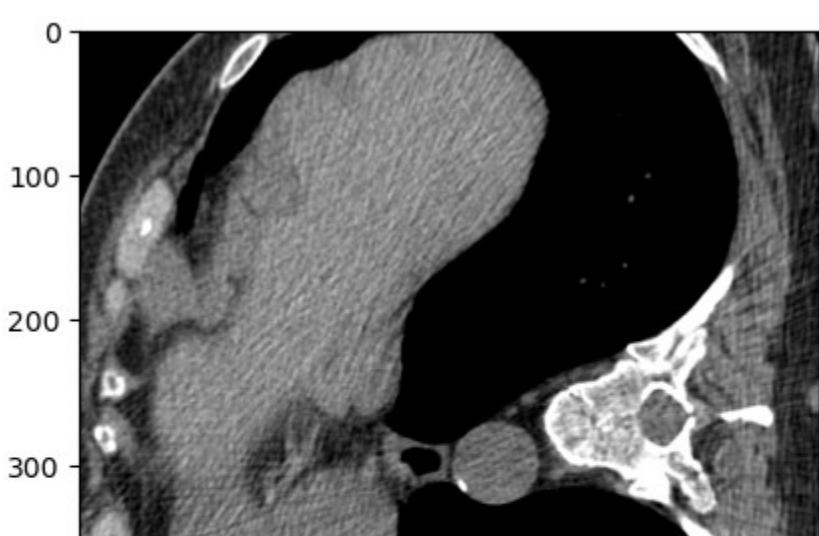
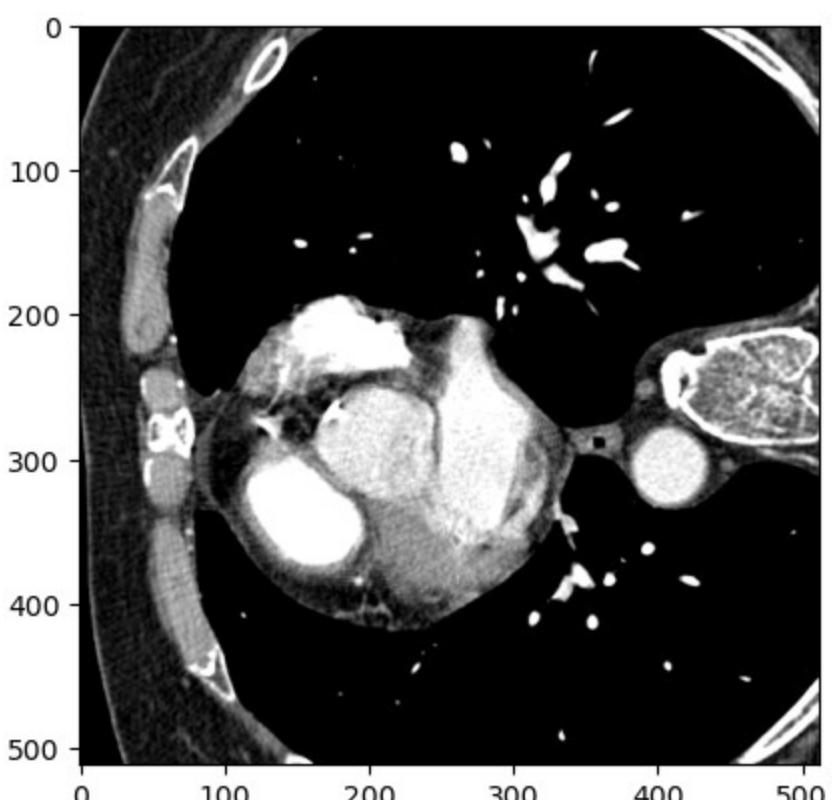
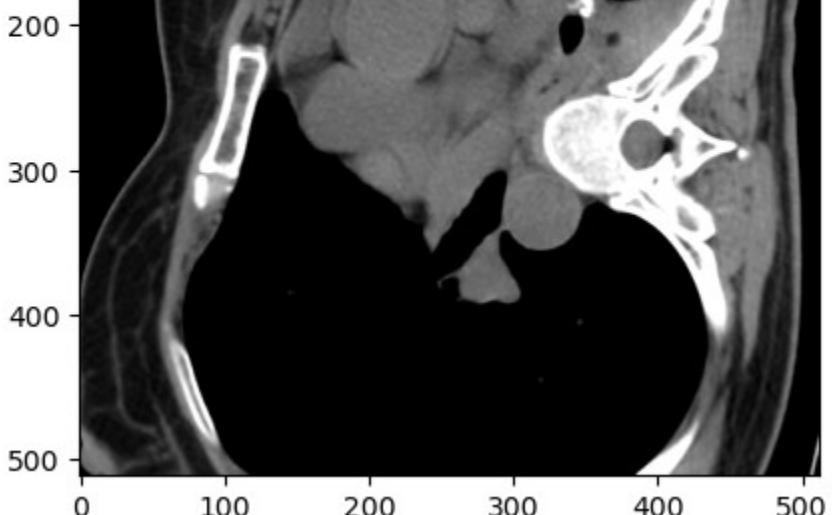


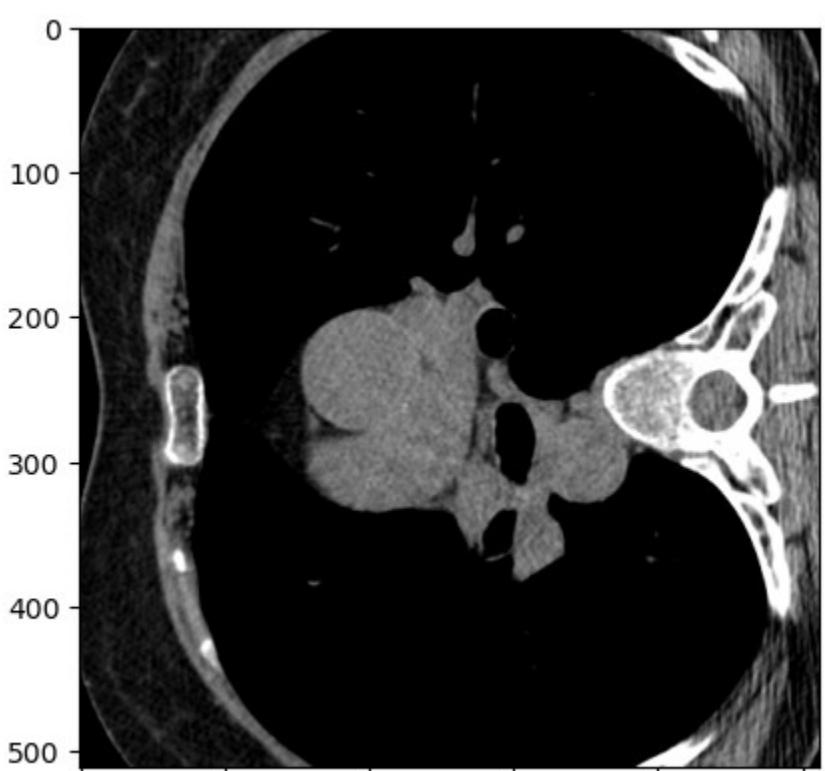
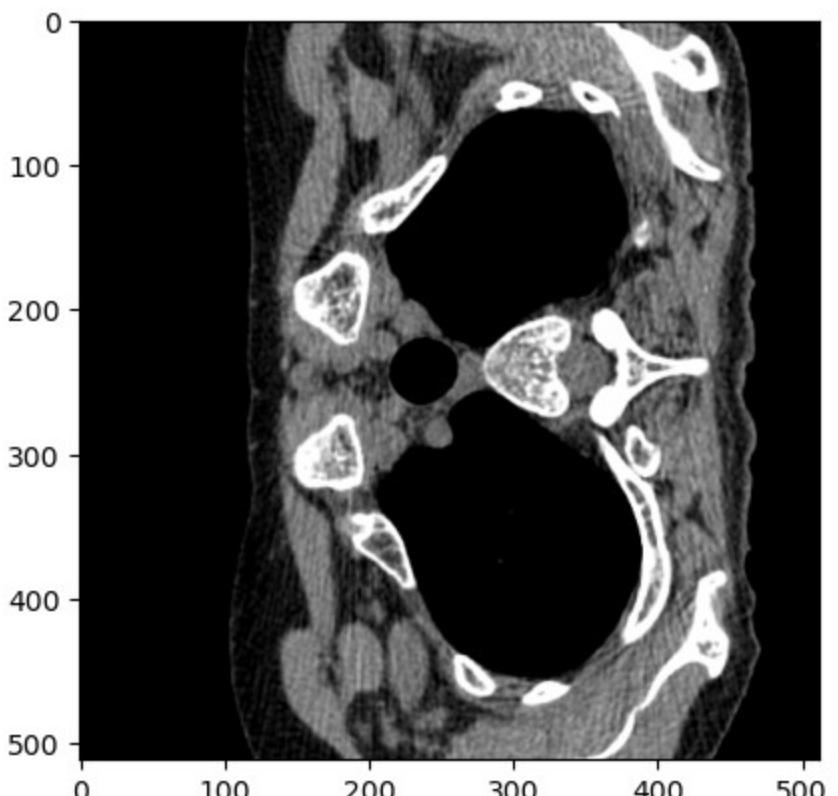
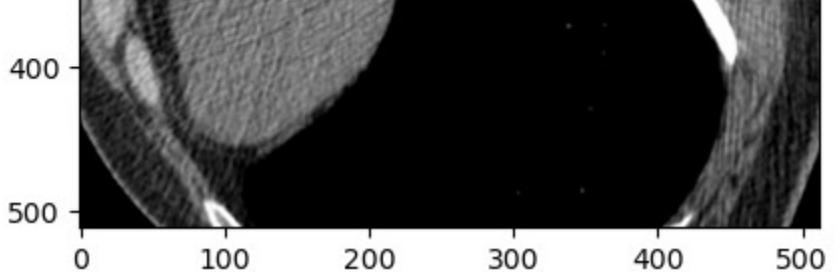


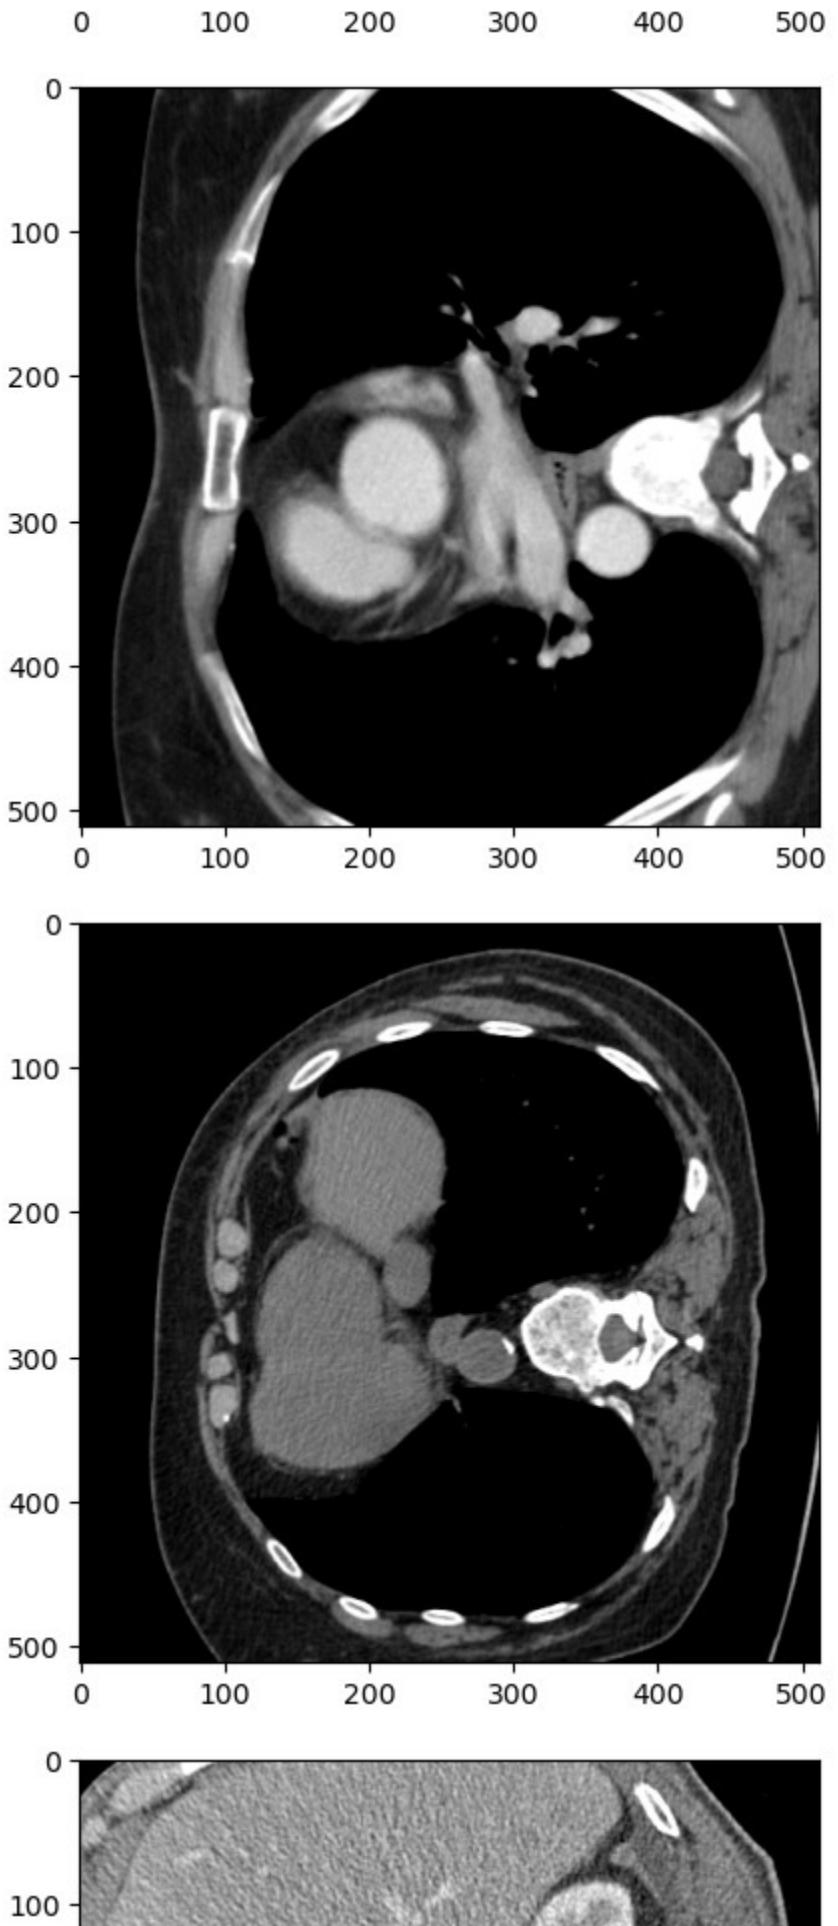


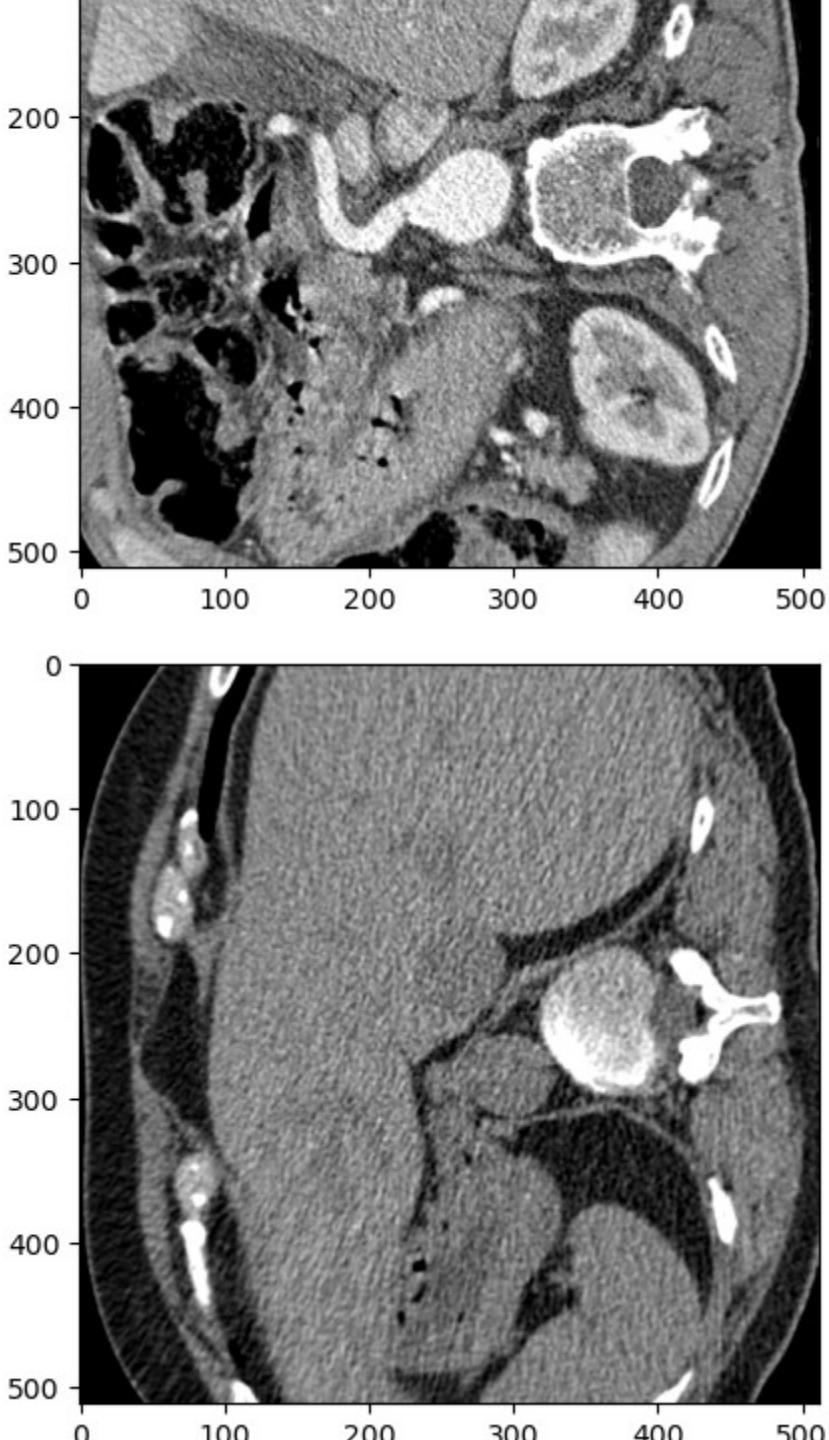












### ▼ Part III : Setup for Inference

```
# Inference
def inference_single_image(model, noisy_image):
    input_image = np.expand_dims(noisy_image, axis=0)
    predicted_image = model.predict(input_image)
    a = np.abs(np.min(predicted_image))
    b = np.max(predicted_image)
```

```
#predicted_image = predicted_image * (b - a) + a
return predicted_image[0]

def inference_batch_images(model, noisy_images):
    input_image = noisy_images

    predicted_image = model.predict(input_image).astype(np.float64)
    return predicted_image

def rgb2gray(rgb):
    return np.expand_dims(np.dot(rgb[...,:3], [0.2989, 0.5870, 0.1140]), axis=-1)

from skimage.metrics import structural_similarity as ssim
from skimage.metrics import peak_signal_noise_ratio
import sys
sys.path.append('..')

from metrics import compute_SSIM, compute_PSNR
from skimage.metrics import mean_squared_error as mse

def calculate_psnr(original_image, reconstructed_image, range=400):
    return peak_signal_noise_ratio(original_image, reconstructed_image, data_range=range)

psnr_value = peak_signal_noise_ratio(original_image, reconstructed_image, data_range=240+160)
return psnr_value

def calculate_ssim(original_image, reconstructed_image, range=400.0):
    ssim_value = ssim(original_image.astype(np.int16), reconstructed_image.astype(np.int16), win_size=11)
    return ssim_value

def calculate_rmse(original_image, reconstructed_image):
    return mse(original_image, reconstructed_image)

def visualize_predictions(model, X_test, n, predictions, model_name):
    random_numbers = list(range(n)) # not very random
    for i in random_numbers:
        gt_image= X_test[i].astype(np.float16)
        predicted_image = predictions[i].astype(np.float16)

        if predicted_image.shape[-1] == 3:
            predicted_image = rgb2gray(predicted_image)

        psnr_recon = calculate_psnr(trunc(denormalize(gt_image)), trunc(denormalize(predicted_image)))
        ssim_recon = calculate_ssim(trunc(denormalize(gt_image)), trunc(denormalize(predicted_image)))
        rmse_recon = calculate_rmse(trunc(denormalize(gt_image)), trunc(denormalize(predicted_image)))

        psnr_recon = round(psnr_recon, 4)
        ssim_recon = round(ssim_recon, 4)
        rmse_recon = round(rmse_recon, 4)

        f, axarr = plt.subplots(1,2, figsize=(21,21))

        axarr[0].imshow(trunc(denormalize(gt_image)), cmap='gray', vmin=-160.0, vmax=240.0)
        axarr[0].set_title("QD Image")
        axarr[0].set_axis_off()
        axarr[1].imshow(trunc(denormalize(predicted_image)), cmap='gray', vmin=-160.0, vmax=240.0)
        axarr[1].set_title("Reconstructed Image")
        axarr[1].set_axis_off()
```

```
    axarr[1].set_title("{{} Predicted Image : PSNR={}\\nSSIM={}\\nRMSE={}}".format(model_name, psnr_reco
    axarr[1].set_axis_off()

    plt.show()

from skimage.metrics import peak_signal_noise_ratio

def get_average_metrics(predicted_images, _noisy_array):
    psnr_original_mean = 0
    psnr_prediction_mean = 0

    ssim_original_mean = 0
    ssim_prediction_mean = 0

    mse_original_mean = 0
    mse_prediction_mean = 0

    i = 0
    for gt_img, predicted_img in zip(noisy_array, predicted_images):
        predicted_img= predicted_images[i]
        if predicted_img.shape[-1] == 3:
            predicted_img = rgb2gray(predicted_img)

        psnr_recon = calculate_psnr(trunc(denormalize(gt_img)), trunc(denormalize(predicted_img)))
        ssim_recon = calculate_ssim(trunc(denormalize(gt_img)), trunc(denormalize(predicted_img)))
        rmse_recon = calculate_rmse(trunc(denormalize(gt_img)), trunc(denormalize(predicted_img)))

        psnr_prediction_mean += psnr_recon
        ssim_prediction_mean += ssim_recon
        mse_prediction_mean += rmse_recon

        i = i + 1

    psnr_prediction_mean/=noisy_array.shape[0]
    ssim_prediction_mean/=noisy_array.shape[0]
    mse_prediction_mean/=noisy_array.shape[0]

    print("Predicted average gt-predicted PSNR ->", psnr_prediction_mean)
    print("Predicted average gt-predicted SSIM ->", ssim_prediction_mean)
    print("Predicted average gt-predicted MSE->", mse_prediction_mean)

    return round(psnr_prediction_mean, 4), round(ssim_prediction_mean, 4), round(mse_prediction_mean, 4)
```

## Part IV : Evaluation of each model

## Model I : Hformer (for base reference)

```
sys.path.append('../denoising-models/hformer_vit/model/')
sys.path.append('../denoising-models/hformer_vit/')
from hformer_model_extended import get_hformer_model, PatchExtractor

hformer_model = get_hformer_model(num_channels_to_be_generated=64, name="hformer_model_extended")
hformer_model.build(input_shape=(None, 64, 64, 1))
hformer_model.load_weights('../denoising-models/hformer_vit/test/experiments/full_dataset/hformer_64_ch'
print('Model summary : ')
print(hformer_model.summary())

Model summary :
Model: "hformer_model_extended"



| Layer (type)                                    | Output Shape | Param # |
|-------------------------------------------------|--------------|---------|
| input_projection_layer (InputProjectionLayer)   | multiple     | 320     |
| output_projection_layer (OutputProjectionLayer) | multiple     | 257     |
| conv_net_block_1 (ConvolutionalBlock)           | multiple     | 36416   |
| conv_net_block_2 (ConvolutionalBlock)           | multiple     | 36416   |
| downsampling_layer_1 (Conv2D)                   | multiple     | 32896   |
| downsampling_layer_2 (Conv2D)                   | multiple     | 131328  |
| upsampling_layer_1 (Conv2DTranspose)            | multiple     | 131200  |
| upsampling_layer_2 (Conv2DTranspose)            | multiple     | 32832   |
| hformer_block_1 (HformerBlock)                  | multiple     | 187392  |
| hformer_block_2 (HformerBlock)                  | multiple     | 735232  |
| hformer_block_3 (HformerBlock)                  | multiple     | 187392  |
| Total params:                                   | 1,511,681    |         |
| Trainable params:                               | 1,511,681    |         |
| Non-trainable params:                           | 0            |         |


```

---

None

```

def reconstruct_image_from_patches(patches, num_patches_per_row):
    patch_size = patches.shape[1] # Assuming square patches
    num_patches = patches.shape[0]

    # Calculate the number of rows
    num_patches_per_col = num_patches // num_patches_per_row

    # Initialize an empty image to store the reconstructed result
    reconstructed_image = np.zeros((num_patches_per_col * patch_size, num_patches_per_row * patch_size))

    # Reshape the patches into a 2D array
    patches_2d = patches.reshape((num_patches_per_col, num_patches_per_row, patch_size, patch_size))
    # Reconstruct the image by placing each patch in its corresponding position

    for i in range(num_patches_per_col):
        for j in range(num_patches_per_row):
            reconstructed_image[i * patch_size:(i + 1) * patch_size, j * patch_size:(j + 1) * patch_size] = patches_2d[i, j]

    return np.expand_dims(reconstructed_image, axis=-1)

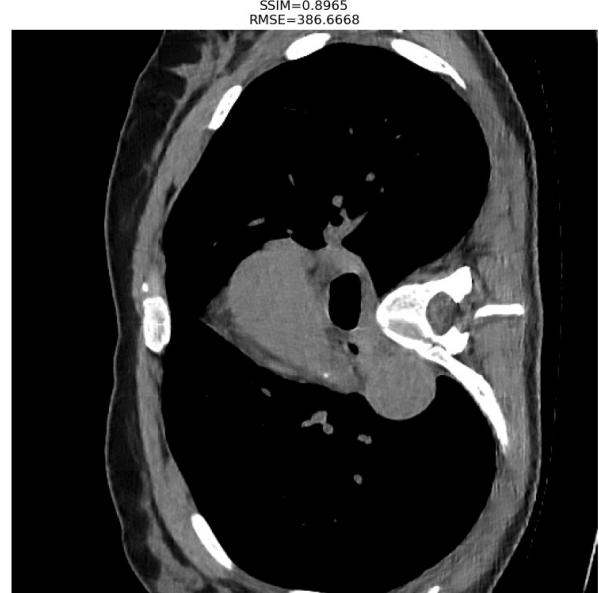
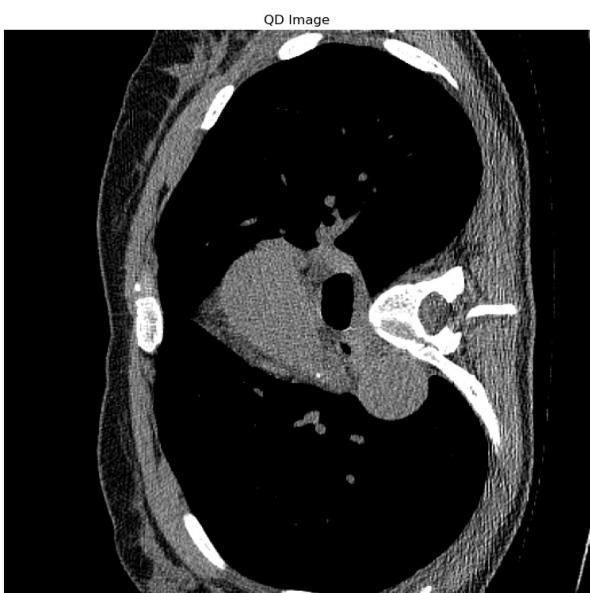
# View the predictions
patch_extractor = PatchExtractor(patch_size=64, stride=64, name="patch_extractor")
noisy_image_patches_array = patch_extractor(noisy_array)

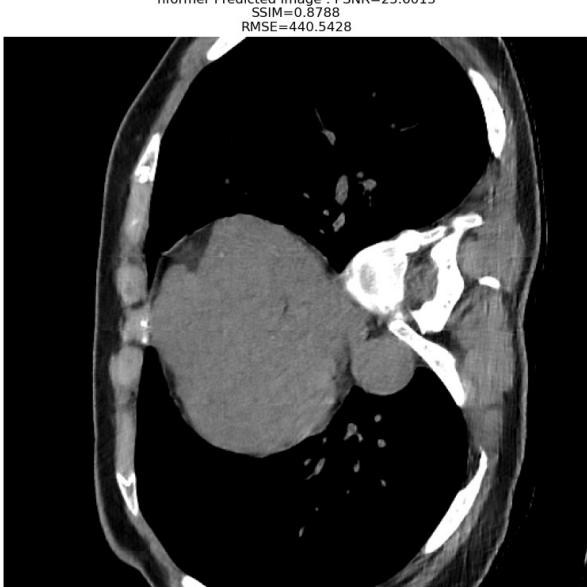
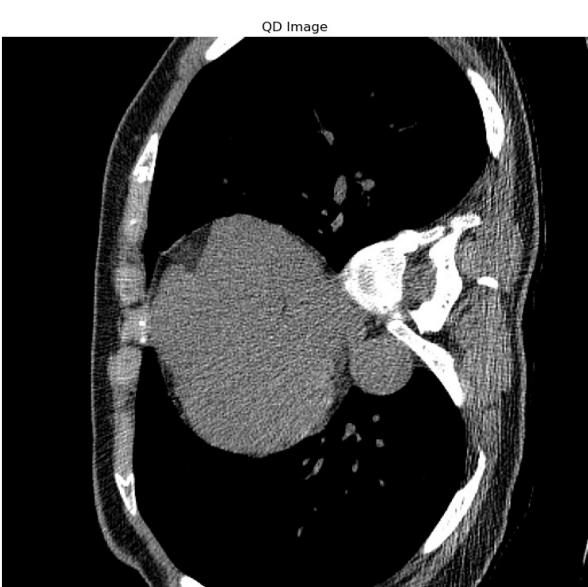
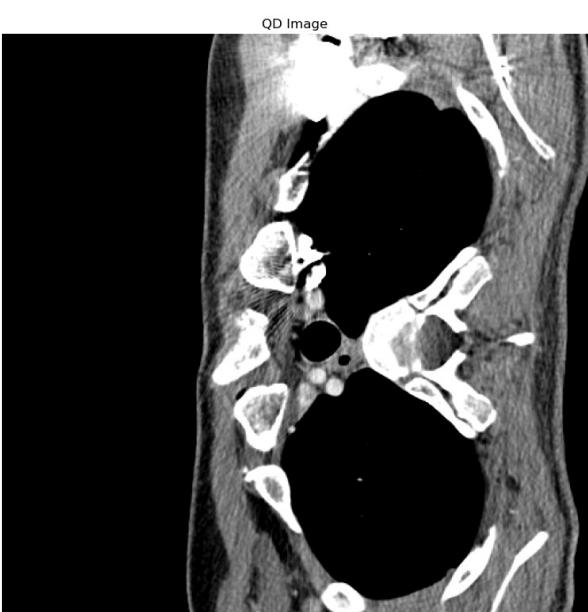
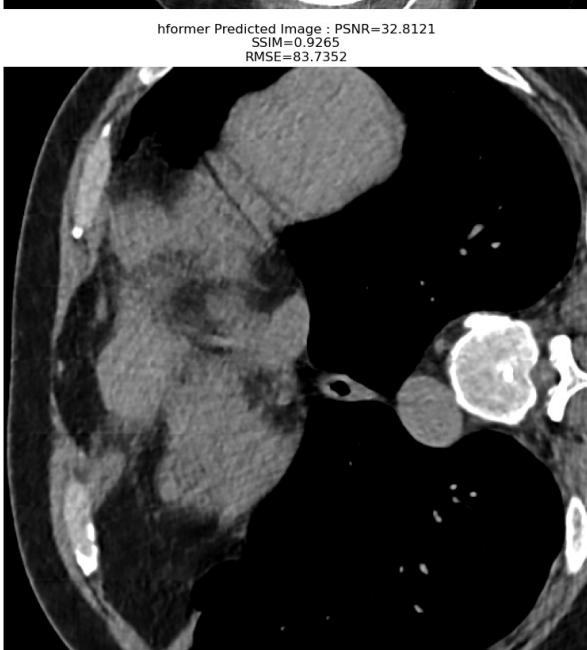
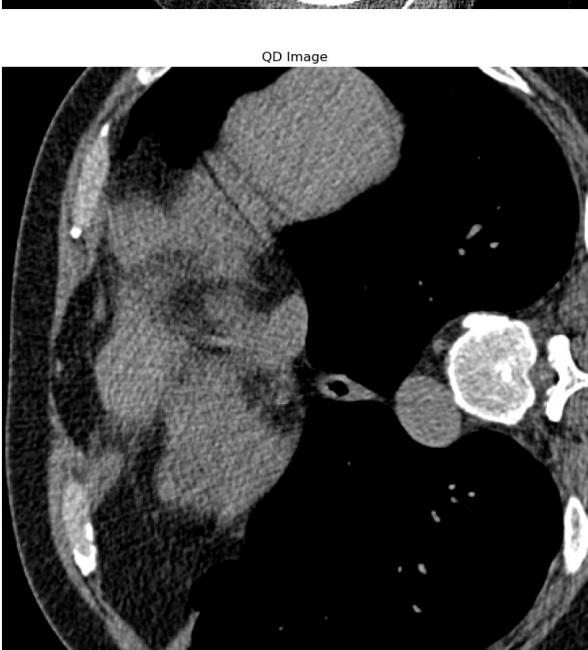
hformer_prediction_patches = hformer_model.predict(noisy_image_patches_array)
hformer_predictions = np.expand_dims(reconstruct_image_from_patches(hformer_prediction_patches[0:64], 8))

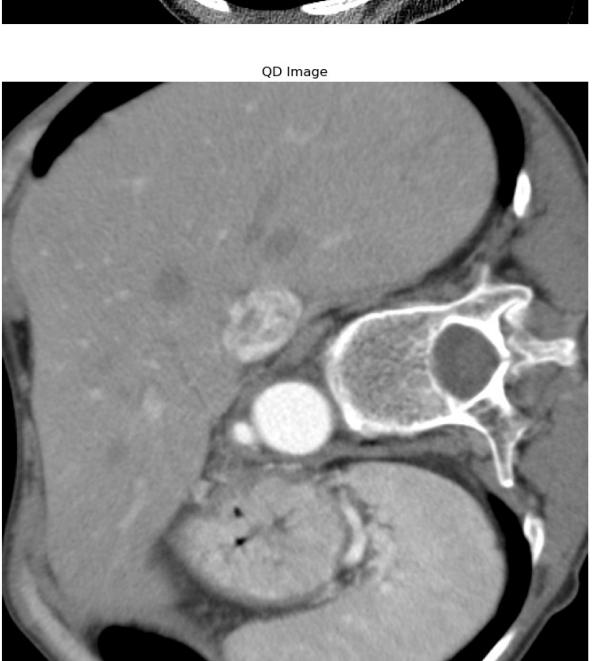
for i in range(1, int(hformer_prediction_patches.shape[0] / 64)):
    reconstructed_image = reconstruct_image_from_patches(hformer_prediction_patches[i * 64 : i * 64 + 64])
    reconstructed_image = np.expand_dims(reconstructed_image, axis=0)

    hformer_predictions = np.append(hformer_predictions, reconstructed_image, axis=0)
visualize_predictions(hformer_predictions, noisy_array, len(noisy_array), hformer_predictions, "hformer"
56/56 [=====] - 57s 980ms/step

```









QD Image



QD Image



QD Image



hformer Predicted Image : PSNR=35.6345  
SSIM=0.9714  
RMSE=43.7192



hformer Predicted Image : PSNR=32.3595  
SSIM=0.8886  
RMSE=92.9333



hformer Predicted Image : PSNR=31.1293  
SSIM=0.902  
RMSE=123.3631





QD Image



hformer Predicted Image : PSNR=19.1266  
SSIM=0.5663  
RMSE=1956.3892



QD Image



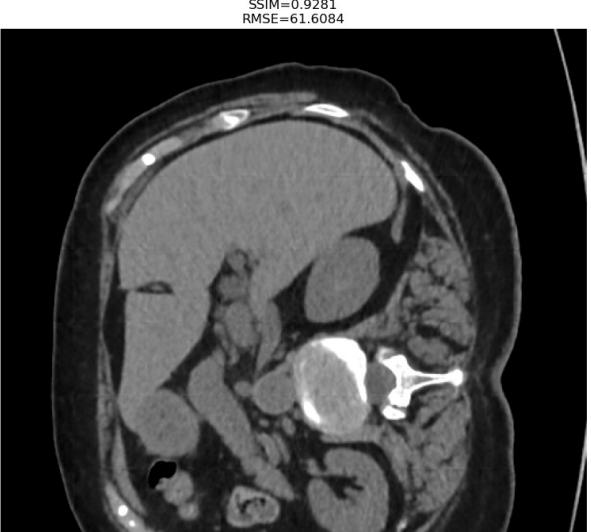
hformer Predicted Image : PSNR=34.4729  
SSIM=0.9316  
RMSE=57.126

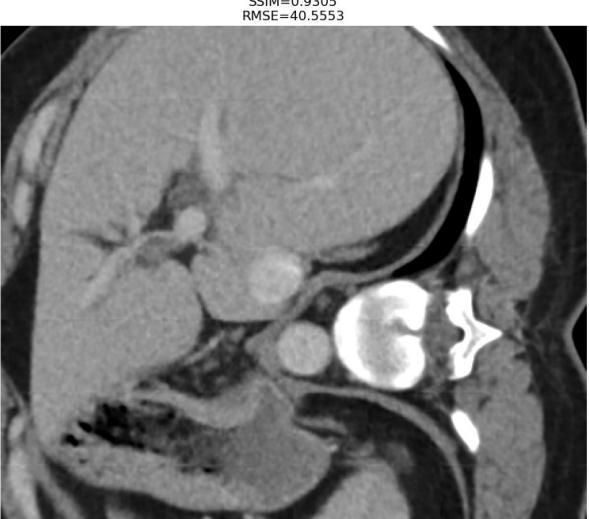


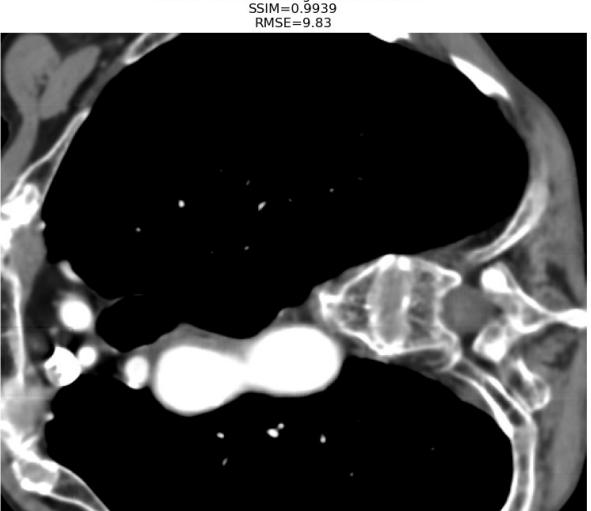
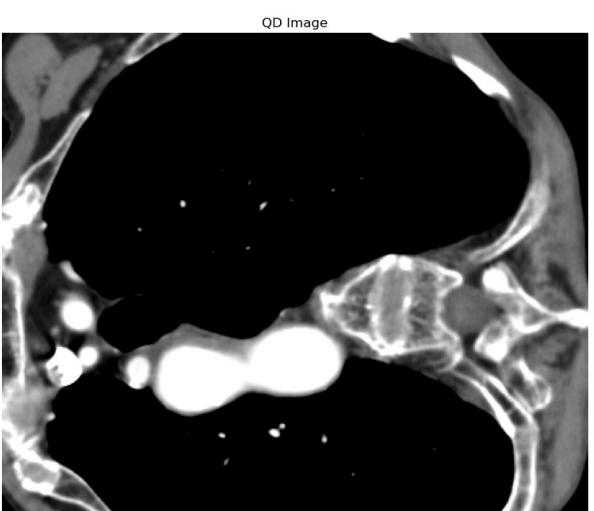
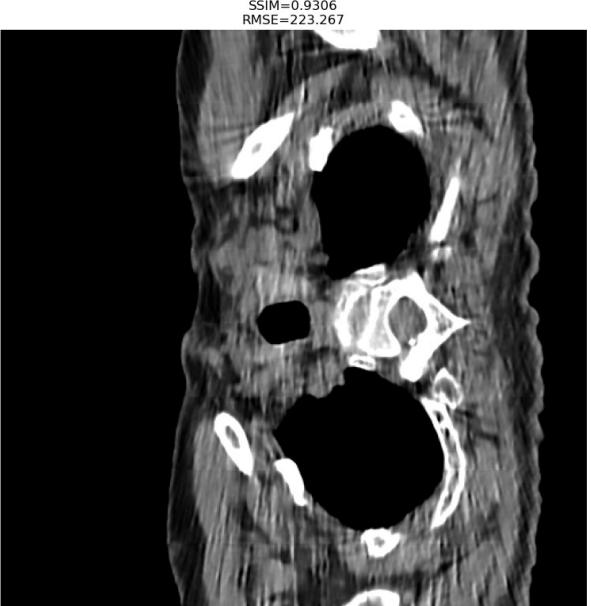
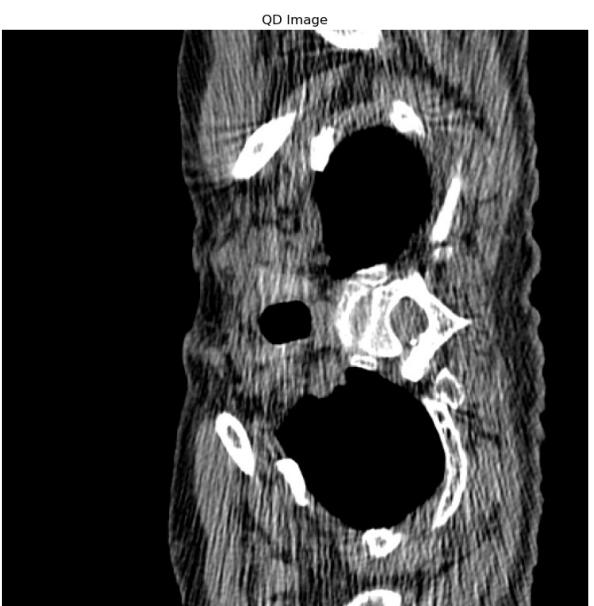
QD Image

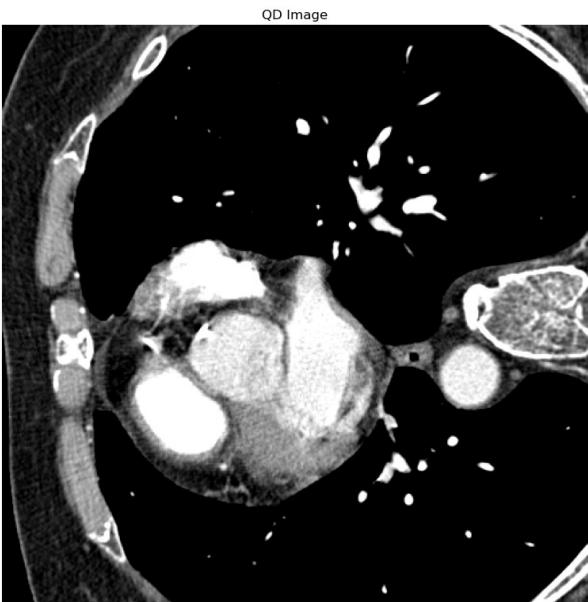


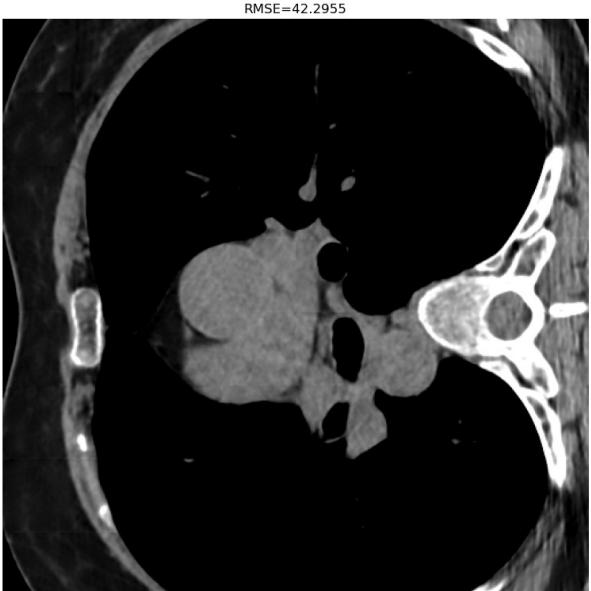
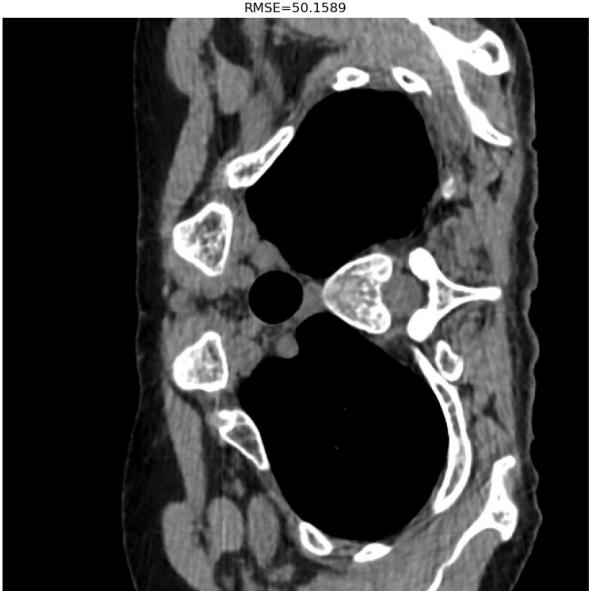
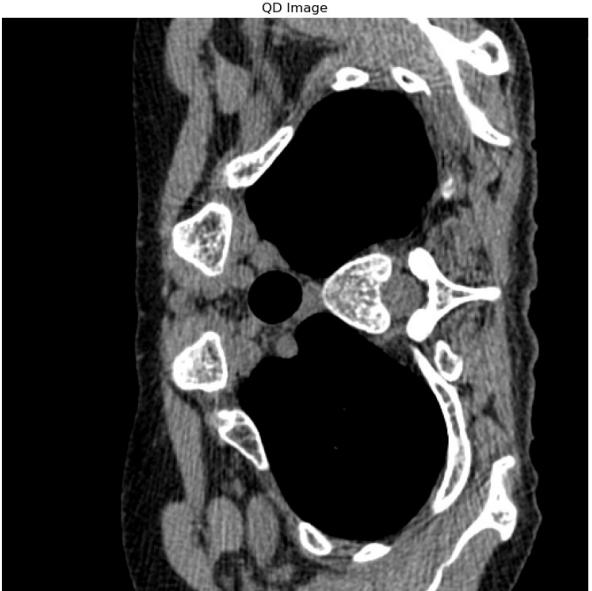
hformer Predicted Image : PSNR=34.1448  
SSIM=0.9281  
RMSE=61.6084













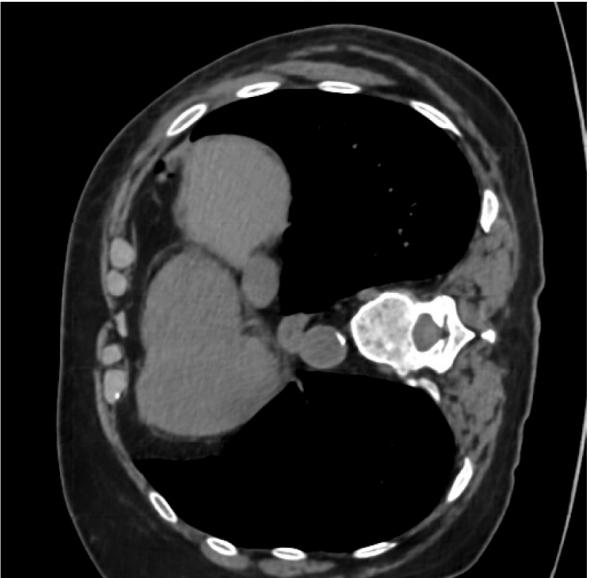
QD Image



hformer Predicted Image : PSNR=37.1773  
SSIM=0.9624  
RMSE=30.6474



QD Image



hformer Predicted Image : PSNR=26.9225  
SSIM=0.8118  
RMSE=324.9903

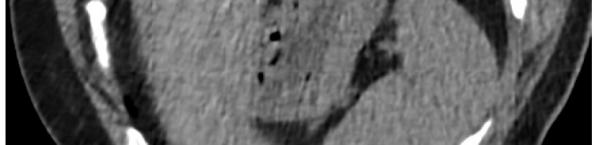
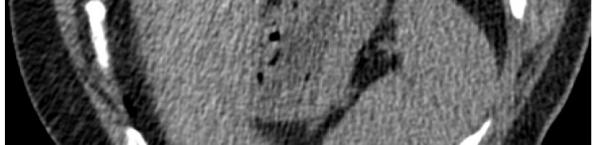


QD Image



hformer Predicted Image : PSNR=29.2817  
SSIM=0.8456  
RMSE=188.7754

























## ▼ Model 2 : XModel

```
from torchinfo import summary
sys.path.append('../denoising-models/novel_model/')
from x_denoiser import XModel
x_model = XModel(num_channels=64)
x_model.load_state_dict(torch.load('../denoising-models/novel_model/weights/x_model/weights/x_model/model_19.pth'))
x_model.eval()
print('Model summary : ')
print(summary(x_model))
```

Model summary :

Layer (type:depth-idx)	Param #
XModel	--
---ConvTranspose2d: 1-1	5
---Conv2d: 1-2	16,448
---InputProjectionLayer: 1-3	--
└---Conv2d: 2-1	640
---OutputProjectionLayer: 1-4	--
└---ConvTranspose2d: 2-2	577
---ConvBlock: 1-5	--
└---ZeroPad2d: 2-3	--
└---DepthWiseSeparableConv2d: 2-4	--
└---Conv2d: 3-1	3,136
└---Conv2d: 3-2	4,096
└---LayerNorm: 2-5	--
└---Conv2d: 2-6	803,072
└---GELU: 2-7	--
└---Conv2d: 2-8	16,448
---ConvBlock: 1-6	--
└---ZeroPad2d: 2-9	--
└---DepthWiseSeparableConv2d: 2-10	--
└---Conv2d: 3-3	3,136
└---Conv2d: 3-4	4,096
└---LayerNorm: 2-11	--
└---Conv2d: 2-12	803,072
└---GELU: 2-13	--

└Conv2d: 2-14	16,448
ConvBlock: 1-7	--
└ZeroPad2d: 2-15	--
└DepthWiseSeparableConv2d: 2-16	--
└Conv2d: 3-5	3,136
└Conv2d: 3-6	4,096
└LayerNorm: 2-17	--
└Conv2d: 2-18	803,072
└GELU: 2-19	--
└Conv2d: 2-20	16,448
ConvBlock: 1-8	--
└ZeroPad2d: 2-21	--
└DepthWiseSeparableConv2d: 2-22	--
└Conv2d: 3-7	3,136
└Conv2d: 3-8	4,096
└LayerNorm: 2-23	--
└Conv2d: 2-24	803,072
└GELU: 2-25	--
└Conv2d: 2-26	16,448
XBlock: 1-9	--
└Linear: 2-27	4,160
└Linear: 2-28	4,160
└Linear: 2-29	4,160
└Conv2d: 2-30	4,160
└LayerNorm: 2-31	--
└Conv2d: 2-32	16,640
└GELU: 2-33	--
└Conv2d: 2-34	16,448
ConvBlock: 1-10	--
└ZeroPad2d: 2-35	--

```
# Get prediction images.
```

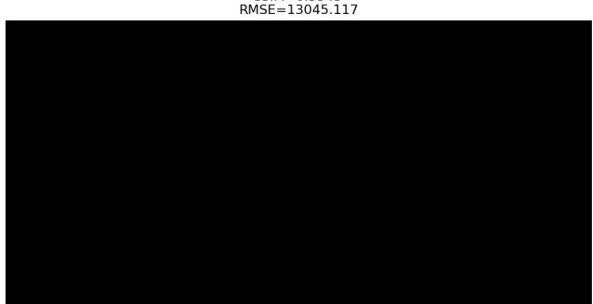
```
x_model_prediction_patches = []

for img in noisy_array:
    img_tensor = torch.from_numpy(img).float()

    with torch.no_grad():
        output_tensor = x_model(img_tensor)
        output_tensor = output_tensor.cpu().numpy()

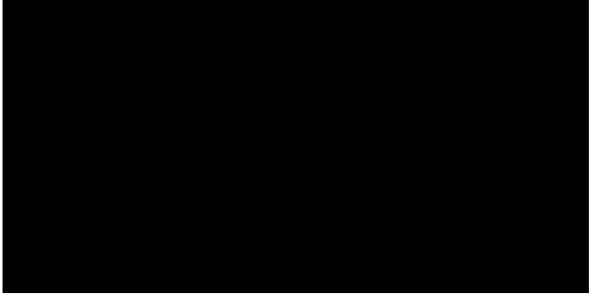
    x_model_prediction_patches.append(output_tensor)

visualize_predictions(x_model, noisy_array, len(noisy_array), np.concatenate(x_model_prediction_patches))
```





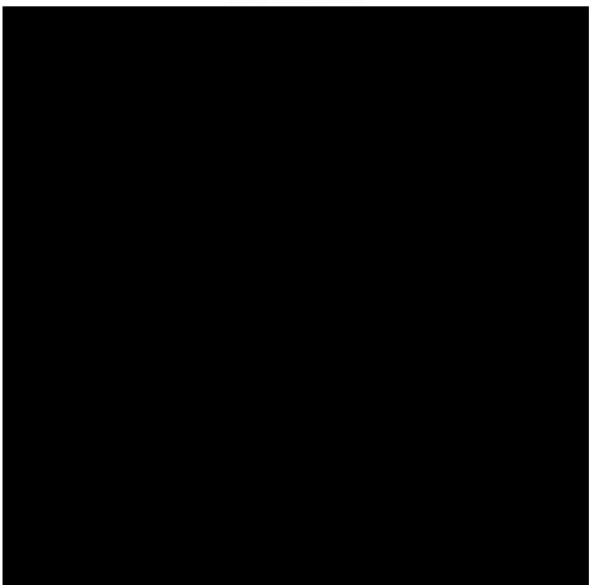
QD Image



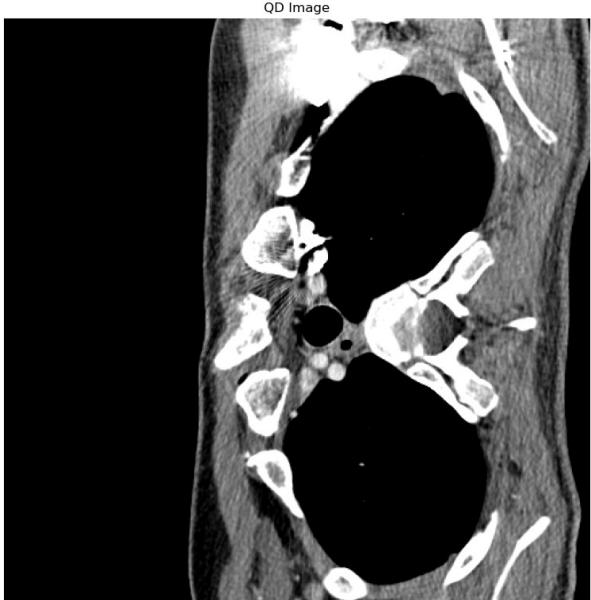
x\_model Predicted Image : PSNR=10.1793  
SSIM=0.3995  
RMSE=15352.7121



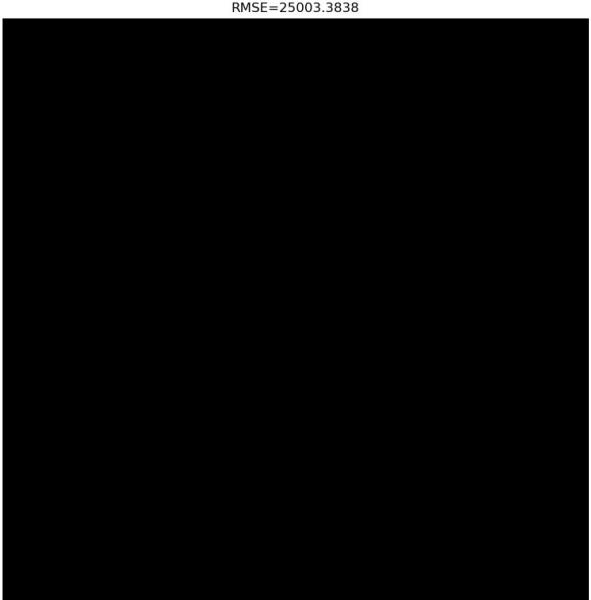
QD Image



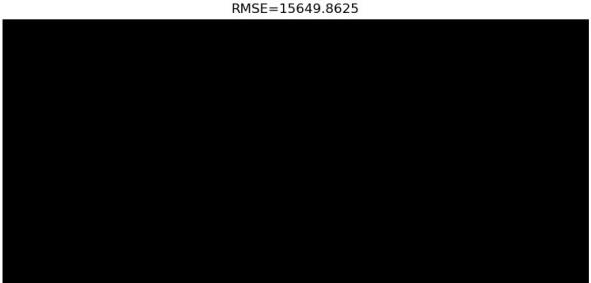
x\_model Predicted Image : PSNR=8.0612  
SSIM=0.4824  
RMSE=25003.3838



QD Image

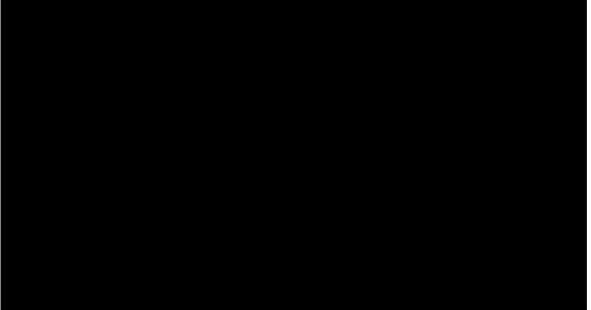


x\_model Predicted Image : PSNR=10.0961  
SSIM=0.5335  
RMSE=15649.8625





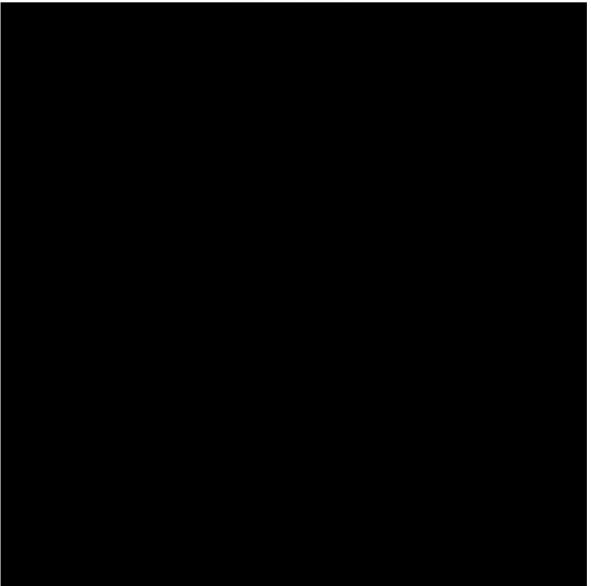
QD Image



x\_model Predicted Image : PSNR=5.2602  
SSIM=-0.2802  
RMSE=47654.2589



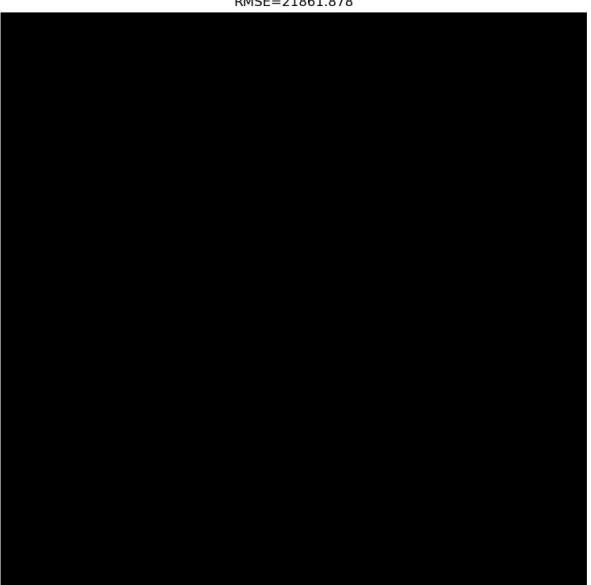
QD Image



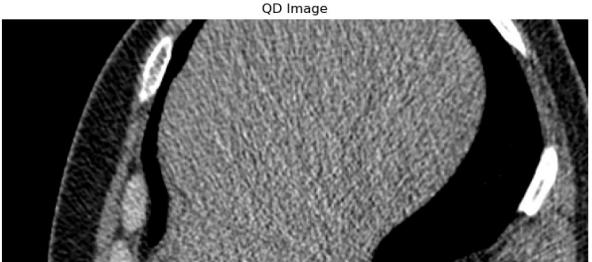
x\_model Predicted Image : PSNR=8.6443  
SSIM=0.4223  
RMSE=21861.878



QD Image

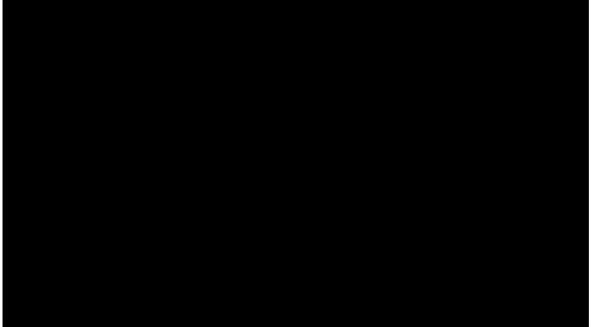


x\_model Predicted Image : PSNR=7.6802  
SSIM=0.1716  
RMSE=27296.2699





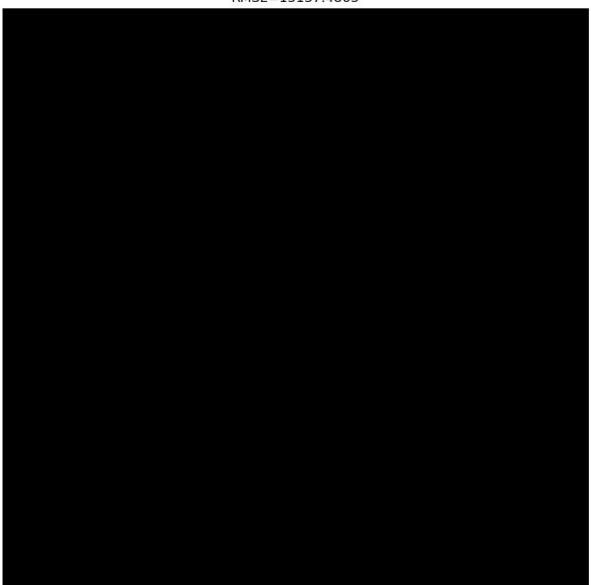
QD Image



x\_model Predicted Image : PSNR=10.8561  
SSIM=0.5422  
RMSE=13137.4863



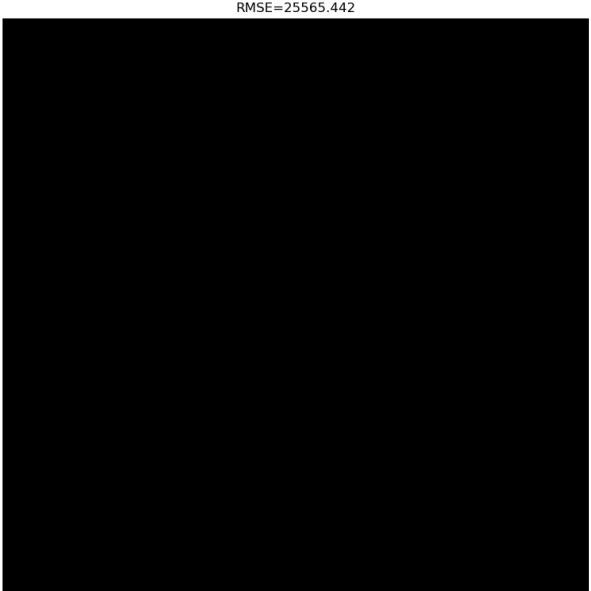
QD Image



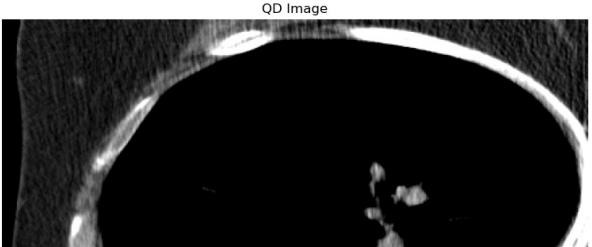
x\_model Predicted Image : PSNR=7.9647  
SSIM=0.0918  
RMSE=25565.442

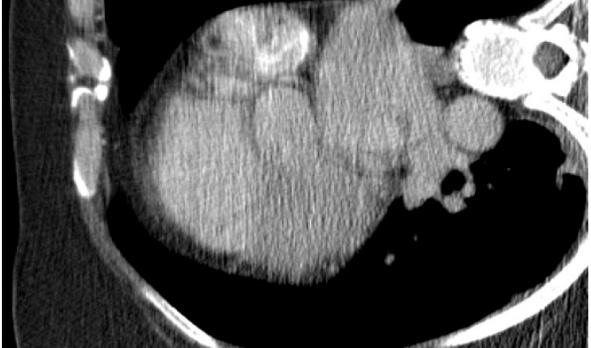


QD Image

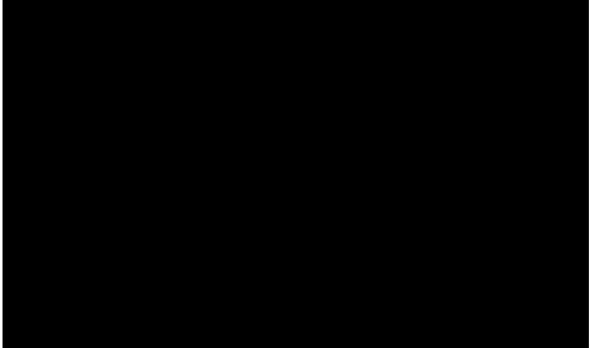


x\_model Predicted Image : PSNR=8.3301  
SSIM=0.4146  
RMSE=23502.0094





QD Image



x\_model Predicted Image : PSNR=6.5611  
SSIM=0.0194  
RMSE=35318.9118



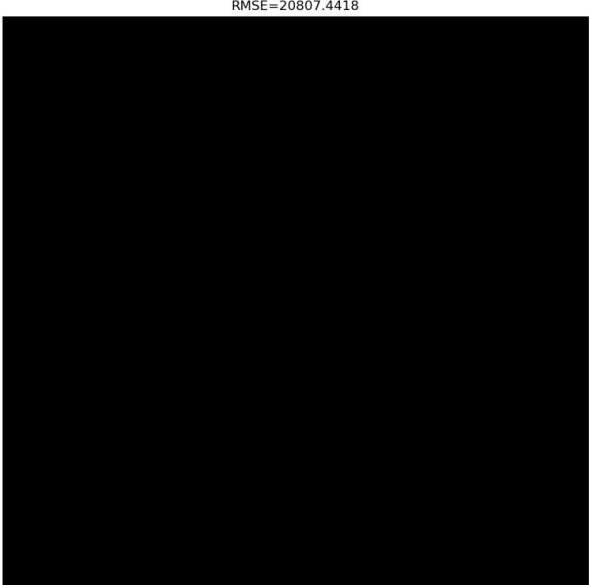
QD Image



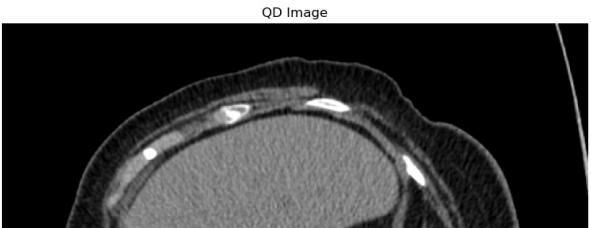
x\_model Predicted Image : PSNR=8.859  
SSIM=0.3801  
RMSE=20807.4418

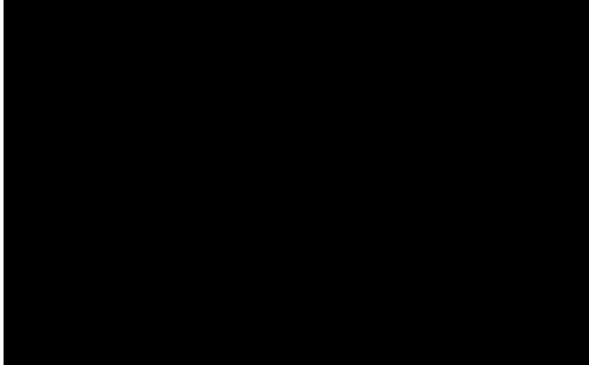


QD Image



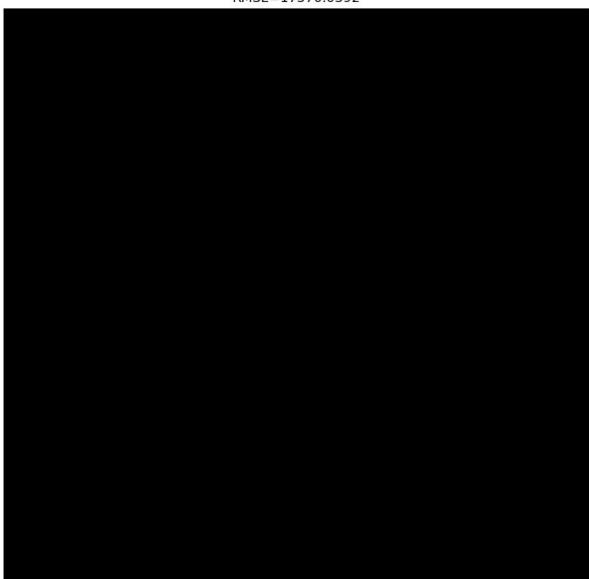
x\_model Predicted Image : PSNR=10.2639  
SSIM=0.3265  
RMSE=15056.7882





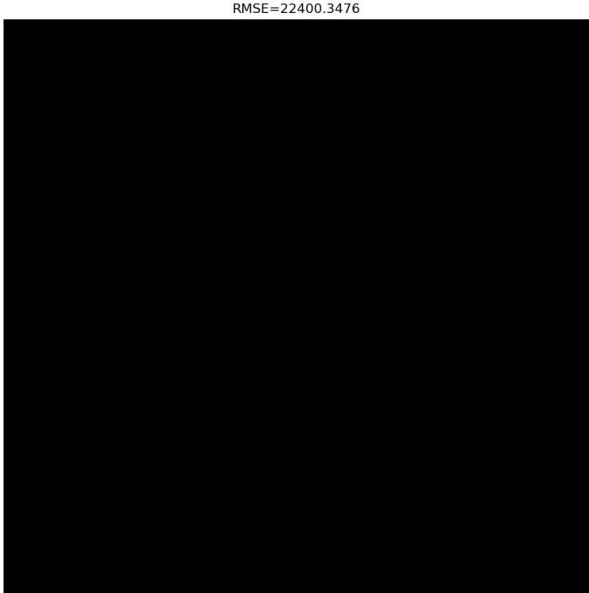
QD Image

x\_model Predicted Image : PSNR=9.5935  
SSIM=0.2867  
RMSE=17570.0392



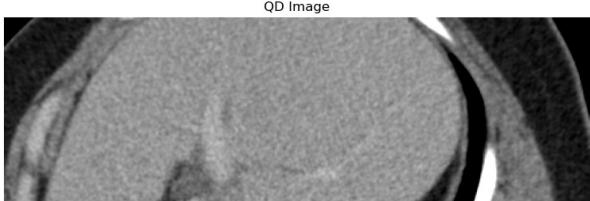
QD Image

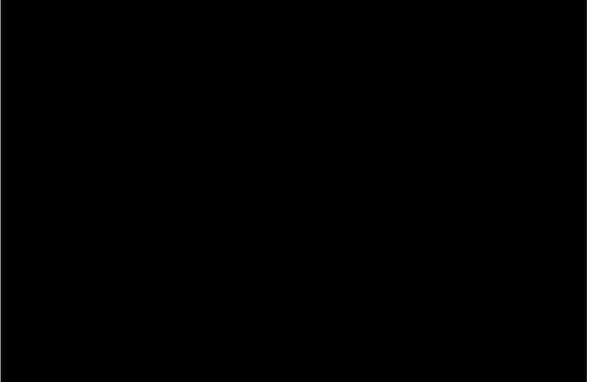
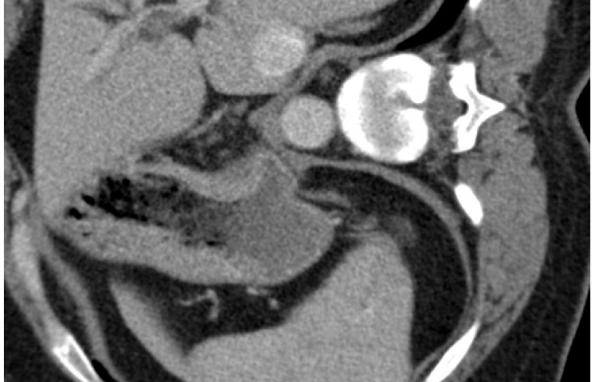
x\_model Predicted Image : PSNR=8.5387  
SSIM=0.379  
RMSE=22400.3476



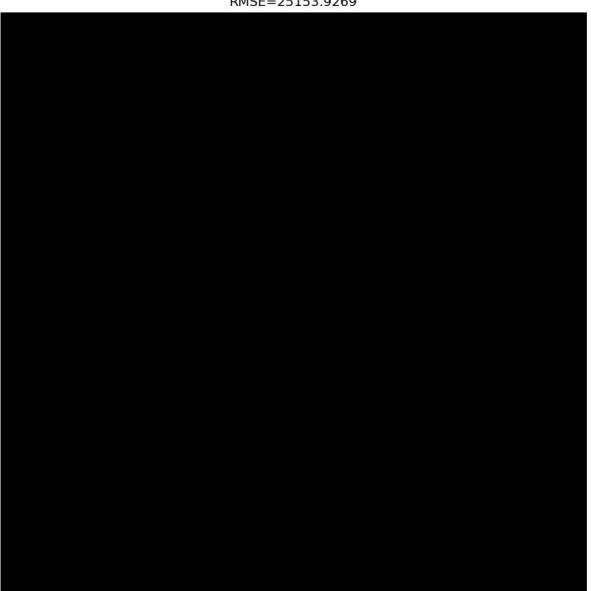
QD Image

x\_model Predicted Image : PSNR=6.4351  
SSIM=-0.0604  
RMSE=36358.9252



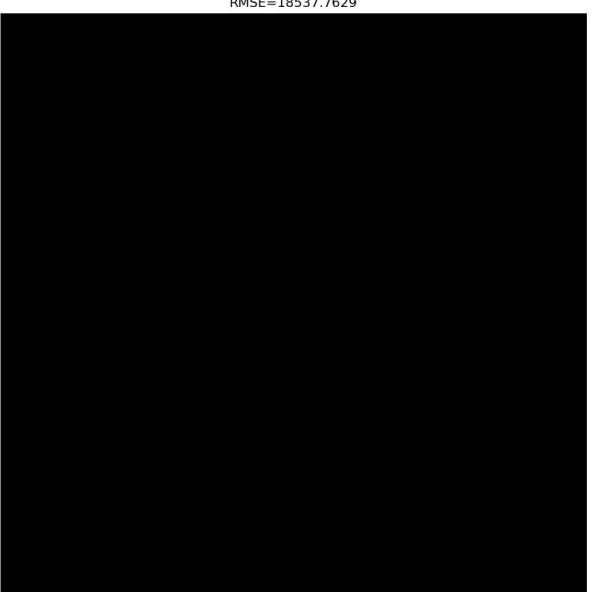
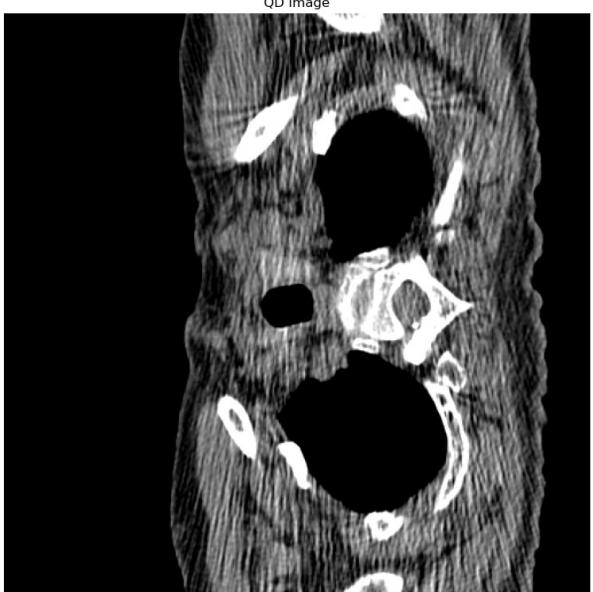


QD Image



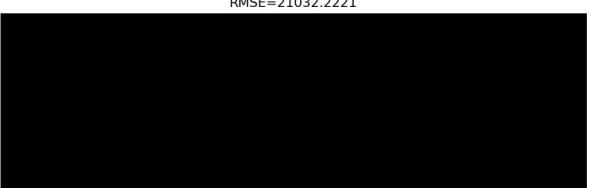
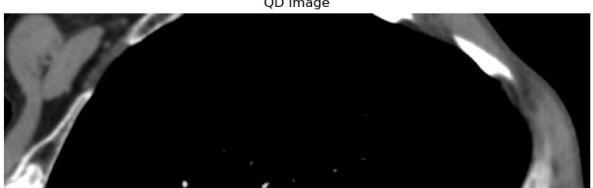
x\_model Predicted Image : PSNR=8.0351  
SSIM=0.3918  
RMSE=25153.9269

QD Image

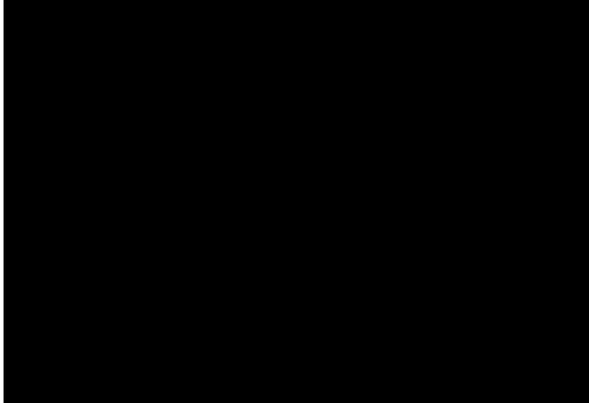
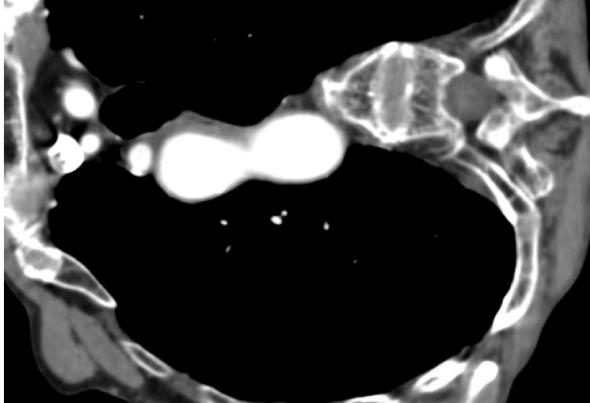


x\_model Predicted Image : PSNR=9.3606  
SSIM=0.4486  
RMSE=18537.7629

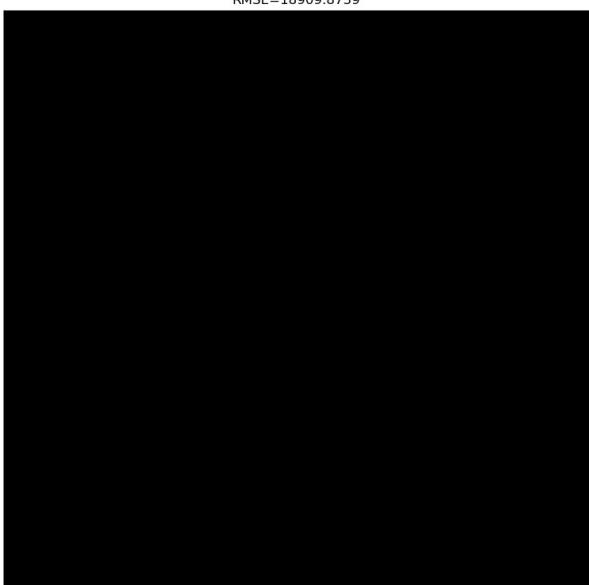
QD Image



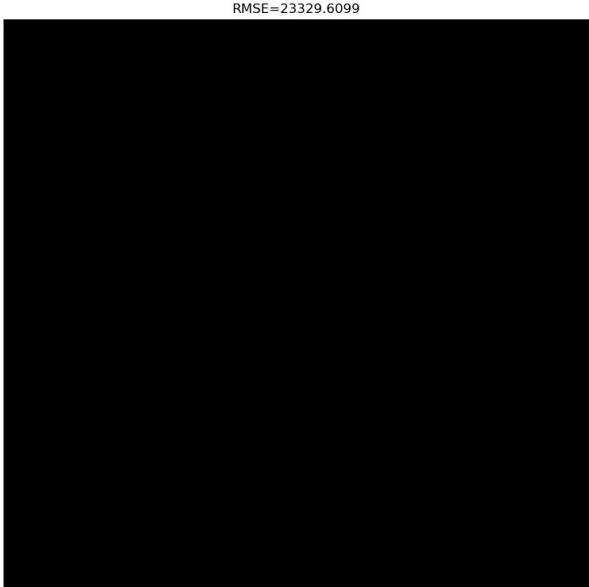
x\_model Predicted Image : PSNR=8.8123  
SSIM=0.5345  
RMSE=21032.2221



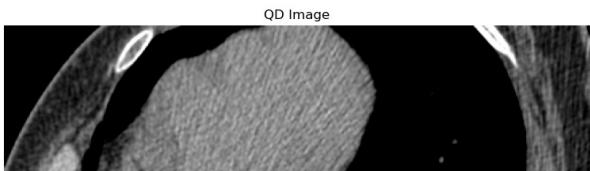
x\_model Predicted Image : PSNR=9.2743  
SSIM=0.4299  
RMSE=18909.8739

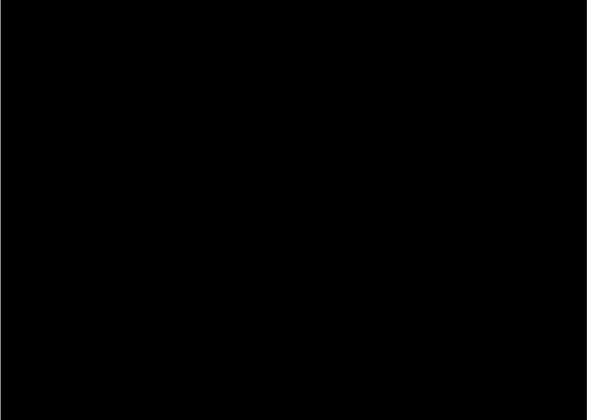
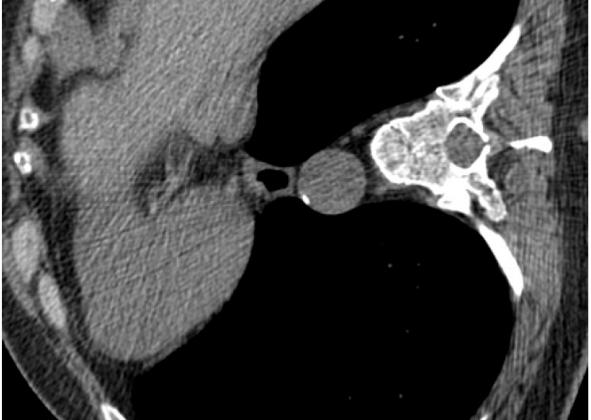


x\_model Predicted Image : PSNR=8.3621  
SSIM=0.4852  
RMSE=23329.6099

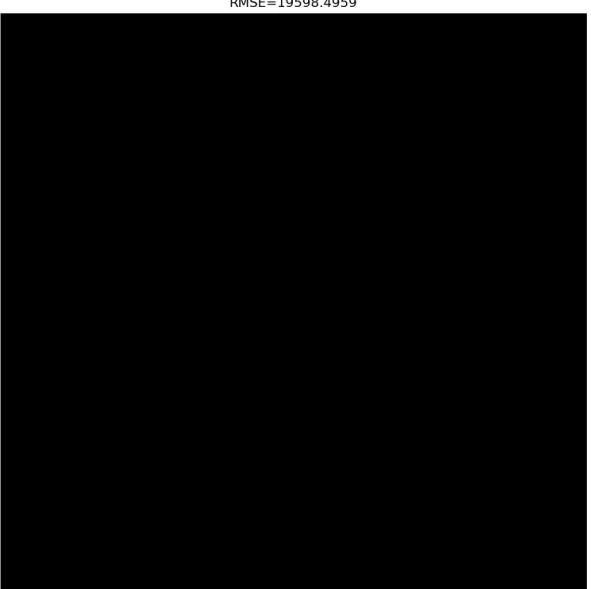
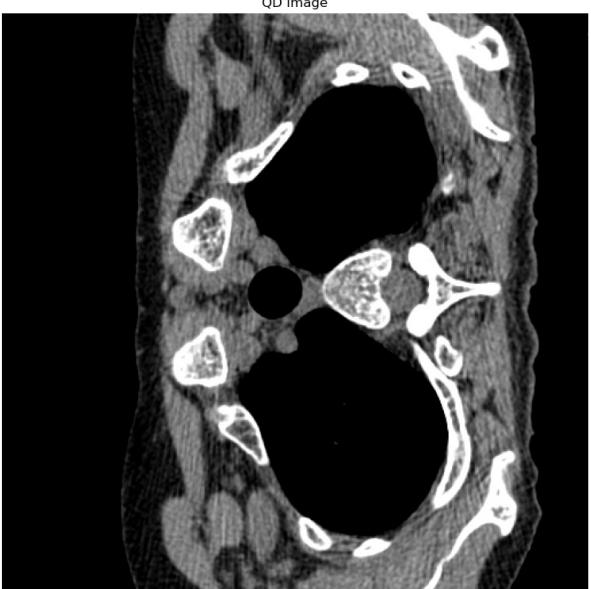


x\_model Predicted Image : PSNR=8.7448  
SSIM=0.3217  
RMSE=21361.9342

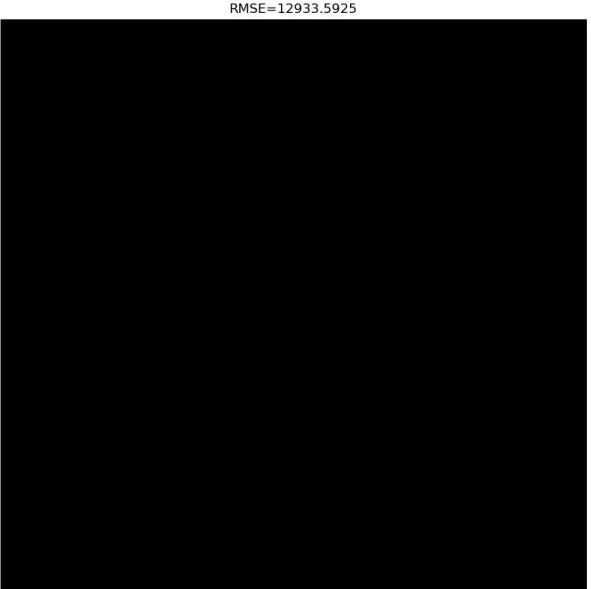




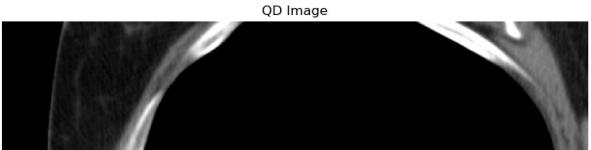
QD Image



QD Image



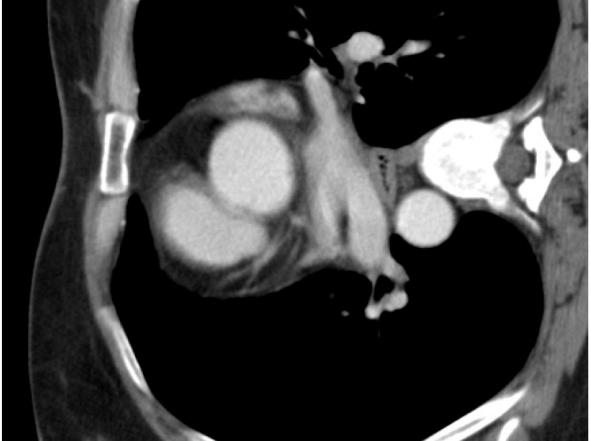
QD Image



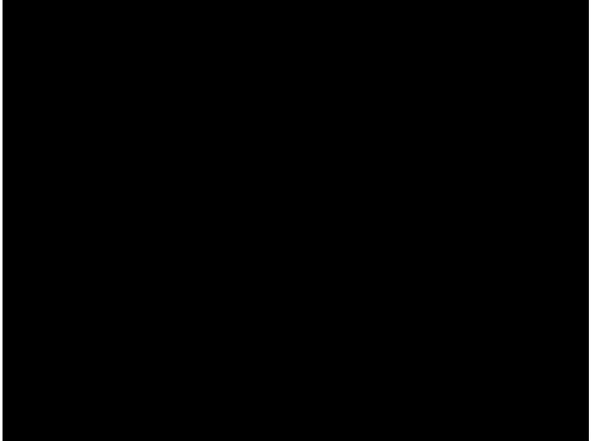
x\_model Predicted Image : PSNR=9.119  
SSIM=0.4541  
RMSE=19598.4959

x\_model Predicted Image : PSNR=10.924  
SSIM=0.5625  
RMSE=12933.5925

x\_model Predicted Image : PSNR=9.4163  
SSIM=0.5015  
RMSE=18301.7504



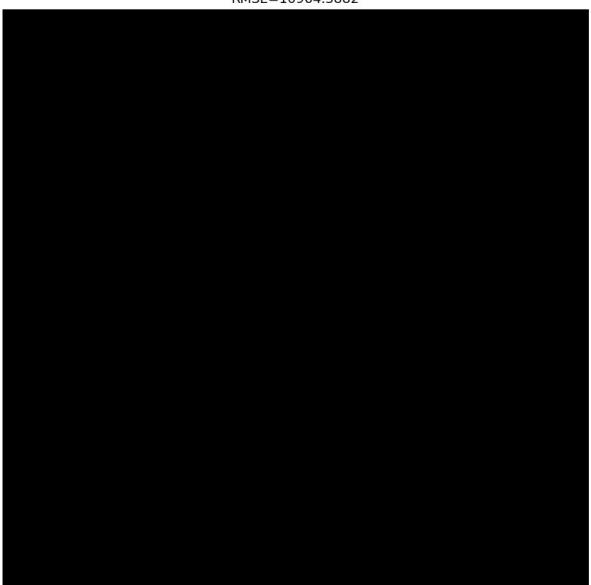
QD Image



x\_model Predicted Image : PSNR=11.6414  
SSIM=0.4954  
RMSE=10964.3882



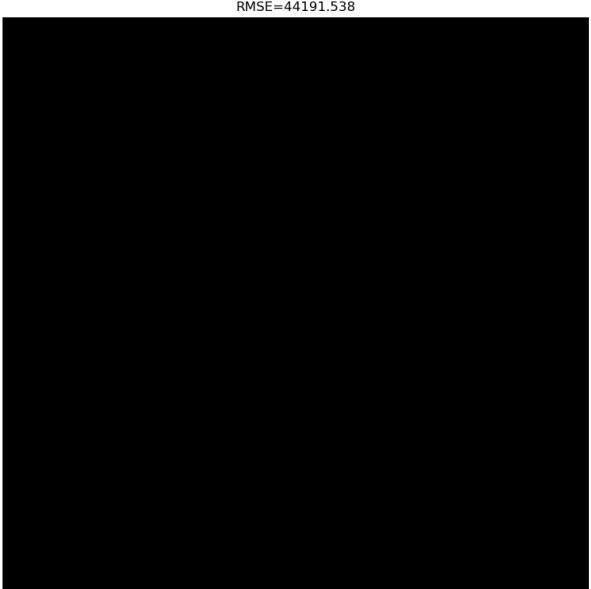
QD Image



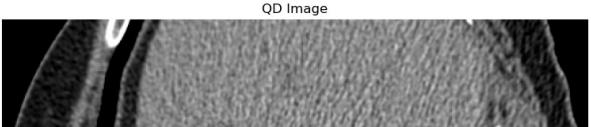
x\_model Predicted Image : PSNR=5.5878  
SSIM=0.0054  
RMSE=44191.538

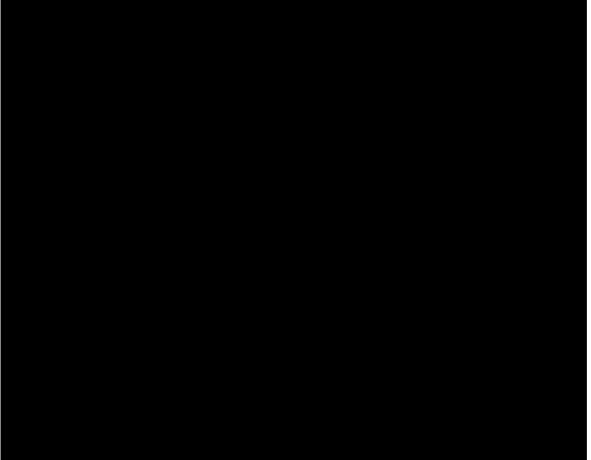


QD Image



x\_model Predicted Image : PSNR=6.8974  
SSIM=0.0068  
RMSE=32687.452

























## ▼ Model 3 : Y model

```
sys.path.append('../denoising-models/mwcnn/')
from y_denoiser import YModel
y_model = YModel(num_channels=56)
y_model.load_state_dict(torch.load('../denoising-models/novel_model/weights/y_model/model_24.pth'))
y_model.eval()
print('Model summary : ')
print(summary(y_model))

Model summary :
=====
Layer (type:depth-idx)           Param #
=====
YModel
├─InputProjectionLayer: 1-1      --
│  └─Conv2d: 2-1                560
├─OutputProjectionLayer: 1-2     --
│  └─ConvTranspose2d: 2-2        505
├─ConvBlock: 1-3                 --
│  └─ZeroPad2d: 2-3             --
│  └─DepthWiseSeparableConv2d: 2-4
│    └─Conv2d: 3-1              2,744
│    └─Conv2d: 3-2              3,136
│  └─LayerNorm: 2-5             --
│  └─Conv2d: 2-6                614,880
│  └─GELU: 2-7                 --
│  └─Conv2d: 2-8                12,600
└─YBlock: 1-4                  --
```

└ Linear: 2-9	3,192
└ Linear: 2-10	3,192
└ Linear: 2-11	3,192
└ Conv2d: 2-12	3,192
└ LayerNorm: 2-13	--
└ Conv2d: 2-14	12,768
└ GELU: 2-15	--
└ Conv2d: 2-16	12,600
Conv2d: 1-5	25,200
ConvBlock: 1-6	--
└ ZeroPad2d: 2-17	--
└ DepthWiseSeparableConv2d: 2-18	--
└ Conv2d: 3-3	5,488
└ Conv2d: 3-4	12,544
└ LayerNorm: 2-19	--
└ Conv2d: 2-20	2,459,072
└ GELU: 2-21	--
└ Conv2d: 2-22	50,288
YBlock: 1-7	--
└ Linear: 2-23	12,656
└ Linear: 2-24	12,656
└ Linear: 2-25	12,656
└ Conv2d: 2-26	12,656
└ LayerNorm: 2-27	--
└ Conv2d: 2-28	50,624
└ GELU: 2-29	--
└ Conv2d: 2-30	50,288
Conv2d: 1-8	100,576
ConvBlock: 1-9	--
└ ZeroPad2d: 2-31	--
└ DepthWiseSeparableConv2d: 2-32	--
└ Conv2d: 3-5	10,976
└ Conv2d: 3-6	50,176
└ LayerNorm: 2-33	--
└ Conv2d: 2-34	9,835,392
└ GELU: 2-35	--
└ Conv2d: 2-36	200,928
YBlock: 1-10	--
└ Linear: 2-37	50,400

```
# Get prediction images.

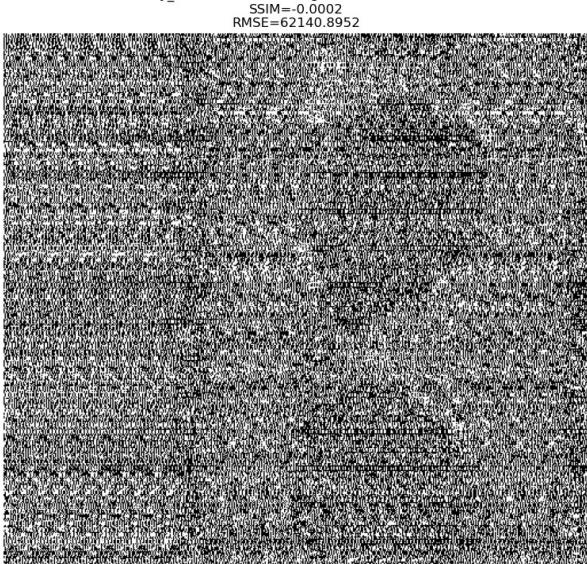
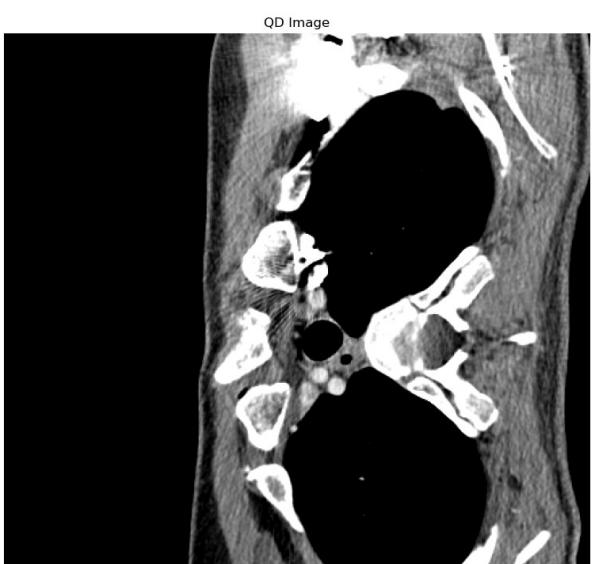
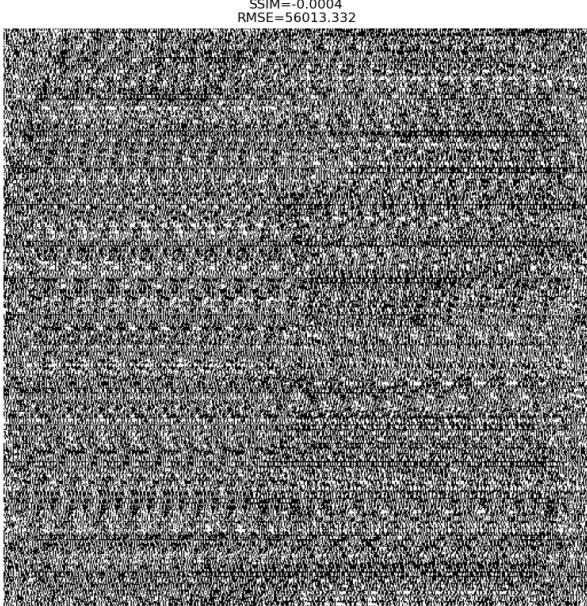
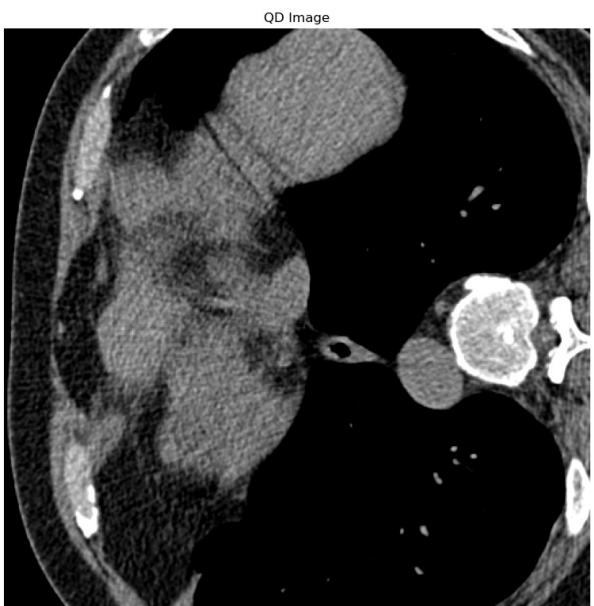
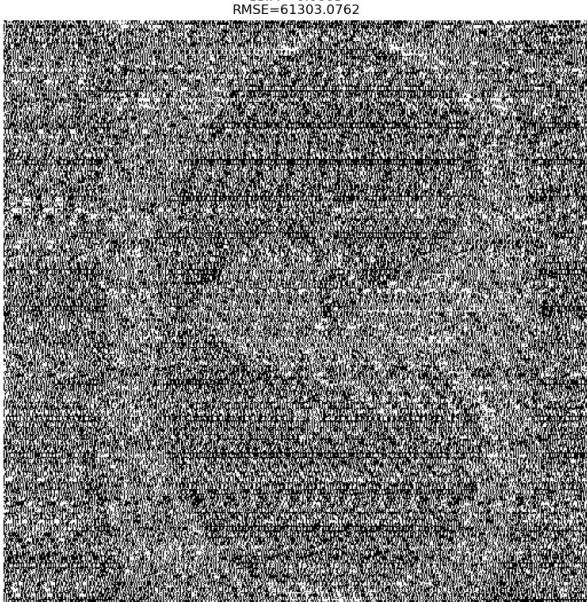
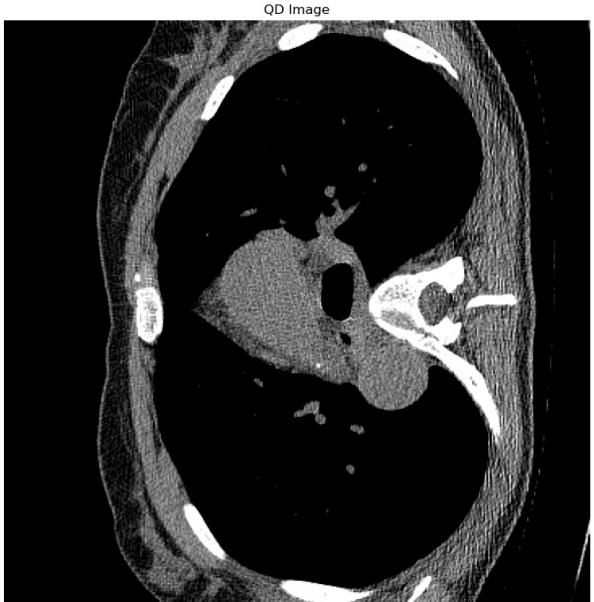
y_model_prediction_patches = []

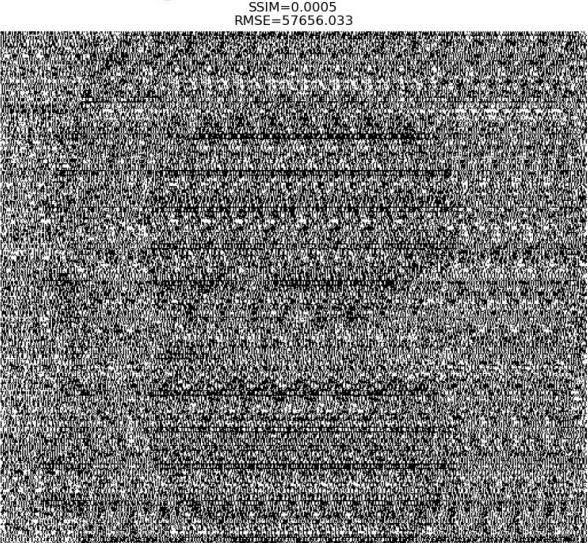
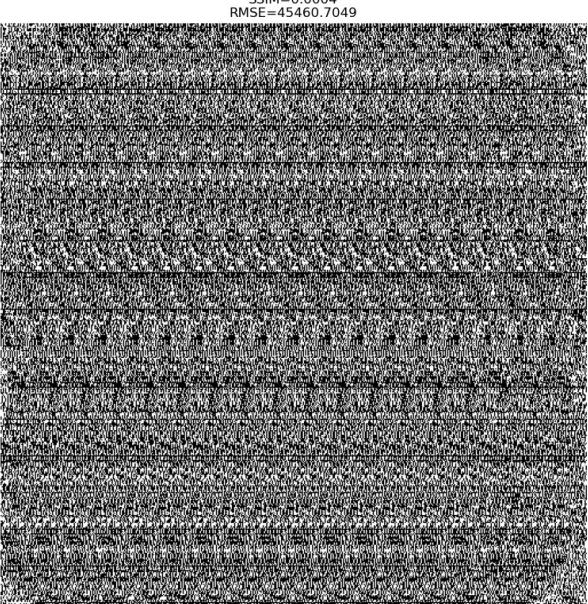
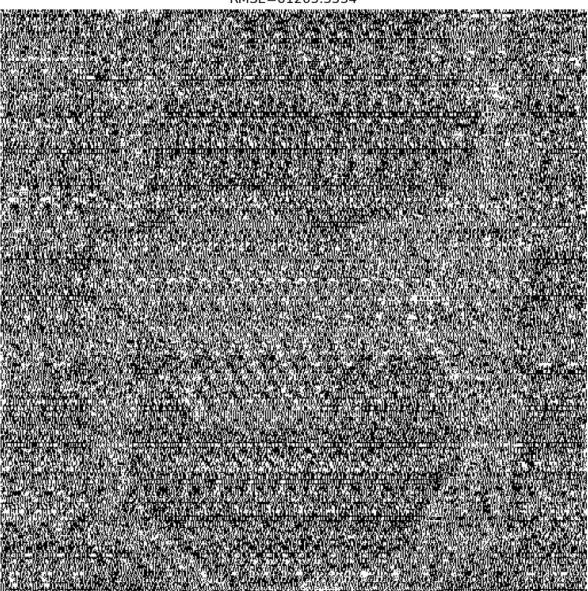
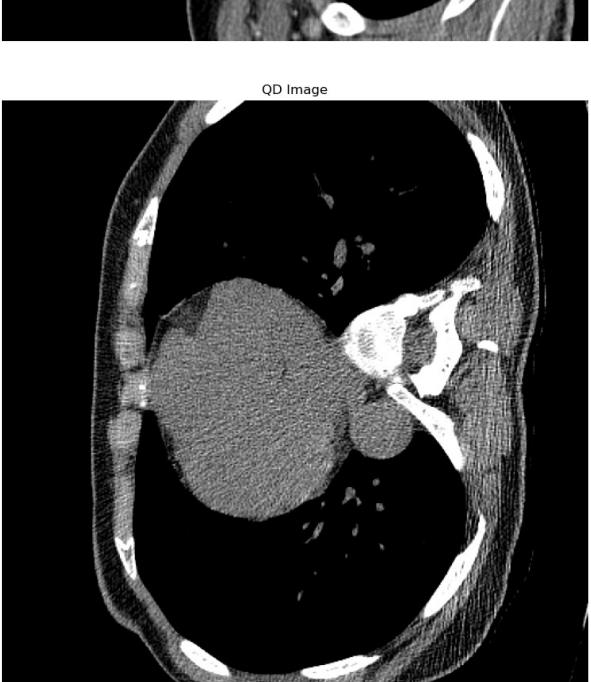
for img in noisy_array:
    img_tensor = torch.unsqueeze(torch.from_numpy(img).float(), dim=0)

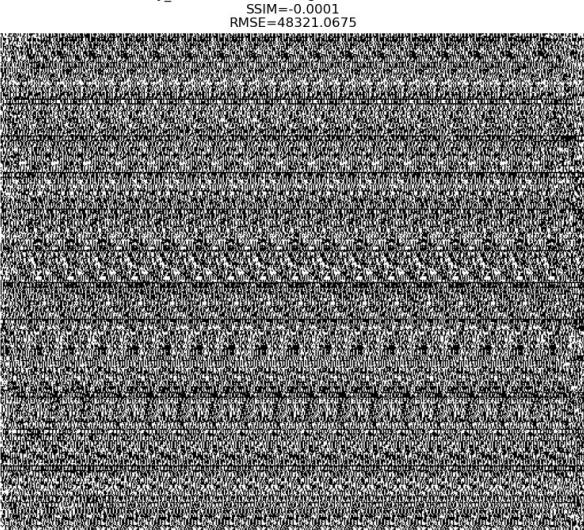
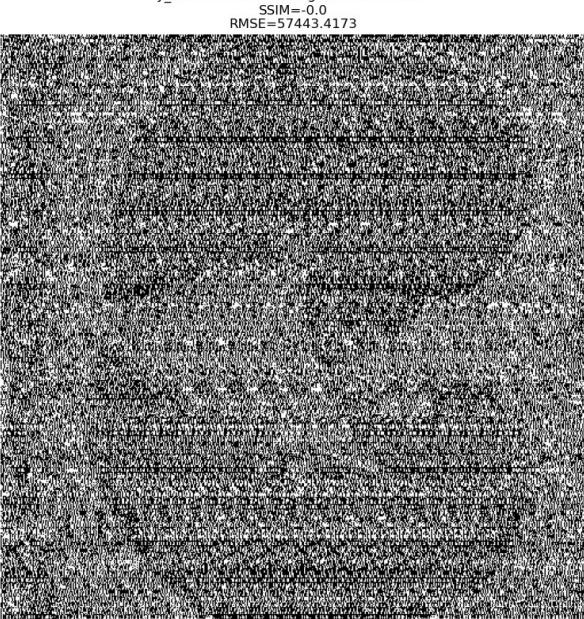
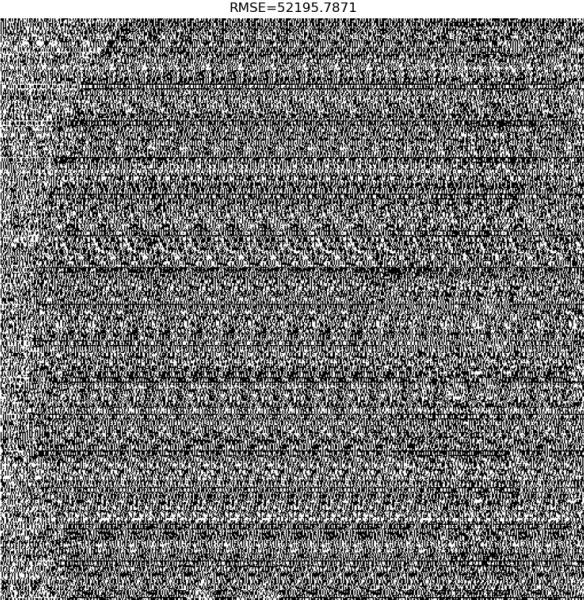
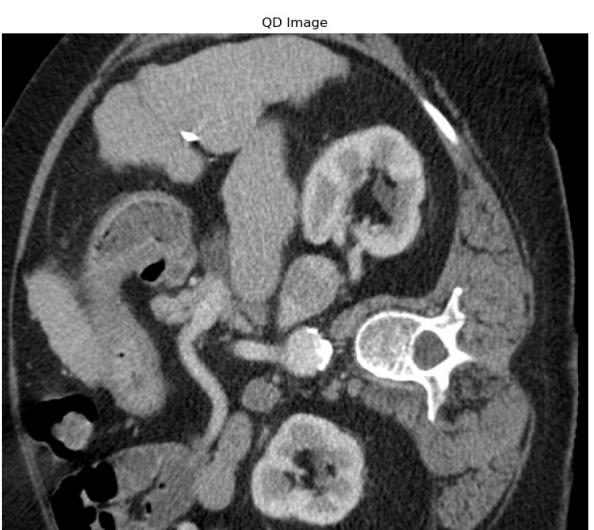
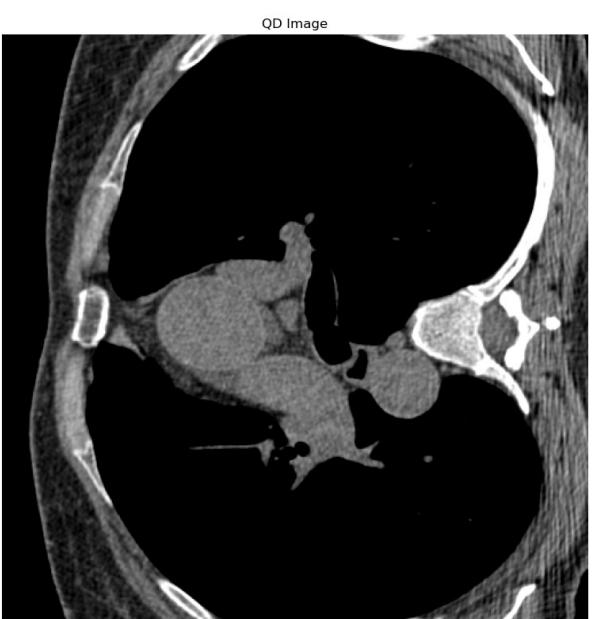
    with torch.no_grad():
        output_tensor = y_model(img_tensor)
        output_tensor = output_tensor.cpu().numpy()

    y_model_prediction_patches.append(output_tensor)
```

```
visualize_predictions(y_model, noisy_array, len(noisy_array), np.concatenate(y_model_prediction_patches
```

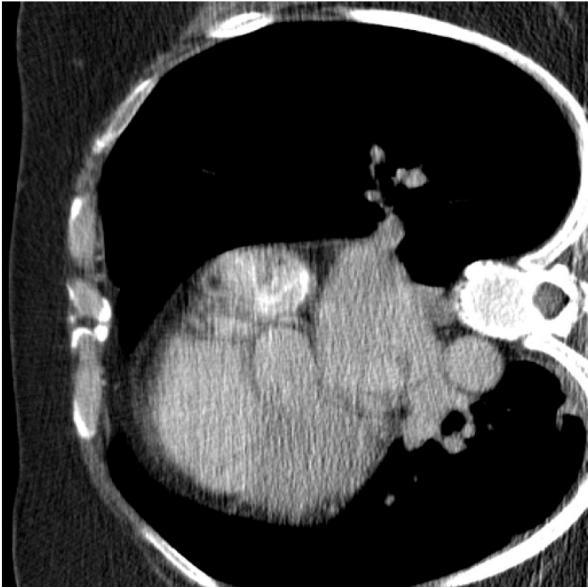








QD Image



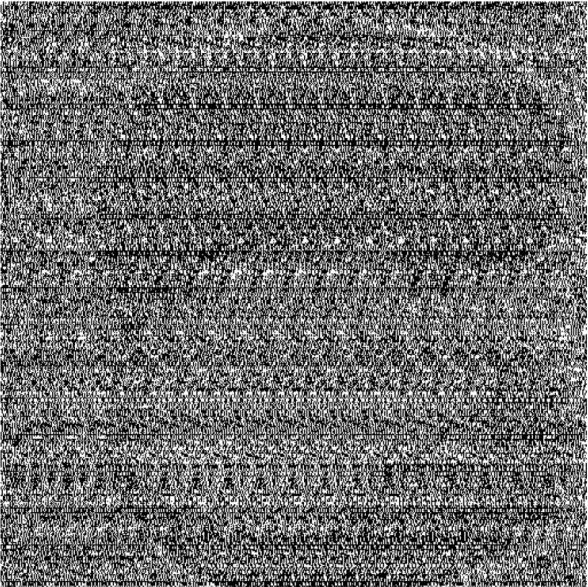
QD Image



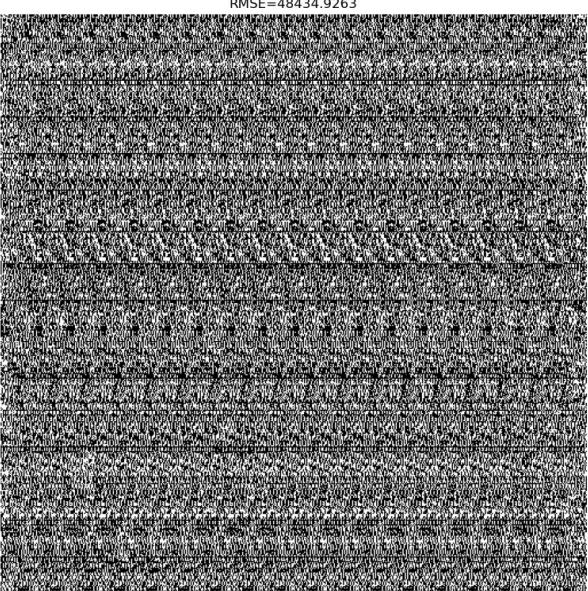
QD Image



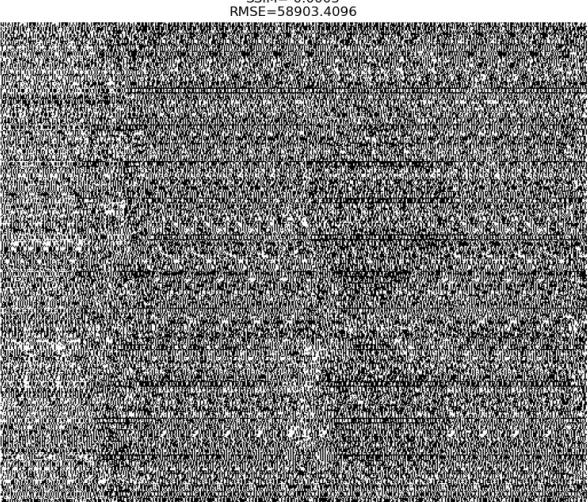
y\_model Predicted Image : PSNR=4.5733  
SSIM=0.0004  
RMSE=55819.5905



y\_model Predicted Image : PSNR=5.1896  
SSIM=0.0002  
RMSE=48434.9263

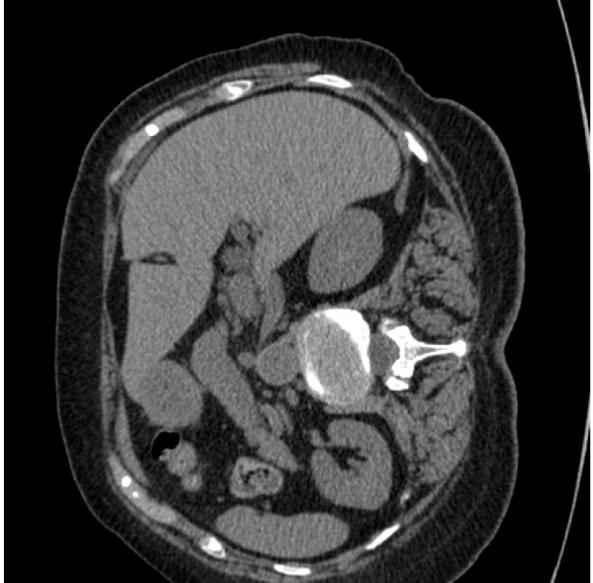


y\_model Predicted Image : PSNR=4.3398  
SSIM=-0.0003  
RMSE=58903.4096





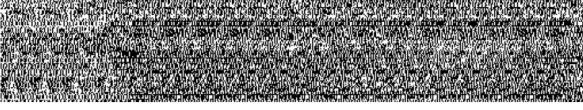
QD Image



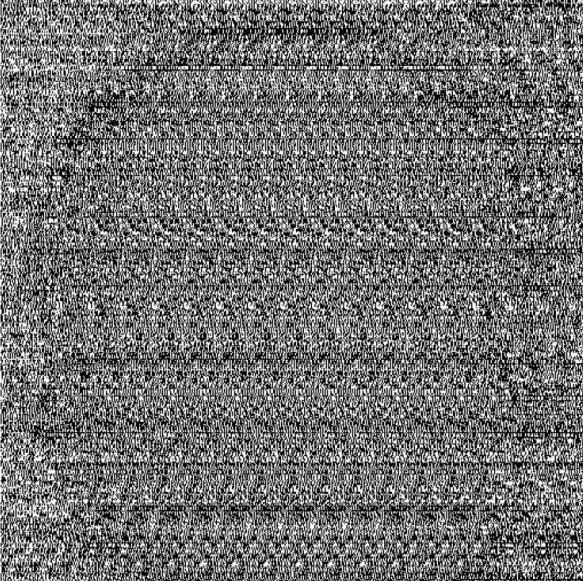
QD Image



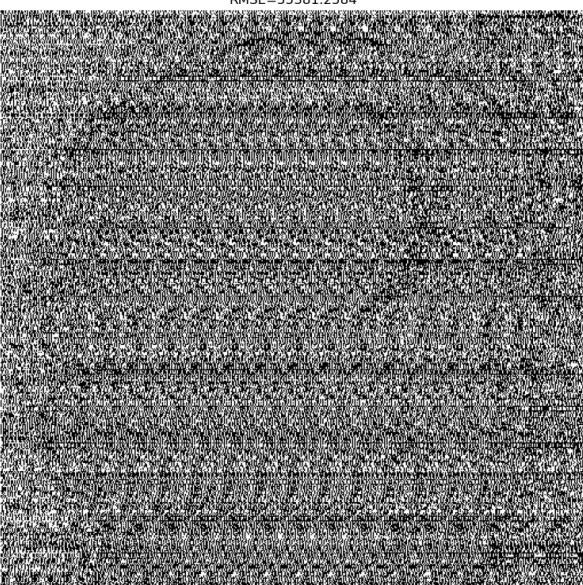
QD Image



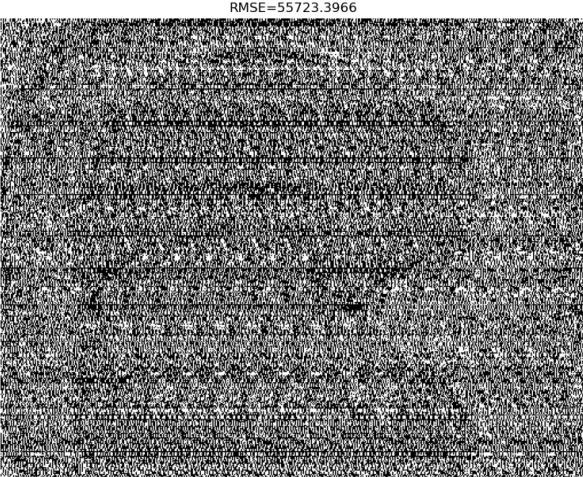
y\_model Predicted Image : PSNR=4.6127  
SSIM=-0.0001  
RMSE=55316.1083

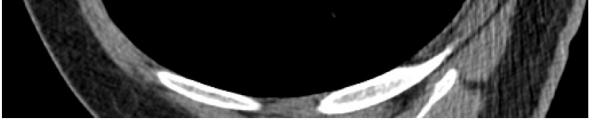


y\_model Predicted Image : PSNR=4.6076  
SSIM=0.0001  
RMSE=55381.2584

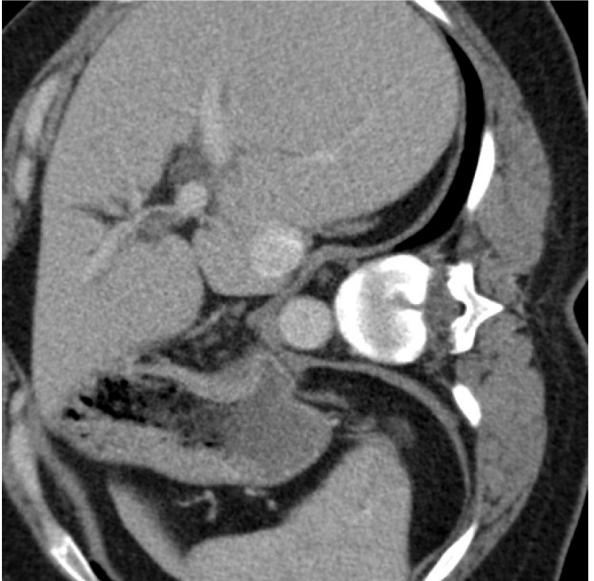


y\_model Predicted Image : PSNR=4.5808  
SSIM=0.0004  
RMSE=55723.3966





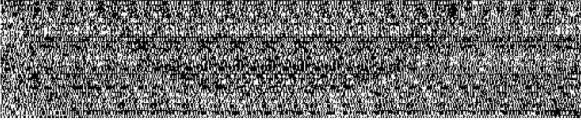
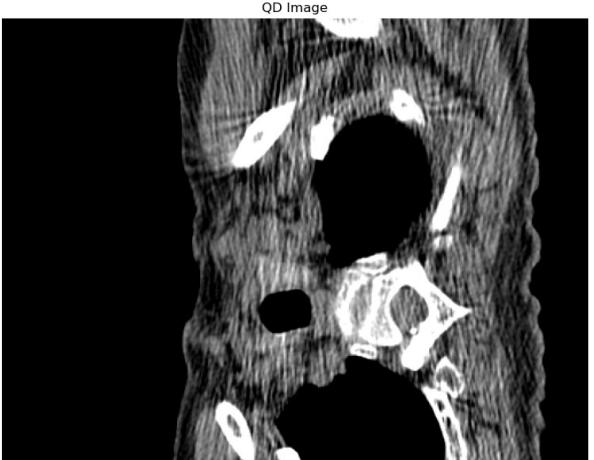
QD Image



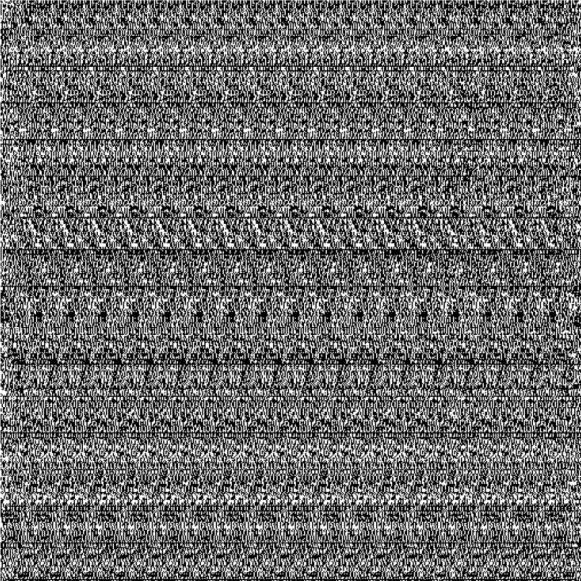
QD Image



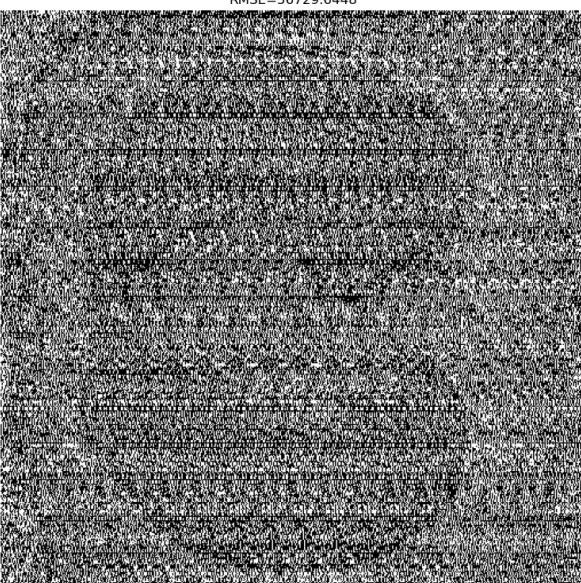
QD Image



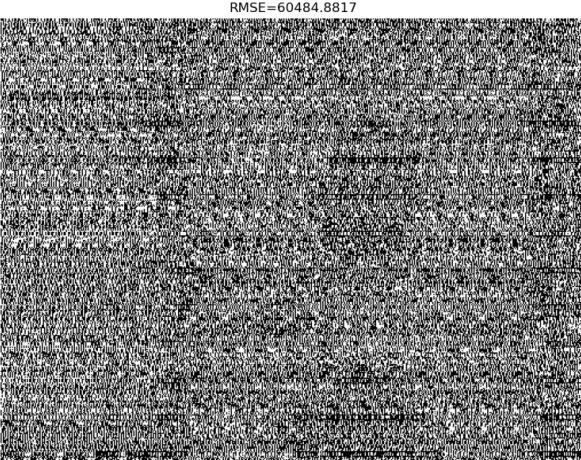
y\_model Predicted Image : PSNR=5.4038  
SSIM=0.0002  
RMSE=46104.0151

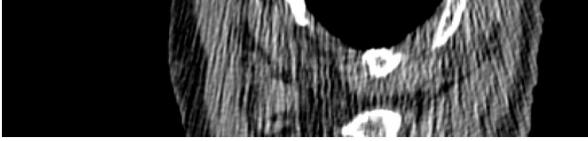


y\_model Predicted Image : PSNR=4.5031  
SSIM=0.0005  
RMSE=56729.6448

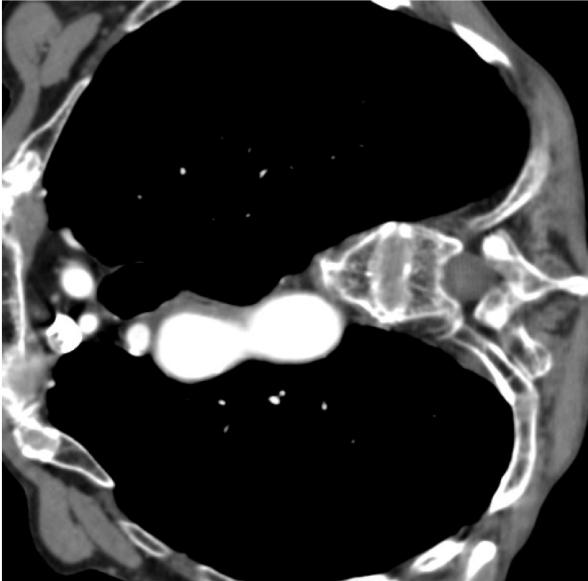


y\_model Predicted Image : PSNR=4.2247  
SSIM=-0.0003  
RMSE=60484.8817





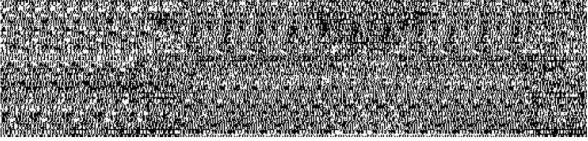
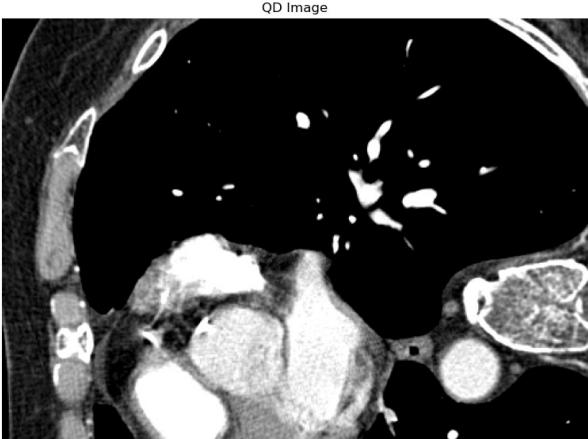
QD Image



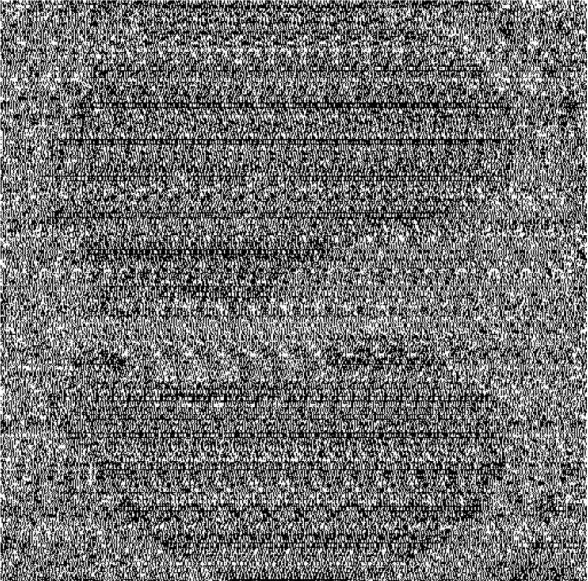
QD Image



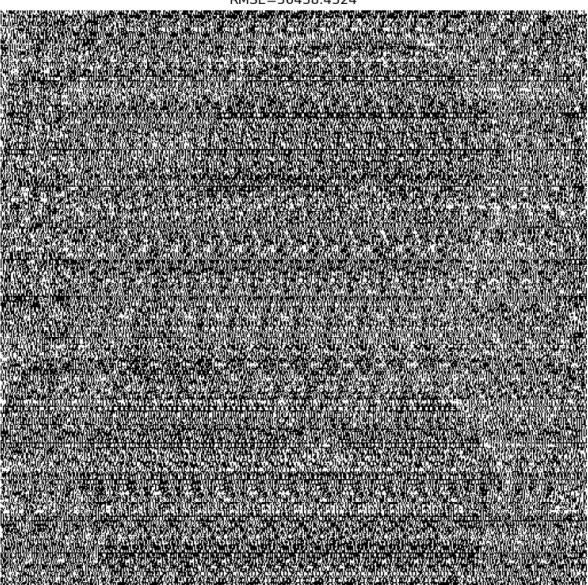
QD Image



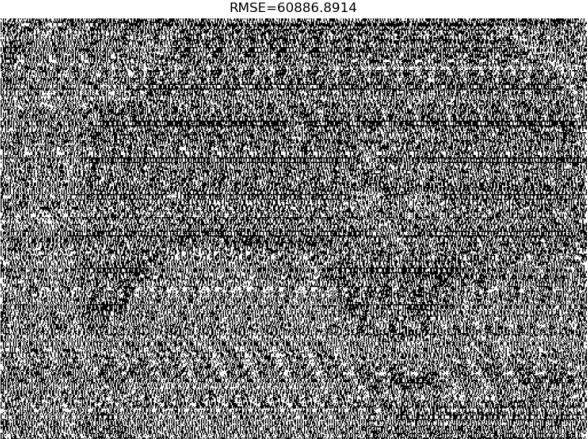
y\_model Predicted Image : PSNR=4.3008  
SSIM=-0.0002  
RMSE=59434.827

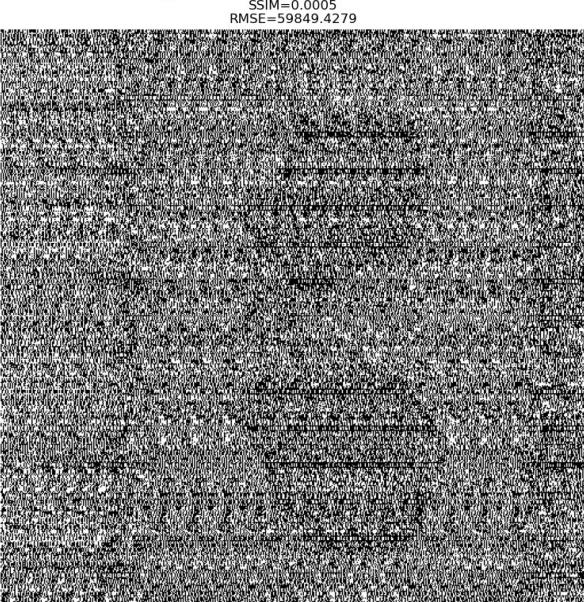
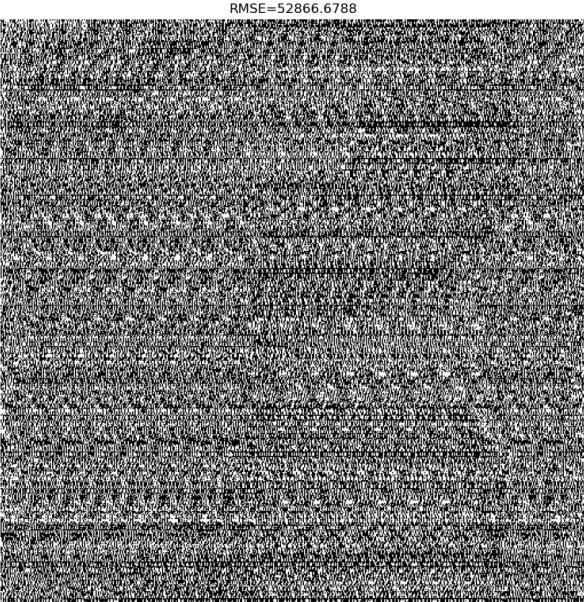
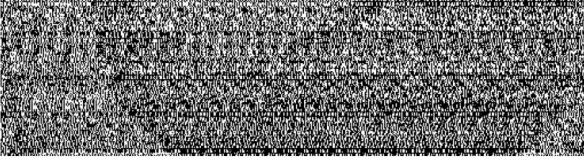
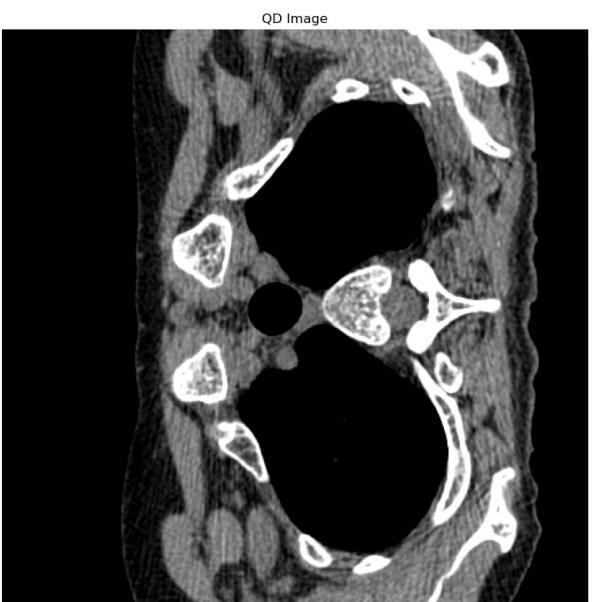
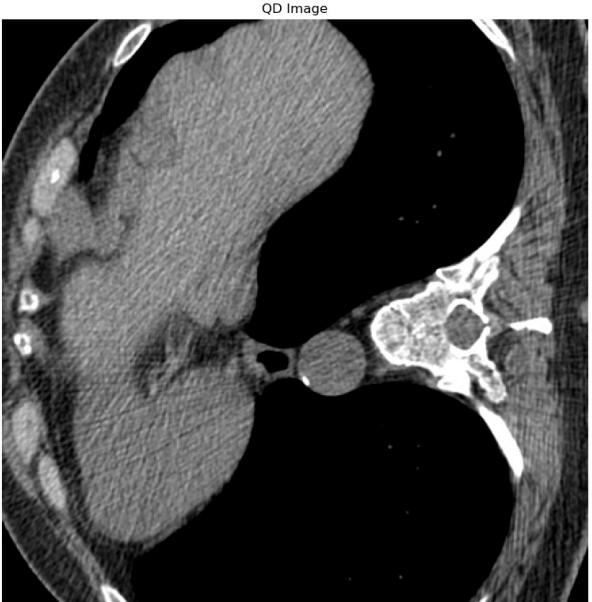


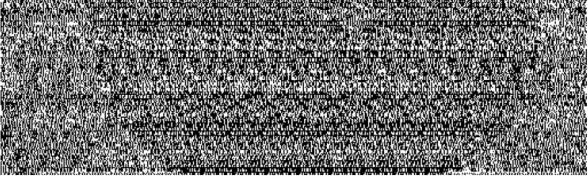
y\_model Predicted Image : PSNR=4.5255  
SSIM=0.0  
RMSE=56438.4324



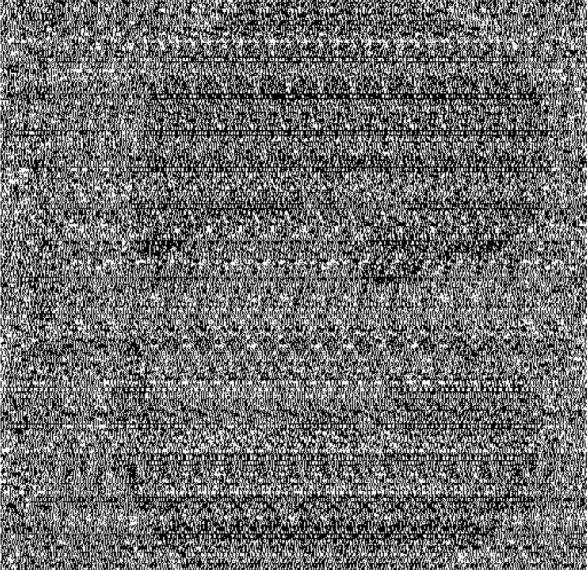
y\_model Predicted Image : PSNR=4.196  
SSIM=0.0003  
RMSE=60886.8914



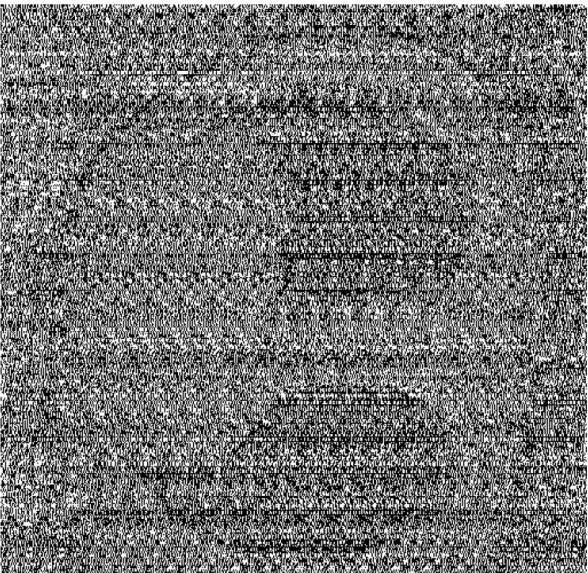




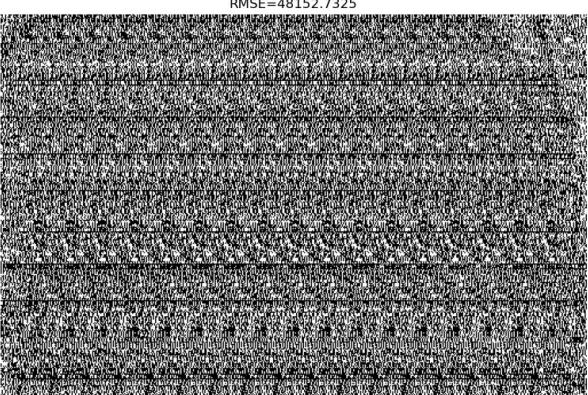
y\_model Predicted Image : PSNR=4.4191  
SSIM=0.0002  
RMSE=57837.8401

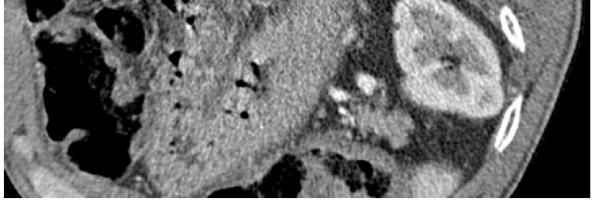


y\_model Predicted Image : PSNR=4.3045  
SSIM=0.0006  
RMSE=59383.5666

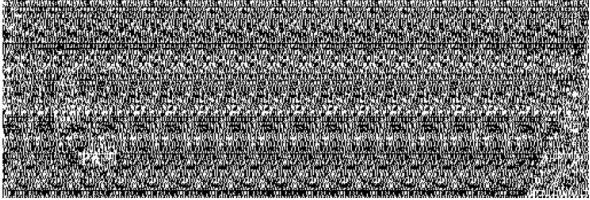


y\_model Predicted Image : PSNR=5.215  
SSIM=0.0006  
RMSE=48152.7325

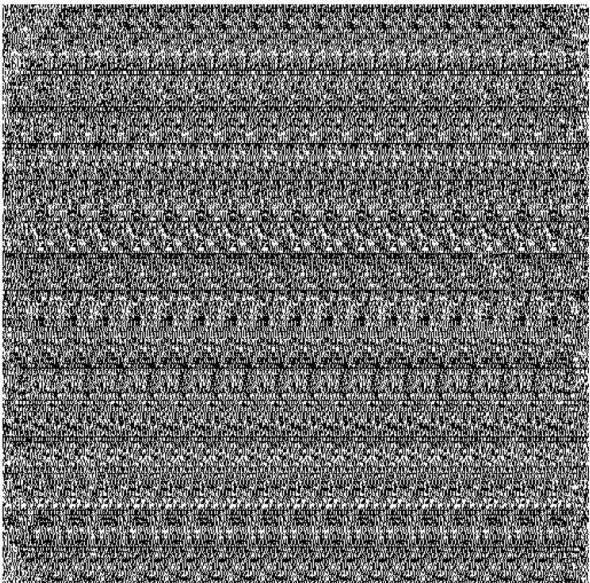




QD Image



y\_model Predicted Image : PSNR=5.4297  
SSIM=0.0006  
RMSE=45829.8564























## ▼ Model 3 : SA Model

```
sys.path.append('../denoising-models/hformer_self_attention')
from torchinfo import summary

from hformer_sa_model import HformerSAModel

sa_model = HformerSAModel(num_channels=64).cuda()
sa_model.load_state_dict(torch.load('../denoising-models/hformer_self_attention/weights/model_75.pth'))
sa_model.eval()
print('model summary\n', summary(sa_model, input_size=(64, 64, 64, 1)))

model summary
=====
Layer (type:depth-idx)           Output Shape          Param #

```

=====			
HformerSAModel	[64, 64, 64, 1]	--	--
└─InputProjectionLayer: 1-1	[64, 64, 64, 64]	--	--
└─Conv2d: 2-1	[64, 64, 64, 64]	640	
└─ConvBlock: 1-2	[64, 64, 64, 64]	--	--
└─ZeroPad2d: 2-2	[64, 64, 18, 18]	--	--
└─DepthWiseSeparableConv2d: 2-3	[64, 64, 18, 18]	--	--
└─Conv2d: 3-1	[64, 64, 18, 18]	3,136	
└─Conv2d: 3-2	[64, 64, 18, 18]	4,096	
└─Conv2d: 2-4	[64, 256, 12, 12]	803,072	
└─GELU: 2-5	[64, 256, 12, 12]	--	--
└─Conv2d: 2-6	[64, 64, 12, 12]	16,448	
└─ZeroPad2d: 2-7	[64, 64, 18, 18]	--	--
└─DepthWiseSeparableConv2d: 2-8	[64, 64, 18, 18]	--	--
└─Conv2d: 3-3	[64, 64, 18, 18]	3,136	
└─Conv2d: 3-4	[64, 64, 18, 18]	4,096	
└─Conv2d: 2-9	[64, 256, 12, 12]	803,072	
└─GELU: 2-10	[64, 256, 12, 12]	--	--
└─Conv2d: 2-11	[64, 64, 12, 12]	16,448	
└─DWTInverse: 2-12	[64, 64, 64, 64]	--	--
└─Conv2d: 1-3	[64, 128, 32, 32]	32,896	
└─HformerBlock: 1-4	[64, 128, 32, 32]	16,512	
└─DepthWiseSeparableConv2d: 2-13	[64, 128, 32, 32]	--	--
└─Conv2d: 3-5	[64, 128, 32, 32]	6,272	
└─Conv2d: 3-6	[64, 128, 32, 32]	16,384	
└─MaxPool2d: 2-14	[64, 128, 16, 16]	--	--
└─Linear: 2-15	[64, 256, 128]	16,384	
└─Linear: 2-16	[64, 256, 128]	16,384	
└─Linear: 2-17	[64, 1024, 128]	16,384	
└─LayerNorm: 2-18	[64, 128, 32, 32]	--	--
└─Conv2d: 2-19	[64, 512, 32, 32]	66,048	
└─GELU: 2-20	[64, 512, 32, 32]	--	--
└─Conv2d: 2-21	[64, 128, 32, 32]	65,664	
└─Conv2d: 1-5	[64, 256, 16, 16]	131,328	
└─HformerBlock: 1-6	[64, 256, 16, 16]	65,792	
└─DepthWiseSeparableConv2d: 2-22	[64, 256, 16, 16]	--	--
└─Conv2d: 3-7	[64, 256, 16, 16]	12,544	
└─Conv2d: 3-8	[64, 256, 16, 16]	65,536	
└─MaxPool2d: 2-23	[64, 256, 8, 8]	--	--
└─Linear: 2-24	[64, 64, 256]	65,536	
└─Linear: 2-25	[64, 64, 256]	65,536	
└─Linear: 2-26	[64, 256, 256]	65,536	
└─LayerNorm: 2-27	[64, 256, 16, 16]	--	--
└─Conv2d: 2-28	[64, 1024, 16, 16]	263,168	
└─GELU: 2-29	[64, 1024, 16, 16]	--	--
└─Conv2d: 2-30	[64, 256, 16, 16]	262,400	
└─ConvTranspose2d: 1-7	[64, 128, 32, 32]	131,200	
└─HformerBlock: 1-8	[64, 128, 32, 32]	16,512	
└─DepthWiseSeparableConv2d: 2-31	[64, 128, 32, 32]	--	--
└─Conv2d: 3-9	[64, 128, 32, 32]	6,272	
└─Conv2d: 3-10	[64, 128, 32, 32]	16,384	
└─MaxPool2d: 2-32	[64, 128, 16, 16]	--	--
└─Linear: 2-33	[64, 256, 128]	16,384	
└─Linear: 2-34	[64, 256, 128]	16,384	
└─Linear: 2-35	[64, 1024, 128]	16,384	

```
sa_prediction_patches = []

with torch.no_grad():
    for i, data in enumerate(noisy_image_patches_array):
        noisy = data

        predictions = sa_model(torch.unsqueeze(torch.from_numpy(noisy.numpy()), dim=0).to('cuda')).cpu()

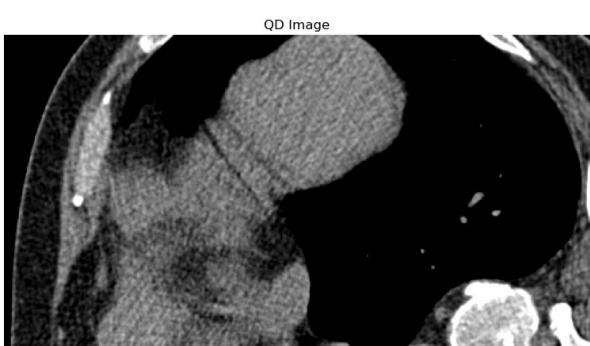
        sa_prediction_patches.append(predictions.detach().cpu())

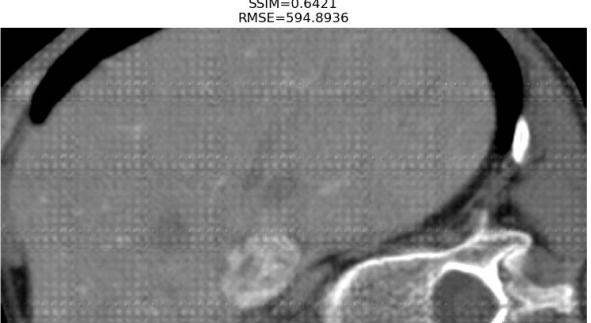
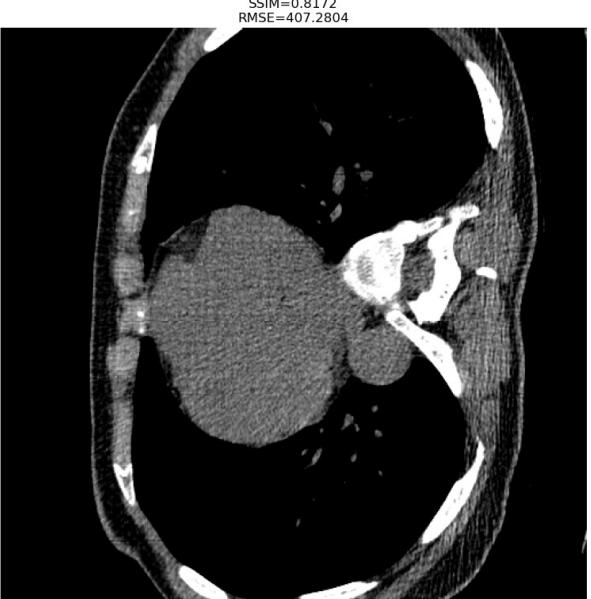
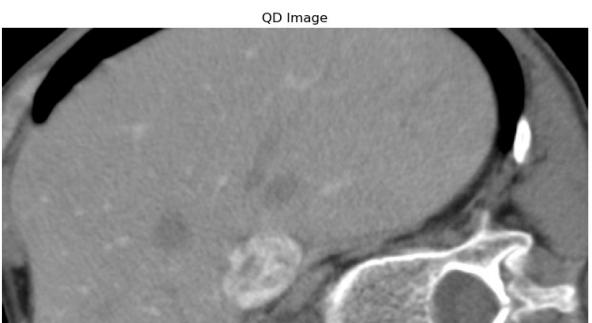
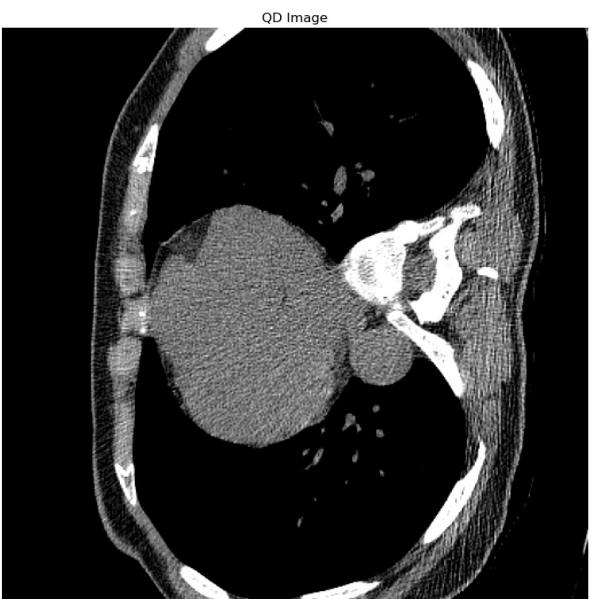
sa_prediction_patches = np.concatenate(sa_prediction_patches, axis=0)
sa_predictions = np.expand_dims(reconstruct_image_from_patches(sa_prediction_patches[0:64], 8), axis=0)

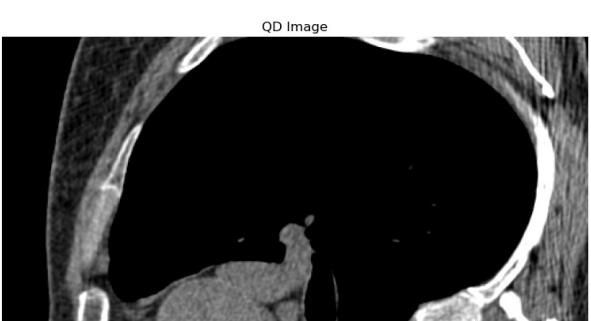
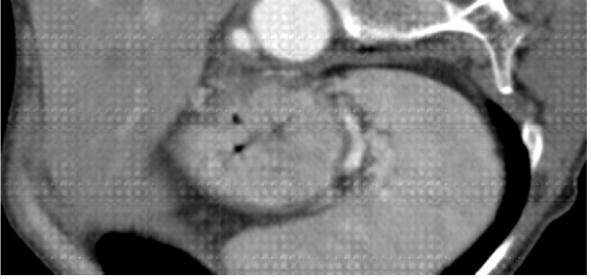
for i in range(1, int(sa_prediction_patches.shape[0] / 64)):
    reconstructed_image = reconstruct_image_from_patches(sa_prediction_patches[i * 64 : i * 64 + 64], 8)
    reconstructed_image = np.expand_dims(reconstructed_image, axis=0)

    sa_predictions= np.append(sa_predictions, reconstructed_image, axis=0)

visualize_predictions(sa_model, noisy_array, len(noisy_array), sa_predictions, "sa model")
```

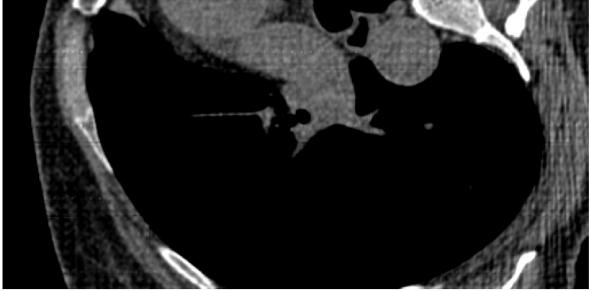








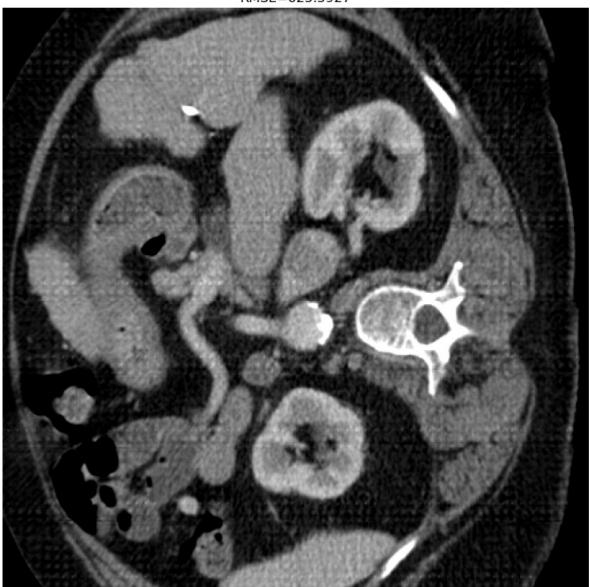
QD Image



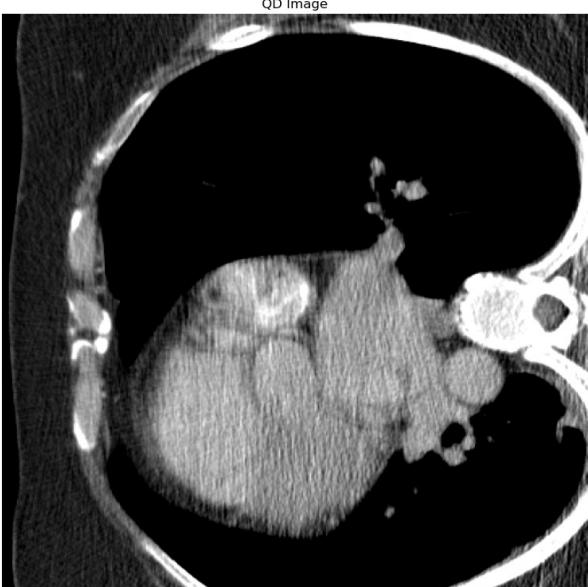
sa model Predicted Image : PSNR=24.0783  
SSIM=0.7006  
RMSE=625.5927



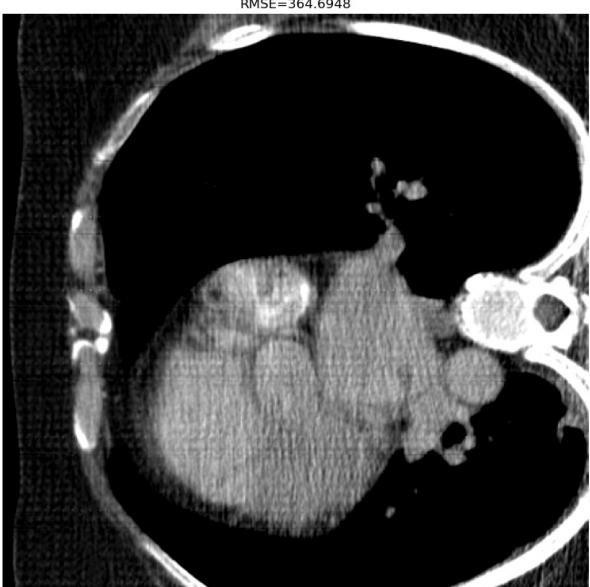
QD Image



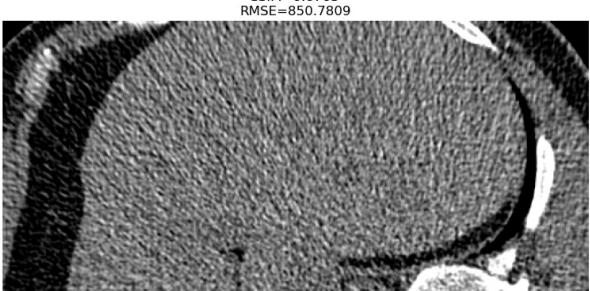
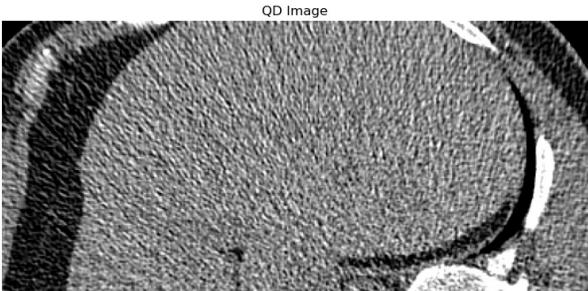
sa model Predicted Image : PSNR=26.4219  
SSIM=0.8758  
RMSE=364.6948



QD Image



sa model Predicted Image : PSNR=22.743  
SSIM=0.6763  
RMSE=850.7809





QD Image



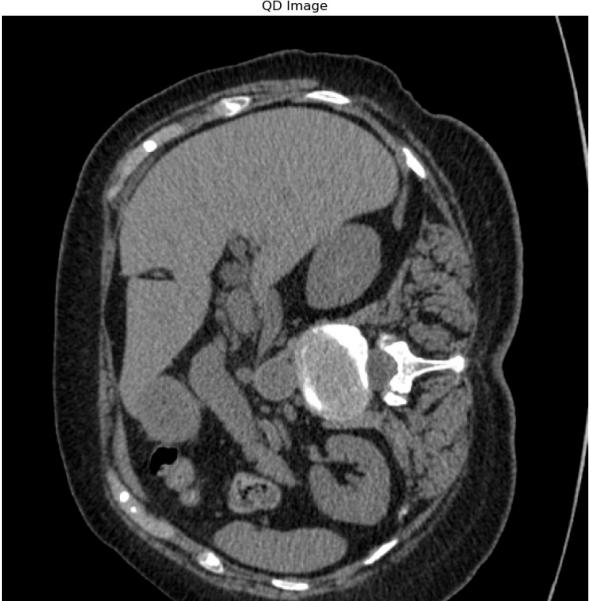
sa model Predicted Image : PSNR=25.9929  
SSIM=0.7982  
RMSE=402.5555



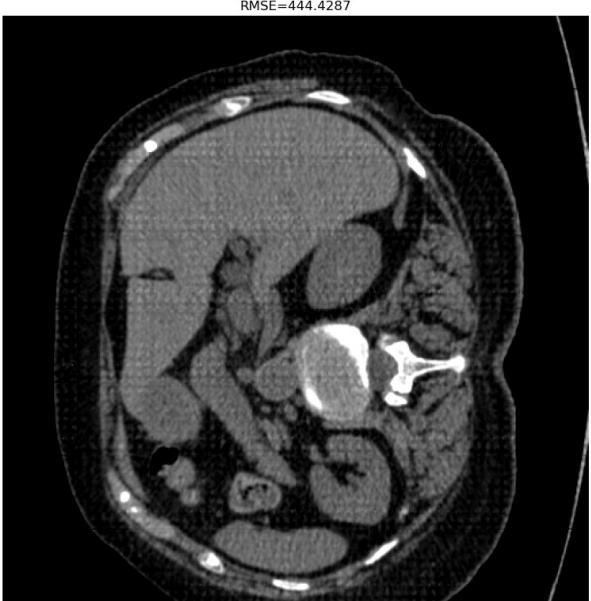
QD Image



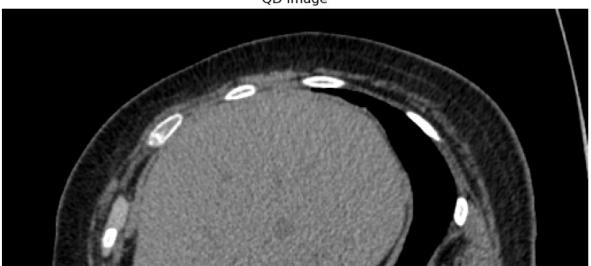
sa model Predicted Image : PSNR=25.5632  
SSIM=0.7296  
RMSE=444.4287

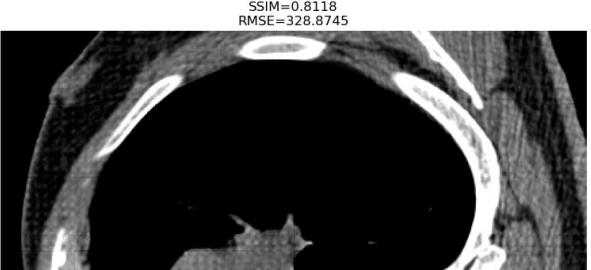
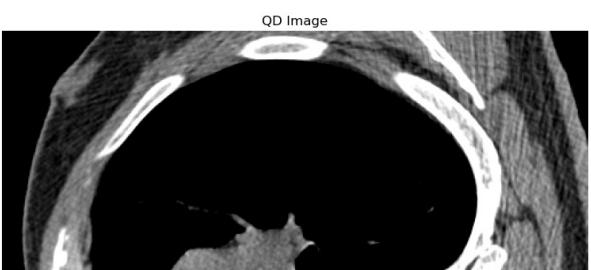
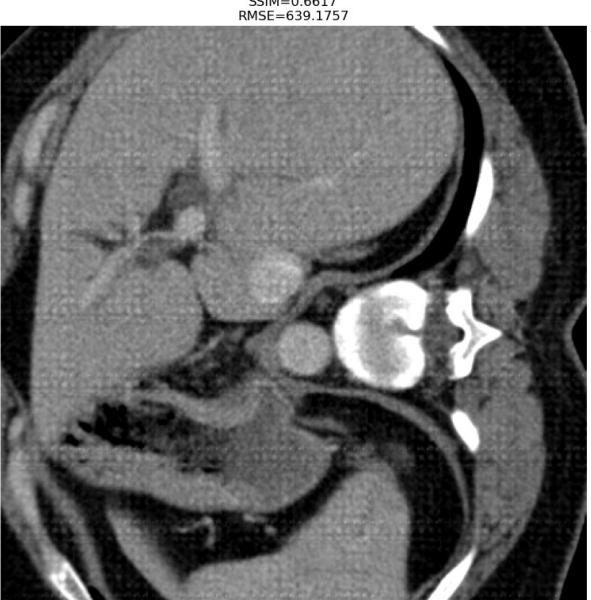
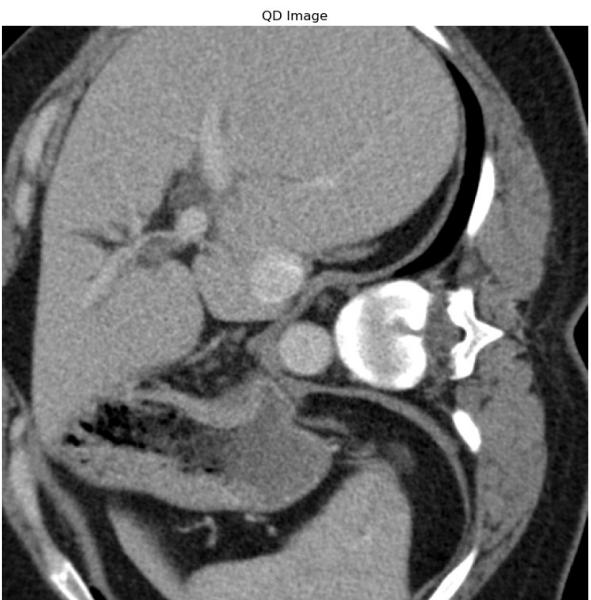
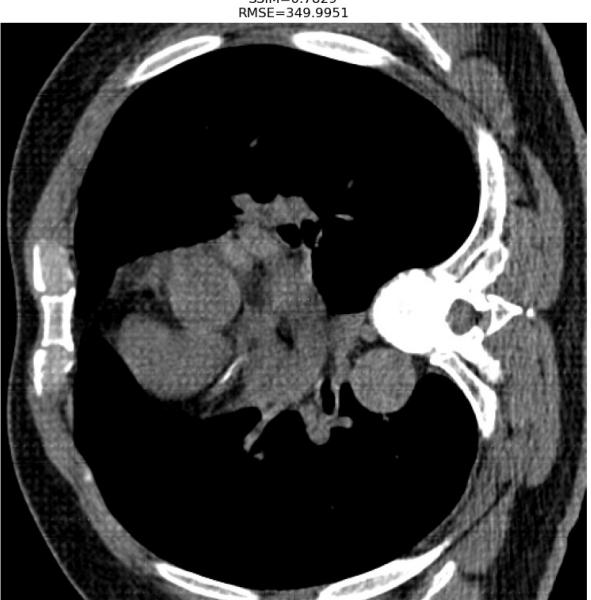
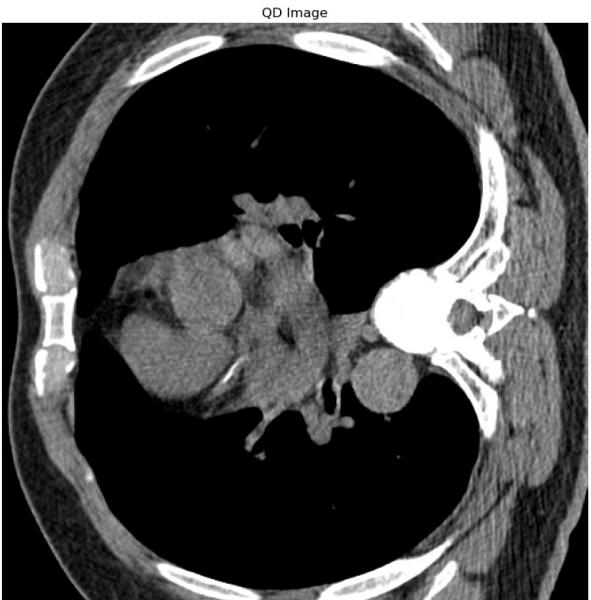
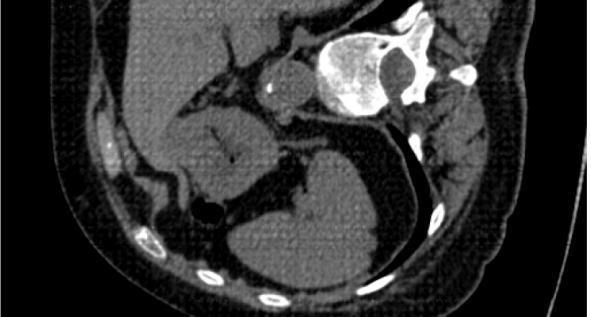
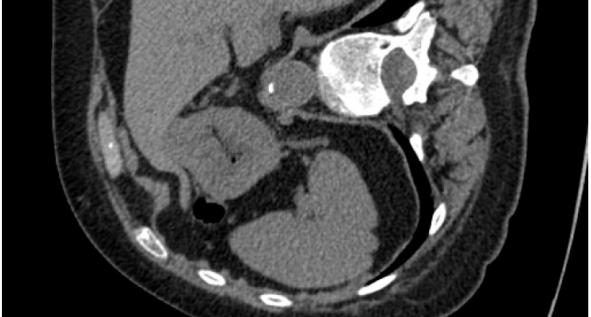


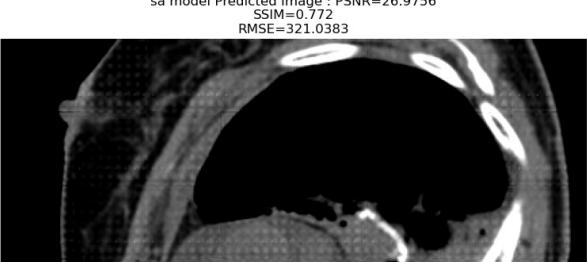
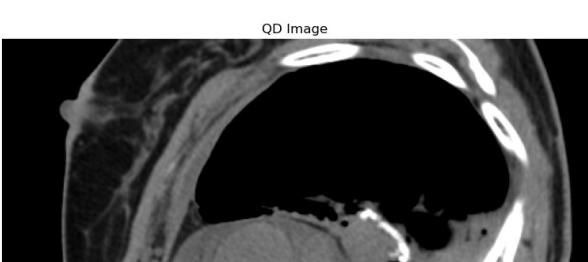
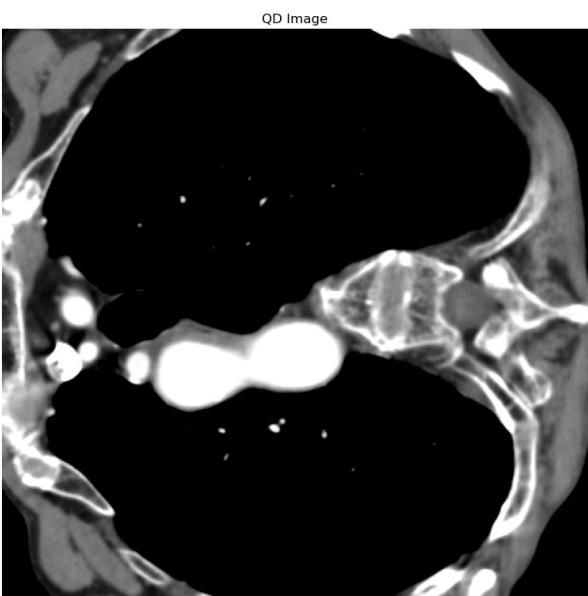
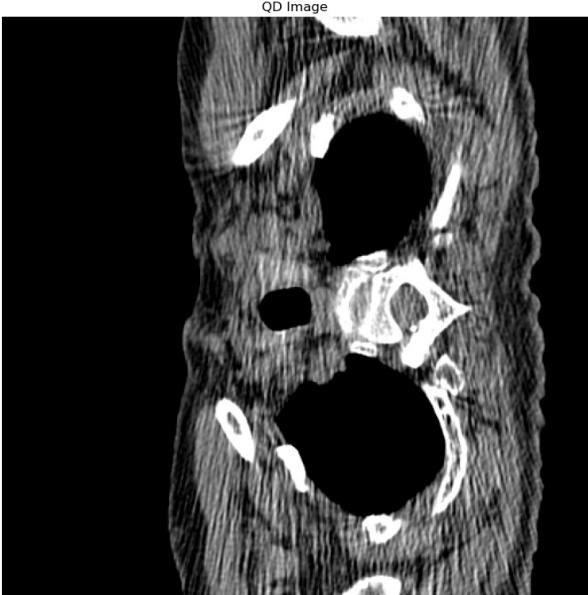
QD Image

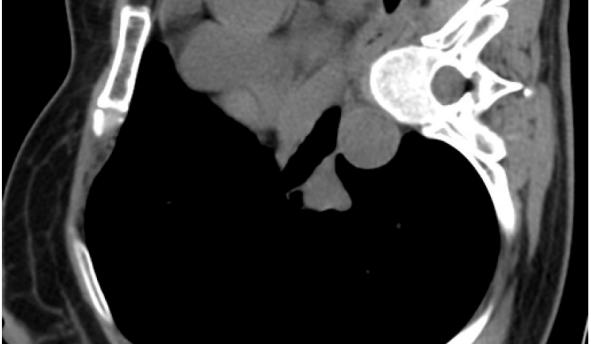


sa model Predicted Image : PSNR=25.5903  
SSIM=0.737  
RMSE=441.658

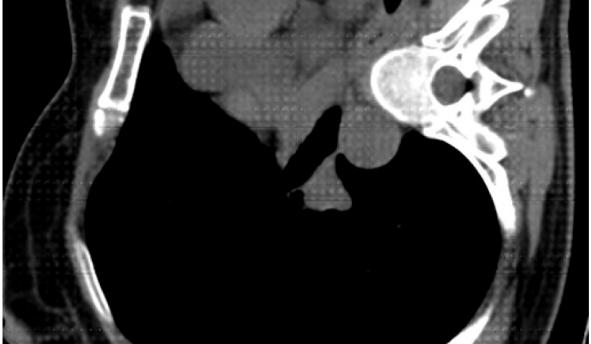




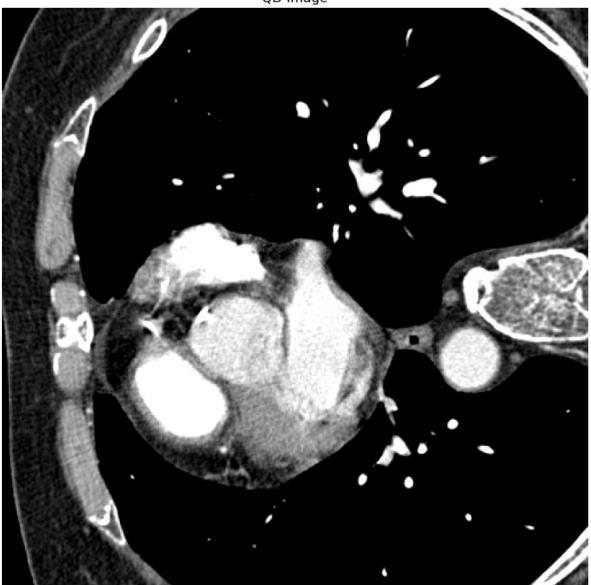




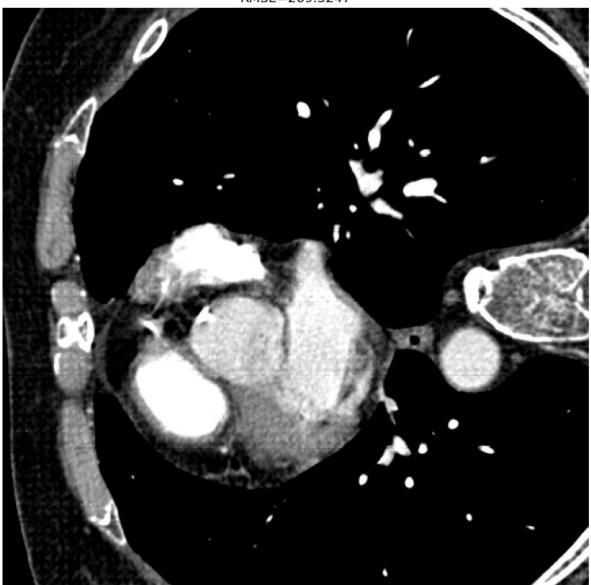
QD Image



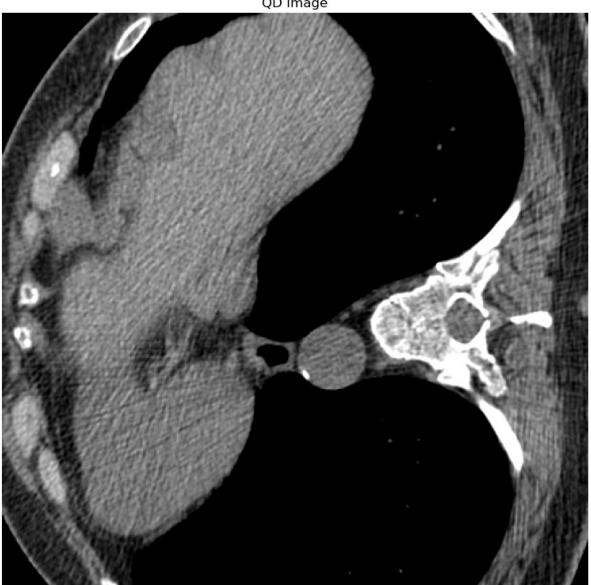
sa model Predicted Image : PSNR=27.7384  
SSIM=0.9143  
RMSE=269.3247



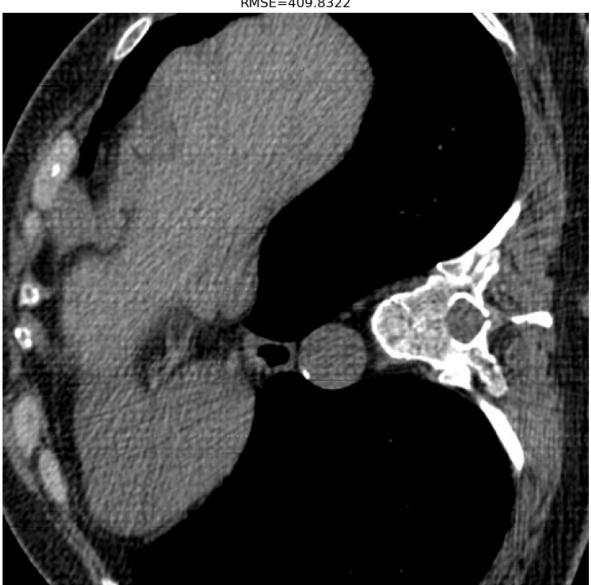
QD Image



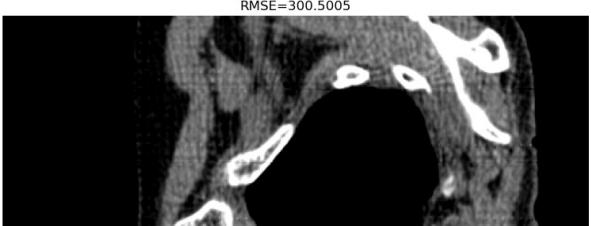
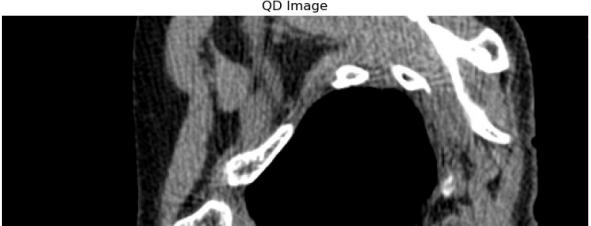
sa model Predicted Image : PSNR=25.9151  
SSIM=0.7503  
RMSE=409.8322

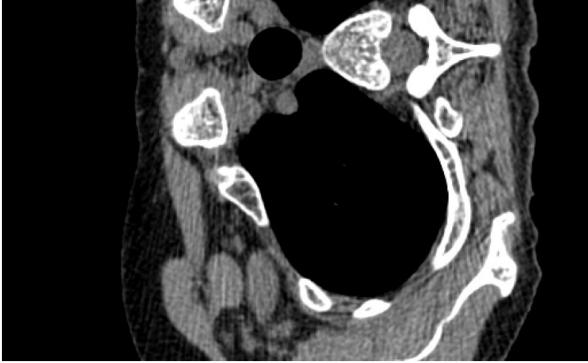


QD Image

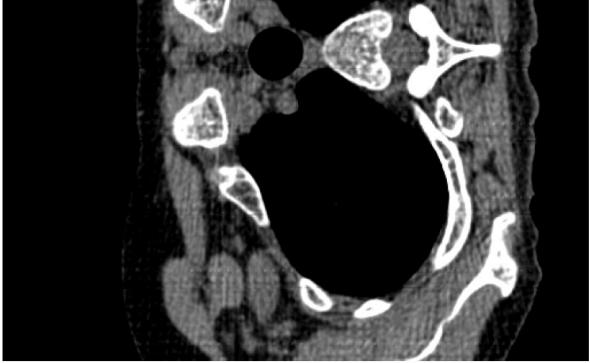


sa model Predicted Image : PSNR=27.2627  
SSIM=0.8112  
RMSE=300.5005





QD Image



sa model Predicted Image : PSNR=28.0553  
SSIM=0.8414  
RMSE=250.3721



QD Image



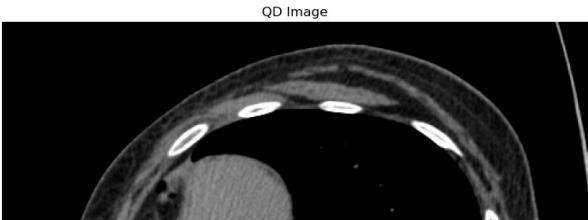
sa model Predicted Image : PSNR=27.2945  
SSIM=0.8404  
RMSE=298.3129



QD Image

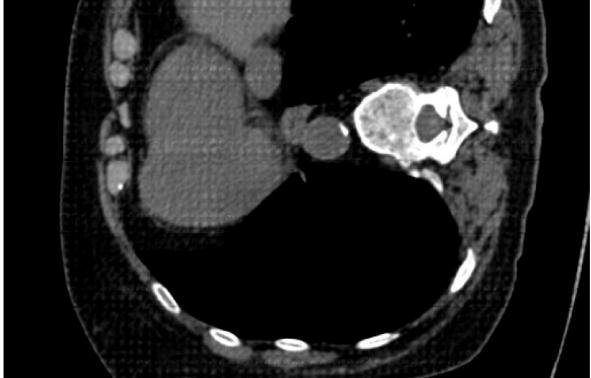


sa model Predicted Image : PSNR=27.0916  
SSIM=0.8126  
RMSE=312.5817





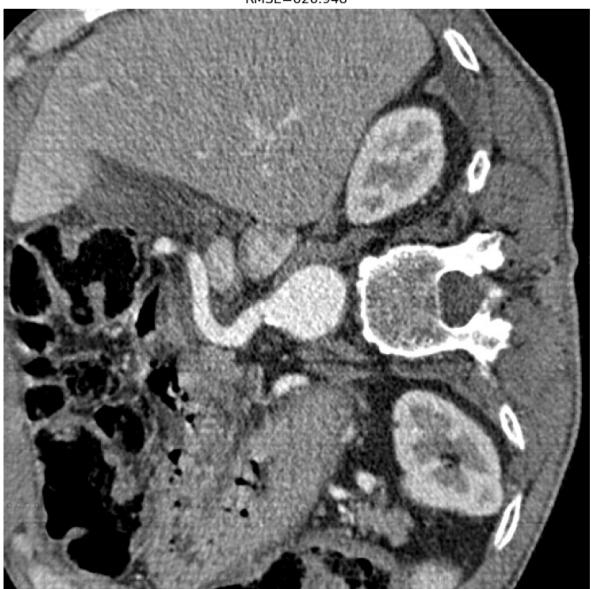
QD Image



sa model Predicted Image : PSNR=24.0689  
SSIM=0.6803  
RMSE=626.948



QD Image



sa model Predicted Image : PSNR=24.0001  
SSIM=0.6469  
RMSE=636.9604



QD Image

























## ✓ Model 4 : Z Model

```

sys.path.append('../denoising-models/hformer_pytorch')
from torchinfo import summary

from z_model import ZModel

z_model = ZModel(num_channels=16).cuda()
z_model.load_state_dict(torch.load('../denoising-models/hformer_pytorch/weights/z_model_159.pth'))
z_model.eval()
print('model summary\n', summary(z_model, input_size=(64, 64, 64, 1)))

Using cache found in C:\Users\Tarun/.cache\torch\hub\mateusbuda_brain-segmentation-p
Using cache found in C:\Users\Tarun/.cache\torch\hub\mateusbuda_brain-segmentation-p
model summary
=====
Layer (type:depth-idx)           Output Shape      Param #
=====
ZModel
├─InputProjectionLayer: 1-1     [64, 64, 64, 1]    --
  └─Conv2d: 2-1                 [64, 16, 64, 64]   --
  └─UNet: 1-2                  [64, 16, 64, 64]   --
    └─Sequential: 2-2          [64, 16, 64, 64]   --
      └─Conv2d: 3-1            [64, 16, 64, 64]   2,304
      └─BatchNorm2d: 3-2       [64, 16, 64, 64]   32
      └─ReLU: 3-3              [64, 16, 64, 64]   --
      └─Conv2d: 3-4            [64, 16, 64, 64]   2,304
      └─BatchNorm2d: 3-5       [64, 16, 64, 64]   32
      └─ReLU: 3-6              [64, 16, 64, 64]   --
      └─MaxPool2d: 2-3         [64, 16, 32, 32]   --
      └─Sequential: 2-4         [64, 32, 32, 32]   --
        └─Conv2d: 3-7          [64, 32, 32, 32]   4,608
        └─BatchNorm2d: 3-8       [64, 32, 32, 32]   64
        └─ReLU: 3-9              [64, 32, 32, 32]   --
        └─Conv2d: 3-10           [64, 32, 32, 32]   9,216
        └─BatchNorm2d: 3-11       [64, 32, 32, 32]   64
        └─ReLU: 3-12              [64, 32, 32, 32]   --
        └─MaxPool2d: 2-5         [64, 32, 16, 16]   --
        └─Sequential: 2-6         [64, 64, 16, 16]   --
          └─Conv2d: 3-13          [64, 64, 16, 16]   18,432
          └─BatchNorm2d: 3-14       [64, 64, 16, 16]   128
          └─ReLU: 3-15              [64, 64, 16, 16]   --
          └─Conv2d: 3-16           [64, 64, 16, 16]   36,864
          └─BatchNorm2d: 3-17       [64, 64, 16, 16]   128
          └─ReLU: 3-18              [64, 64, 16, 16]   --
          └─MaxPool2d: 2-7         [64, 64, 8, 8]   --
          └─Sequential: 2-8         [64, 128, 8, 8]   --
            └─Conv2d: 3-19          [64, 128, 8, 8]   73,728
            └─BatchNorm2d: 3-20       [64, 128, 8, 8]   256
            └─ReLU: 3-21              [64, 128, 8, 8]   --

```

└Conv2d: 3-22	[64, 128, 8, 8]	147,456
└BatchNorm2d: 3-23	[64, 128, 8, 8]	256
└ReLU: 3-24	[64, 128, 8, 8]	--
└MaxPool2d: 2-9	[64, 128, 4, 4]	--
└Sequential: 2-10	[64, 256, 4, 4]	--
└Conv2d: 3-25	[64, 256, 4, 4]	294,912
└BatchNorm2d: 3-26	[64, 256, 4, 4]	512
└ReLU: 3-27	[64, 256, 4, 4]	--
└Conv2d: 3-28	[64, 256, 4, 4]	589,824
└BatchNorm2d: 3-29	[64, 256, 4, 4]	512
└ReLU: 3-30	[64, 256, 4, 4]	--
└ConvTranspose2d: 2-11	[64, 128, 8, 8]	131,200
└Sequential: 2-12	[64, 128, 8, 8]	--
└Conv2d: 3-31	[64, 128, 8, 8]	294,912
└BatchNorm2d: 3-32	[64, 128, 8, 8]	256
└ReLU: 3-33	[64, 128, 8, 8]	--
└Conv2d: 3-34	[64, 128, 8, 8]	147,456
└BatchNorm2d: 3-35	[64, 128, 8, 8]	256
└ReLU: 3-36	[64, 128, 8, 8]	--
└ConvTranspose2d: 2-13	[64, 64, 16, 16]	32,832

```

z_prediction_patches = []

with torch.no_grad():
    for i, data in enumerate(noisy_image_patches_array):
        noisy = data

        predictions = z_model(torch.unsqueeze(torch.from_numpy(noisy.numpy()), dim=0).to('cuda')).cpu()

        z_prediction_patches.append(predictions.detach().cpu())

z_prediction_patches = np.concatenate(z_prediction_patches, axis=0)
z_predictions = np.expand_dims(reconstruct_image_from_patches(z_prediction_patches[0:64], 8), axis=0)

for i in range(1, int(z_prediction_patches.shape[0] / 64)):
    reconstructed_image = reconstruct_image_from_patches(z_prediction_patches[i * 64 : i * 64 + 64], 8)
    reconstructed_image = np.expand_dims(reconstructed_image, axis=0)

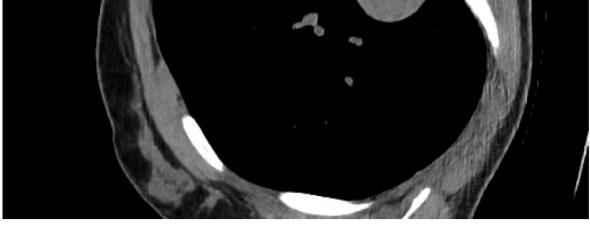
    z_predictions = np.append(z_predictions, reconstructed_image, axis=0)
visualize_predictions(z_model, noisy_array, len(noisy_array), z_predictions, "z model")

```

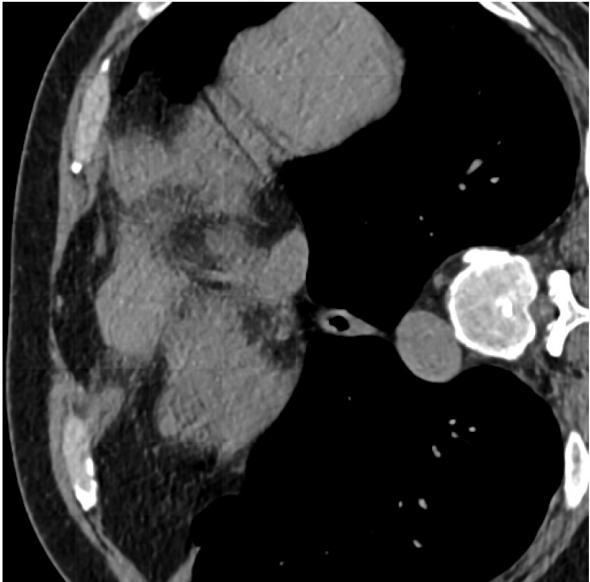




QD Image



z model Predicted Image : PSNR=32.8672  
SSIM=0.9278  
RMSE=82.6796

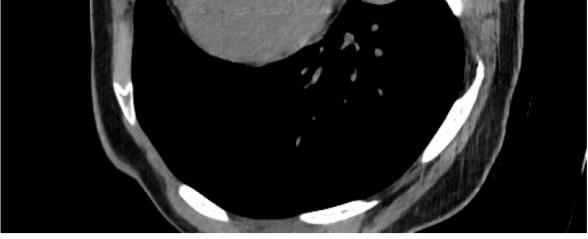


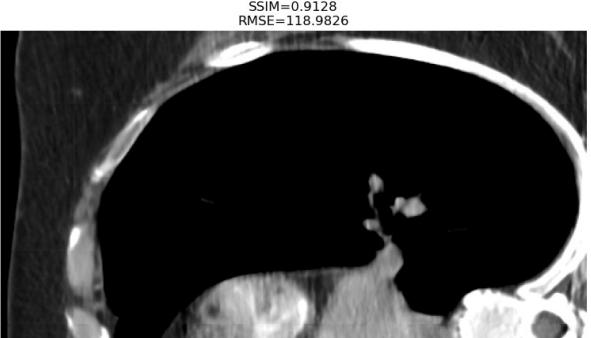
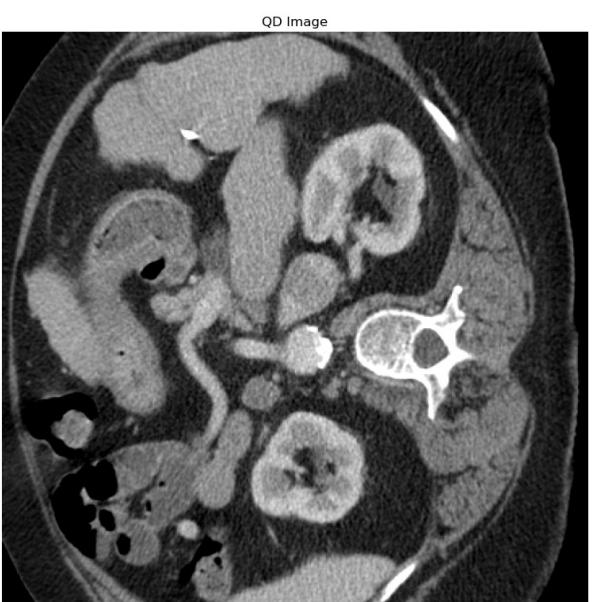
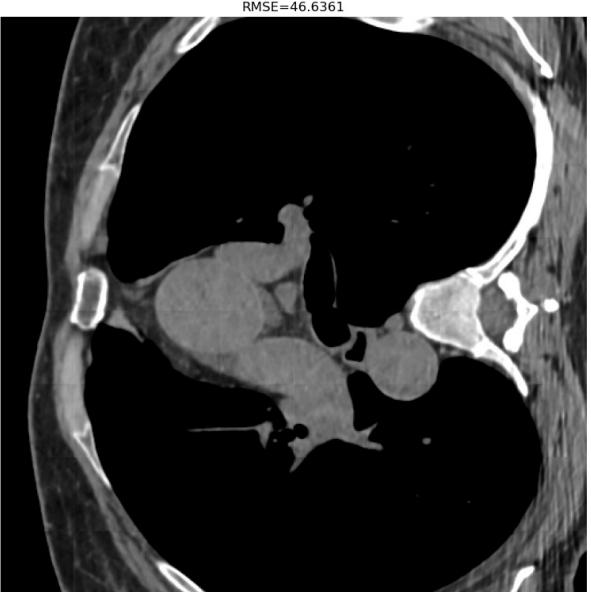
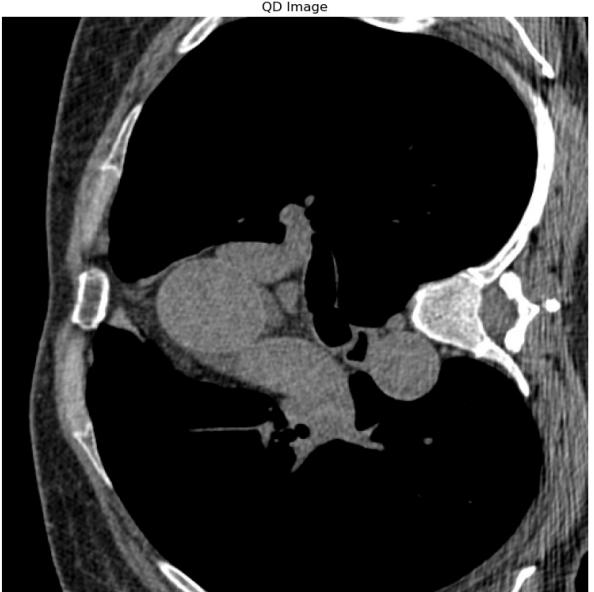
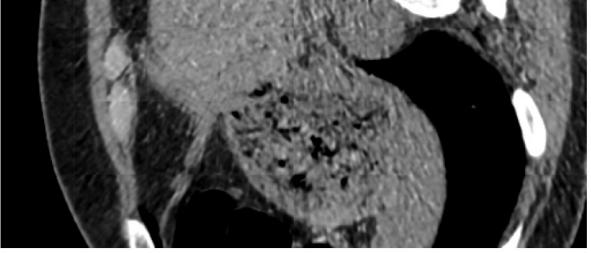
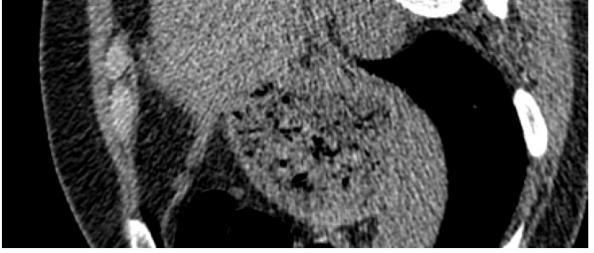
z model Predicted Image : PSNR=34.0757  
SSIM=0.9649  
RMSE=62.5972



z model Predicted Image : PSNR=25.6435  
SSIM=0.8803  
RMSE=436.2883

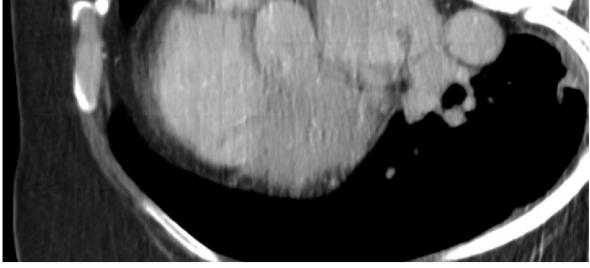








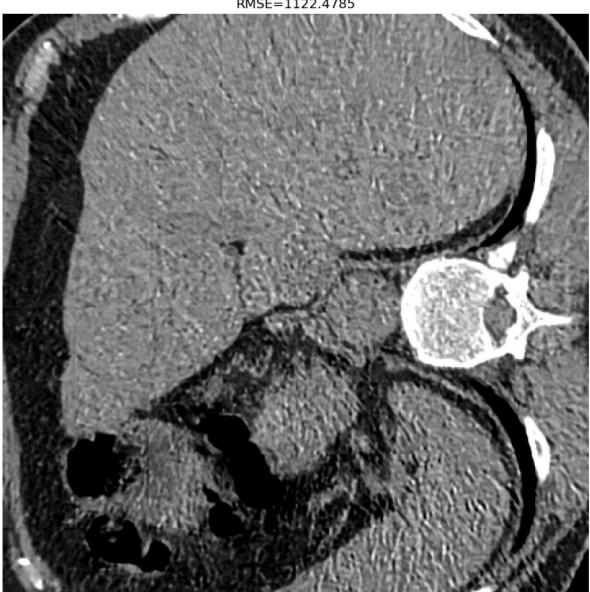
QD Image



z model Predicted Image : PSNR=21.5394  
SSIM=0.7742  
RMSE=1122.4785



QD Image



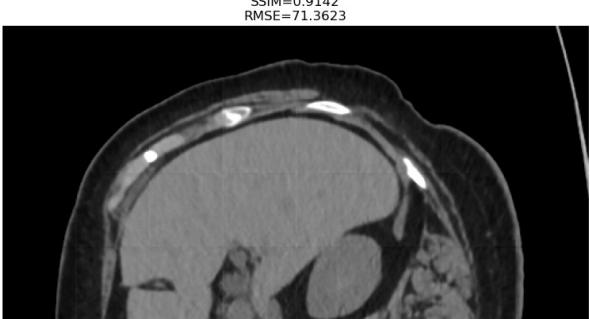
z model Predicted Image : PSNR=34.3931  
SSIM=0.9388  
RMSE=58.1853



QD Image



z model Predicted Image : PSNR=33.5065  
SSIM=0.9142  
RMSE=71.3623





QD Image



z model Predicted Image : PSNR=33.3628  
SSIM=0.9015  
RMSE=73.7633



QD Image



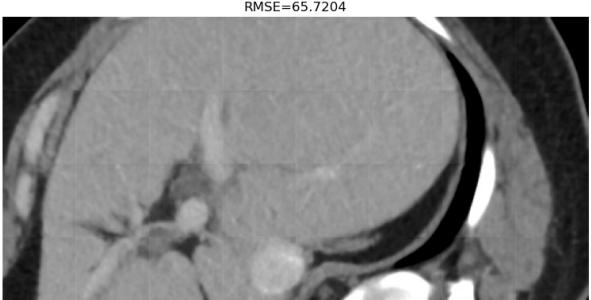
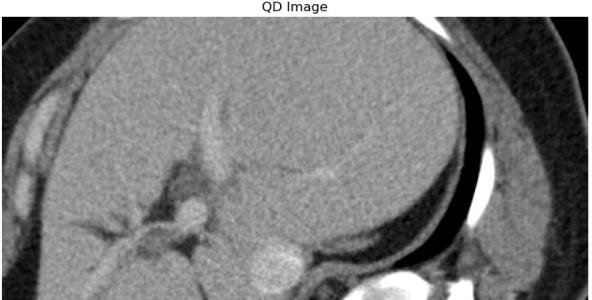
z model Predicted Image : PSNR=33.9679  
SSIM=0.9497  
RMSE=64.1691

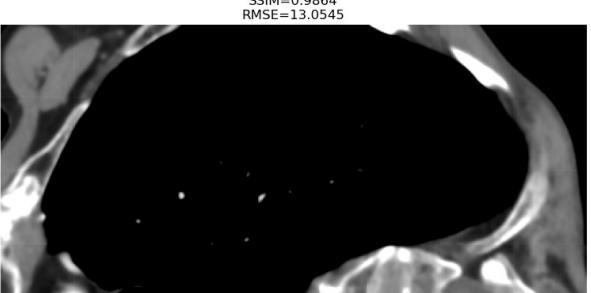
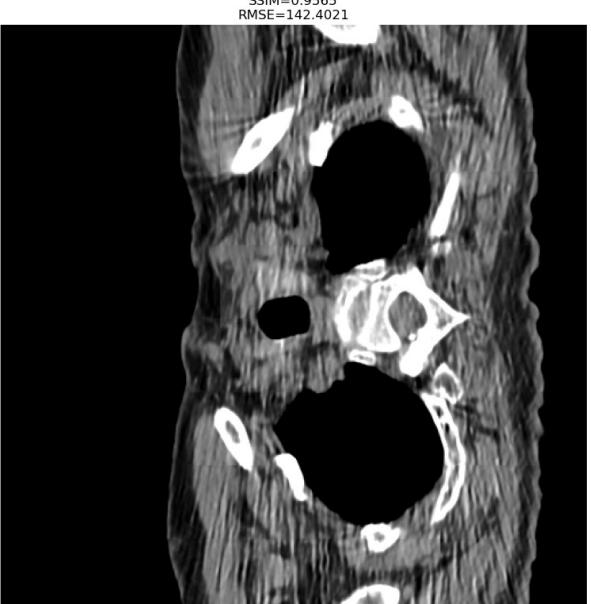
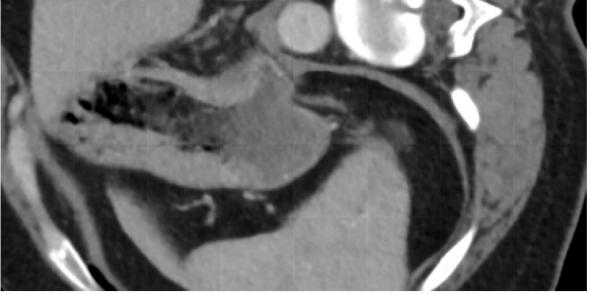
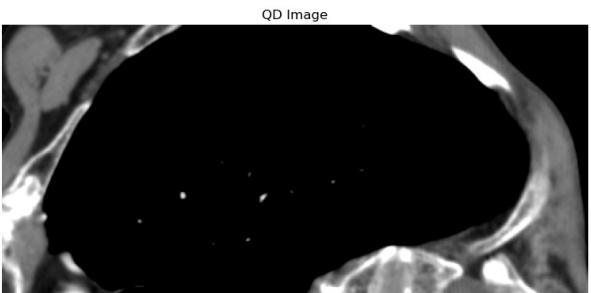
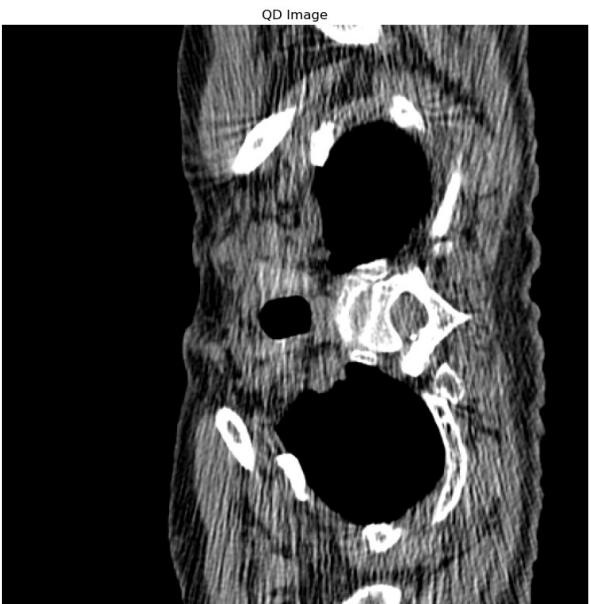
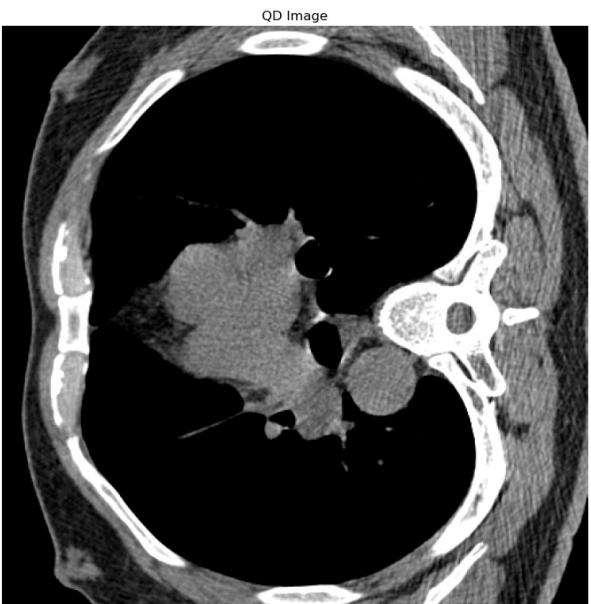
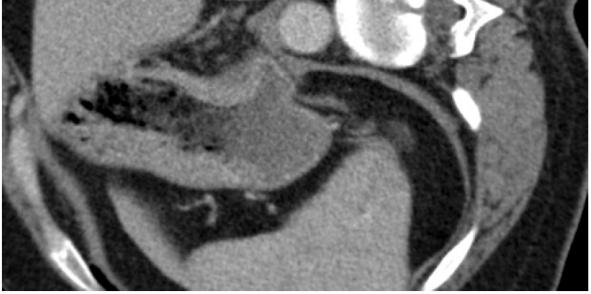


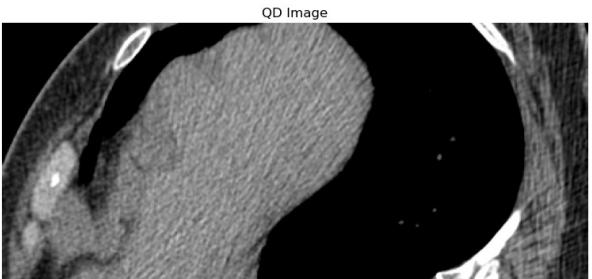
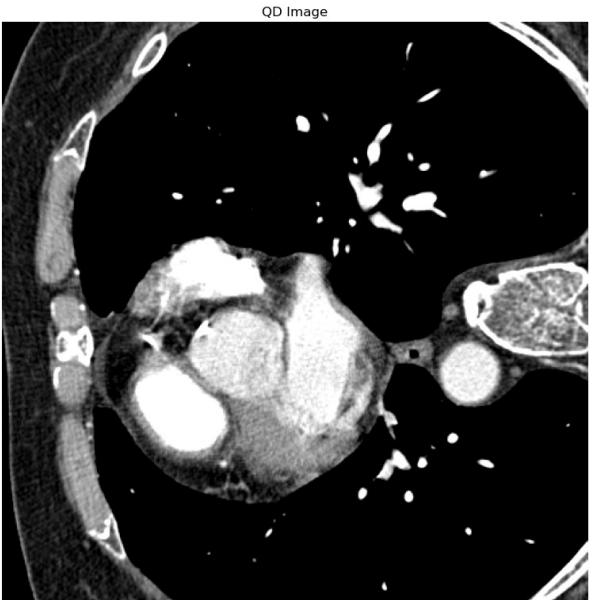
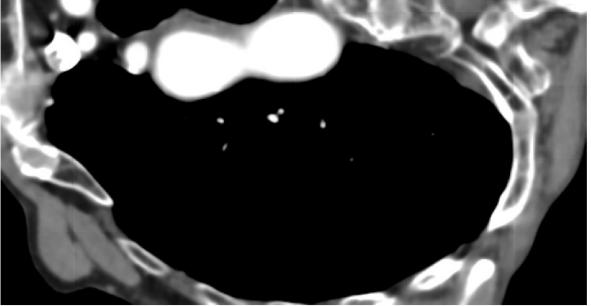
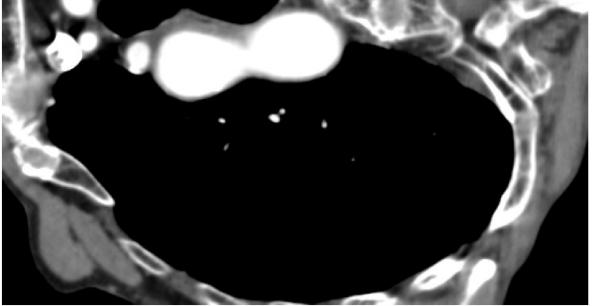
QD Image

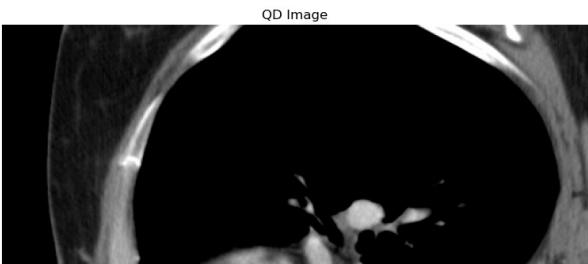
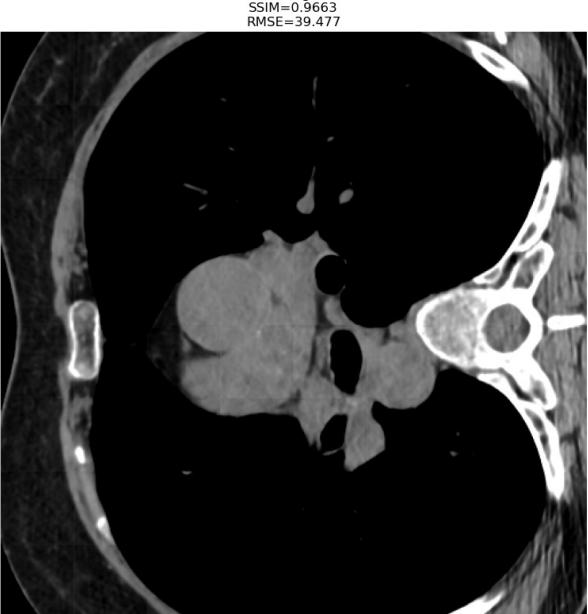
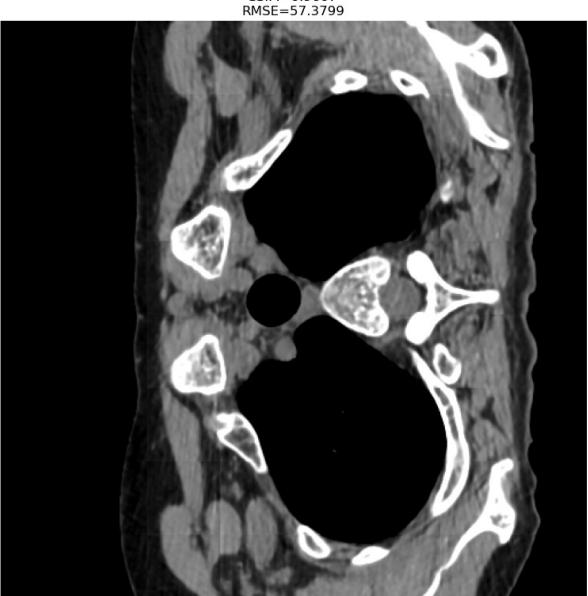
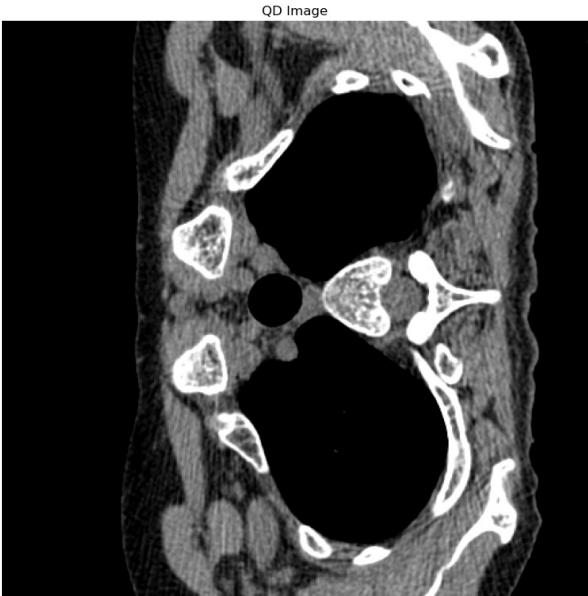
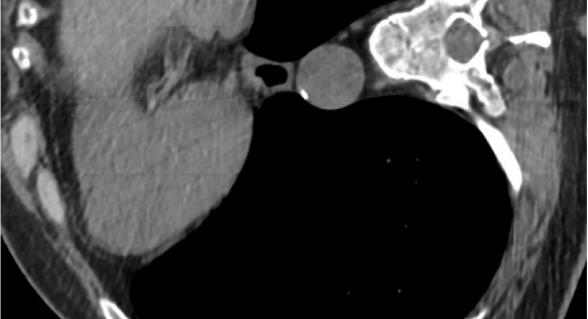
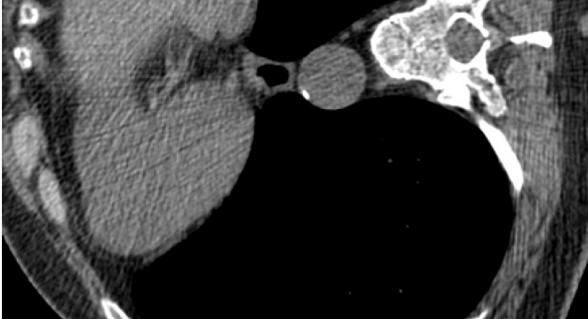


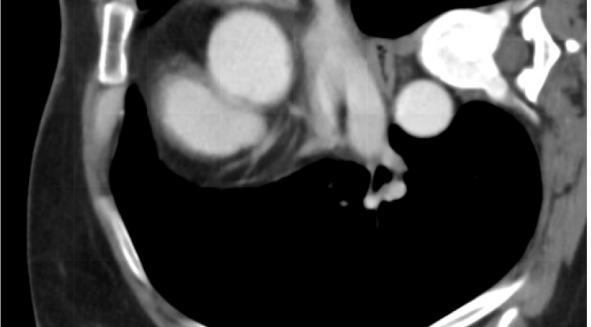
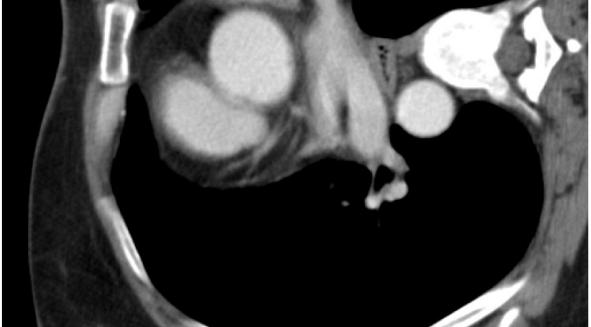
z model Predicted Image : PSNR=33.8642  
SSIM=0.8909  
RMSE=65.7204





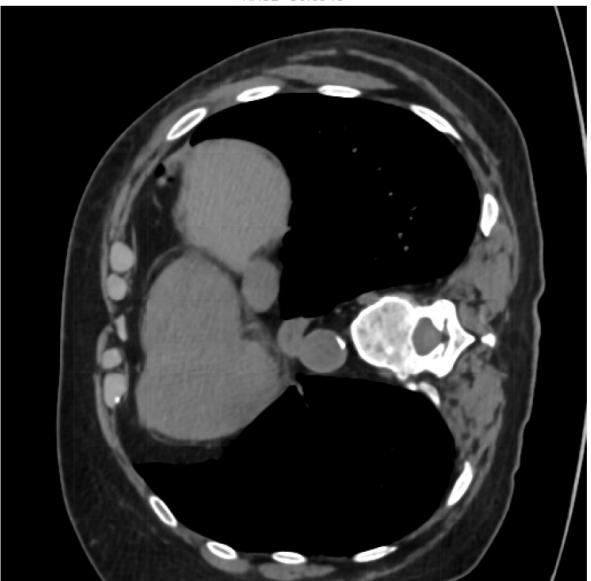
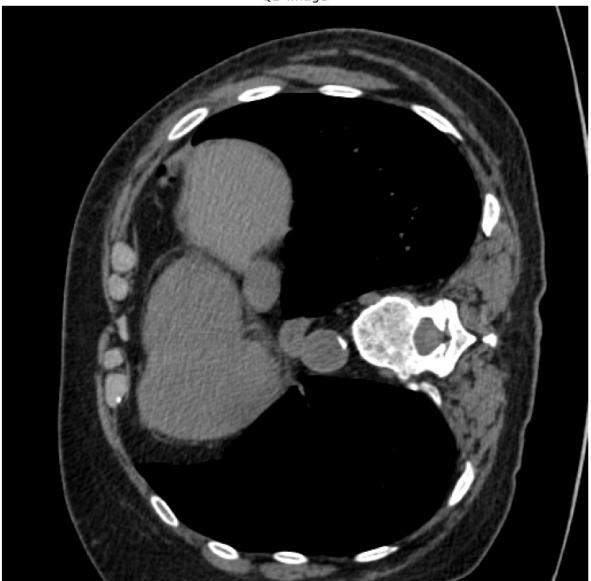






QD Image

z model Predicted Image : PSNR=36.4668  
SSIM=0.9574  
RMSE=36.0948



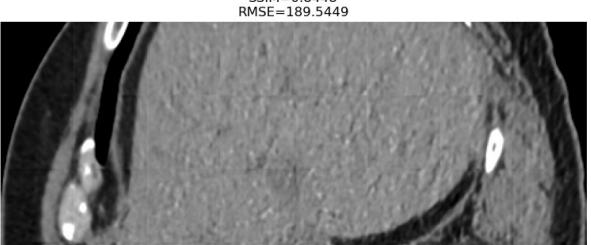
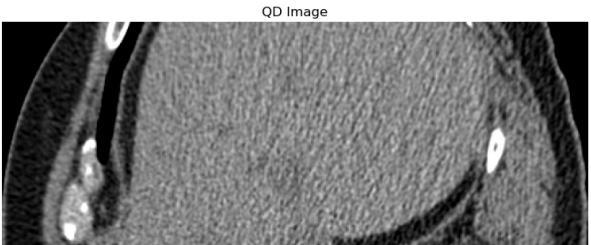
QD Image

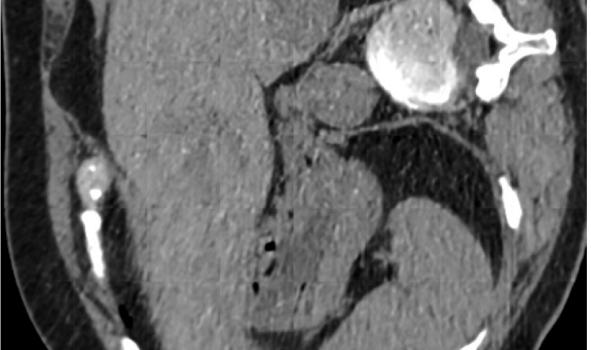
z model Predicted Image : PSNR=27.7374  
SSIM=0.8463  
RMSE=269.3885



QD Image

z model Predicted Image : PSNR=29.2641  
SSIM=0.8448  
RMSE=189.5449

























## ❖ Model 5 : W Model

```
sys.path.append('../denoising-models/hformer_pytorch')
from torchinfo import summary

from w_model import WModel

w_model = WModel(num_channels=32).cuda()
w_model.load_state_dict(torch.load('../denoising-models/hformer_pytorch/weights/model_278.pth'))
w_model.eval()
print('model summary\n', summary(w_model, input_size=(64, 64, 64, 1)))

Using cache found in C:\Users\Tarun/.cache\torch\hub\mateuszbuda_brain-segmentation-p
Using cache found in C:\Users\Tarun/.cache\torch\hub\mateuszbuda_brain-segmentation-p
model summary
=====
Layer (type:depth-idx)           Output Shape      Param #
=====
WModel                           [64, 64, 64, 1]    --
|─InputProjectionLayer: 1-1     [64, 32, 64, 64]   --
|  └Conv2d: 2-1                 [64, 32, 64, 64]   320
|─UNet: 1-2                    [64, 32, 64, 64]   --
|  └Sequential: 2-2             [64, 32, 64, 64]   --
|    └Conv2d: 3-1               [64, 32, 64, 64]   9,216
|    └BatchNorm2d: 3-2          [64, 32, 64, 64]   64
|    └ReLU: 3-3                 [64, 32, 64, 64]   --
|    └Conv2d: 3-4               [64, 32, 64, 64]   9,216
|    └BatchNorm2d: 3-5          [64, 32, 64, 64]   64
|    └ReLU: 3-6                 [64, 32, 64, 64]   --
|    └MaxPool2d: 2-3            [64, 32, 32, 32]   --
|    └Sequential: 2-4             [64, 64, 32, 32]   --
|      └Conv2d: 3-7              [64, 64, 32, 32]   18,432
|      └BatchNorm2d: 3-8          [64, 64, 32, 32]   128
```

	└ CONV2D	└ RELU	└ MAXPOOL2D
	└ ReLU: 3-9	[64, 64, 32, 32]	--
	└ Conv2d: 3-10	[64, 64, 32, 32]	36,864
	└ BatchNorm2d: 3-11	[64, 64, 32, 32]	128
	└ ReLU: 3-12	[64, 64, 32, 32]	--
	└ MaxPool2d: 2-5	[64, 64, 16, 16]	--
	└ Sequential: 2-6	[64, 128, 16, 16]	--
	└ Conv2d: 3-13	[64, 128, 16, 16]	73,728
	└ BatchNorm2d: 3-14	[64, 128, 16, 16]	256
	└ ReLU: 3-15	[64, 128, 16, 16]	--
	└ Conv2d: 3-16	[64, 128, 16, 16]	147,456
	└ BatchNorm2d: 3-17	[64, 128, 16, 16]	256
	└ ReLU: 3-18	[64, 128, 16, 16]	--
	└ MaxPool2d: 2-7	[64, 128, 8, 8]	--
	└ Sequential: 2-8	[64, 256, 8, 8]	--
	└ Conv2d: 3-19	[64, 256, 8, 8]	294,912
	└ BatchNorm2d: 3-20	[64, 256, 8, 8]	512
	└ ReLU: 3-21	[64, 256, 8, 8]	--
	└ Conv2d: 3-22	[64, 256, 8, 8]	589,824
	└ BatchNorm2d: 3-23	[64, 256, 8, 8]	512
	└ ReLU: 3-24	[64, 256, 8, 8]	--
	└ MaxPool2d: 2-9	[64, 256, 4, 4]	--
	└ Sequential: 2-10	[64, 512, 4, 4]	--
	└ Conv2d: 3-25	[64, 512, 4, 4]	1,179,648
	└ BatchNorm2d: 3-26	[64, 512, 4, 4]	1,024
	└ ReLU: 3-27	[64, 512, 4, 4]	--
	└ Conv2d: 3-28	[64, 512, 4, 4]	2,359,296
	└ BatchNorm2d: 3-29	[64, 512, 4, 4]	1,024
	└ ReLU: 3-30	[64, 512, 4, 4]	--
	└ ConvTranspose2d: 2-11	[64, 256, 8, 8]	524,544
	└ Sequential: 2-12	[64, 256, 8, 8]	--
	└ Conv2d: 3-31	[64, 256, 8, 8]	1,179,648
	└ BatchNorm2d: 3-32	[64, 256, 8, 8]	512
	└ ReLU: 3-33	[64, 256, 8, 8]	--
	└ Conv2d: 3-34	[64, 256, 8, 8]	589,824
	└ BatchNorm2d: 3-35	[64, 256, 8, 8]	512
	└ ReLU: 3-36	[64, 256, 8, 8]	--
	└ ConvTranspose2d: 2-13	[64, 128, 16, 16]	131,200

```
w_prediction_patches = []

with torch.no_grad():
    for i, data in enumerate(noisy_image_patches_array):
        noisy = data

        predictions = w_model(torch.unsqueeze(torch.from_numpy(noisy.numpy()), dim=0).to('cuda')).cpu()

        w_prediction_patches.append(predictions.detach().cpu())

w_prediction_patches = np.concatenate(w_prediction_patches, axis=0)

w_predictions = np.expand_dims(reconstruct_image_from_patches(w_prediction_patches[0:64], 8), axis=0)

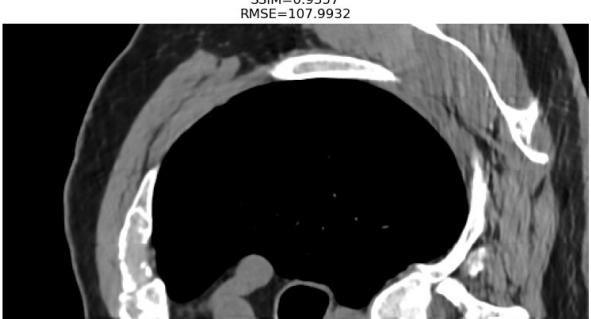
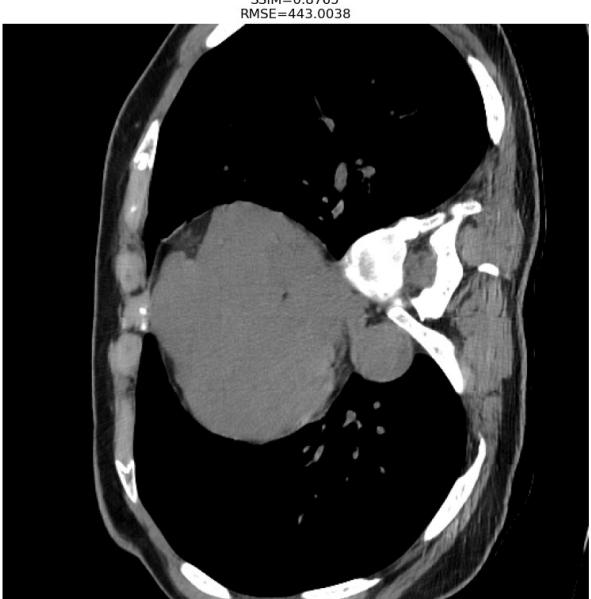
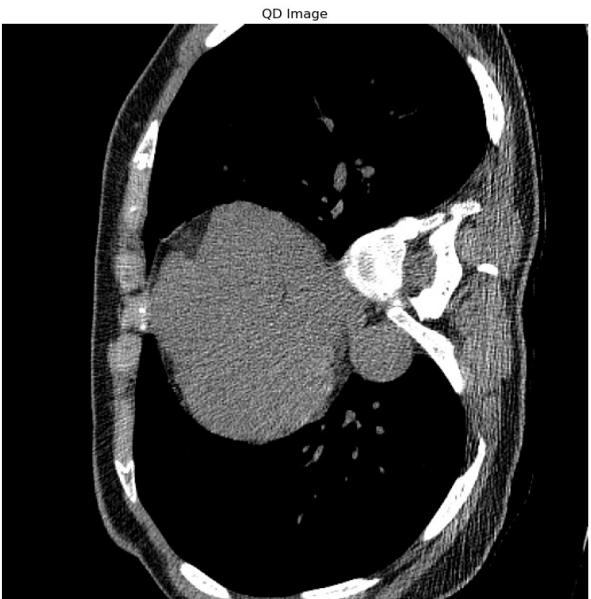
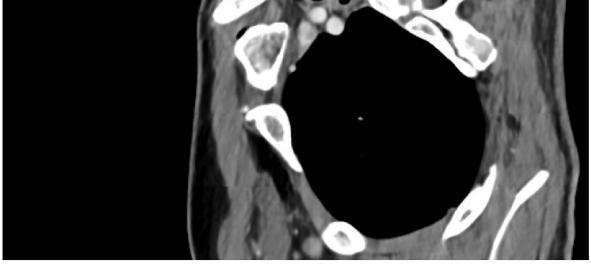
for i in range(1, int(w_prediction_patches.shape[0] / 64)):
```

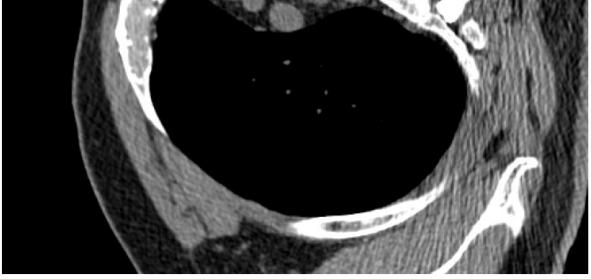
```
reconstructed_image = reconstruct_image_from_patches(w_prediction_patches[i * 64 : i * 64 + 64], num_patches)
reconstructed_image = np.expand_dims(reconstructed_image, axis=0)

w_predictions= np.append(w_predictions, reconstructed_image, axis=0)

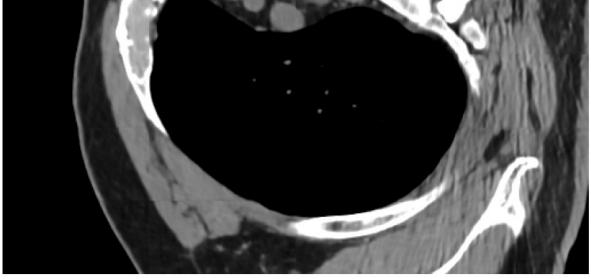
visualize_predictions(w_model, noisy_array, len(noisy_array), w_predictions, "w model")
```







QD Image



w model Predicted Image : PSNR=27.1027  
SSIM=0.8404  
RMSE=311.7794



QD Image



w model Predicted Image : PSNR=34.8294  
SSIM=0.9612  
RMSE=52.6238

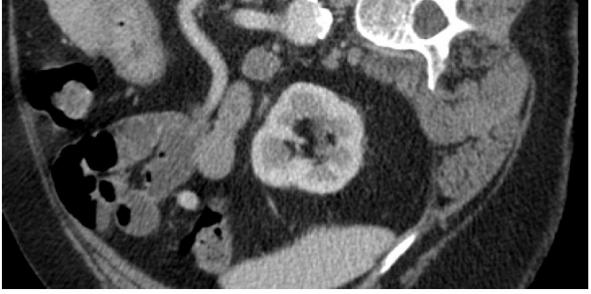


QD Image

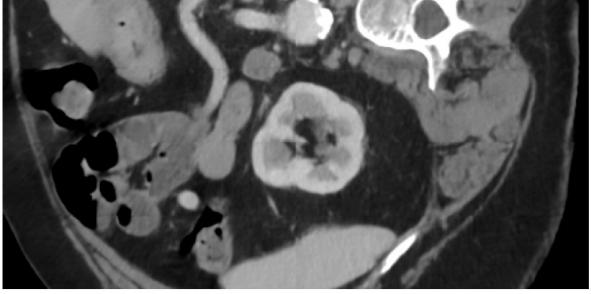


w model Predicted Image : PSNR=31.3435  
SSIM=0.8579  
RMSE=117.4285





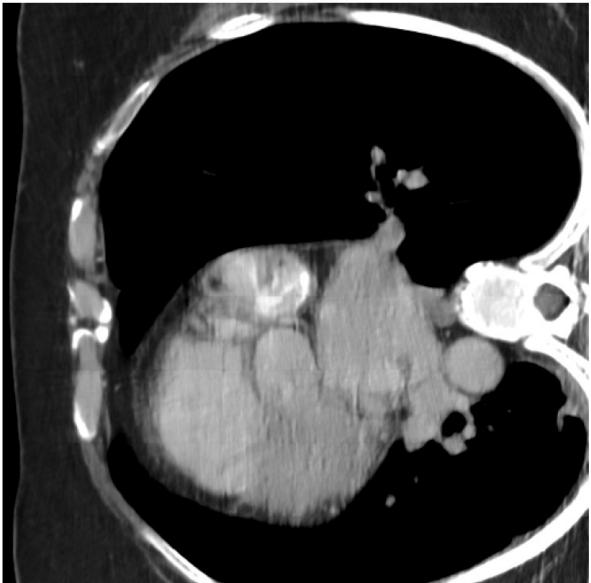
QD Image



w model Predicted Image : PSNR=30.8824  
SSIM=0.8948  
RMSE=130.5819



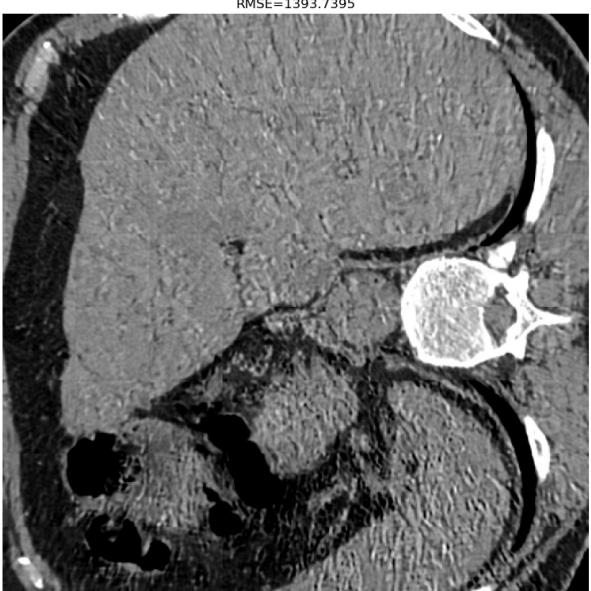
QD Image



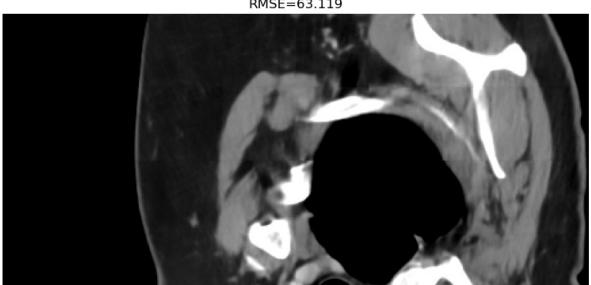
w model Predicted Image : PSNR=20.5994  
SSIM=0.7082  
RMSE=1393.7395

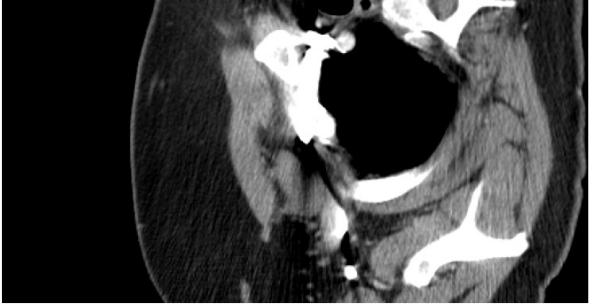


QD Image

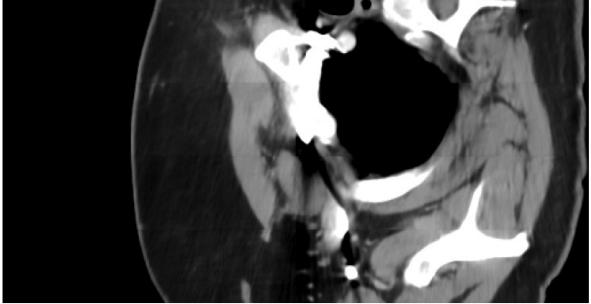


w model Predicted Image : PSNR=34.0396  
SSIM=0.9246  
RMSE=63.119





QD Image



w model Predicted Image : PSNR=33.2459  
SSIM=0.903  
RMSE=75.7763



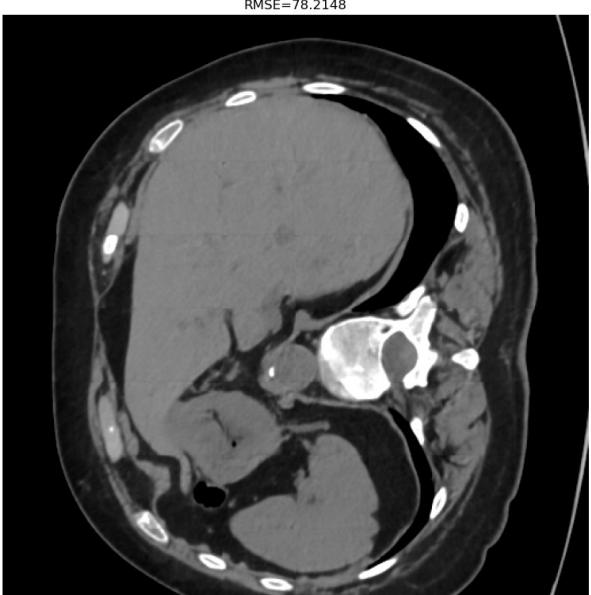
QD Image



w model Predicted Image : PSNR=33.1083  
SSIM=0.8874  
RMSE=78.2148

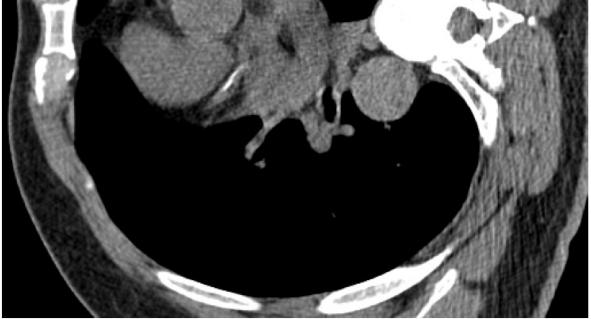


QD Image



w model Predicted Image : PSNR=33.3518  
SSIM=0.9434  
RMSE=73.9496

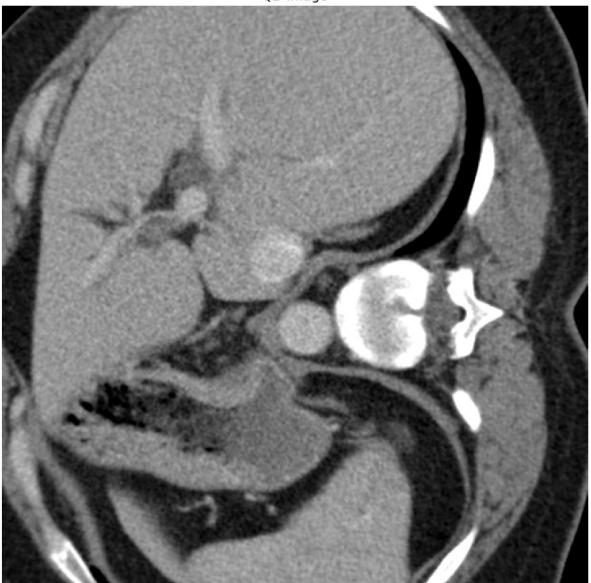




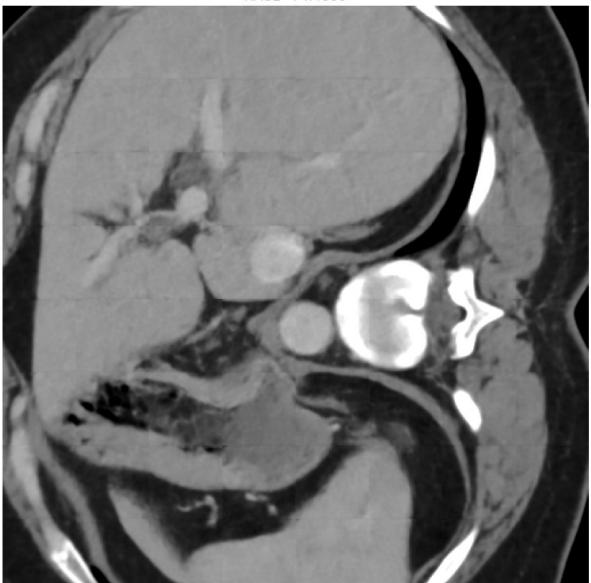
QD Image



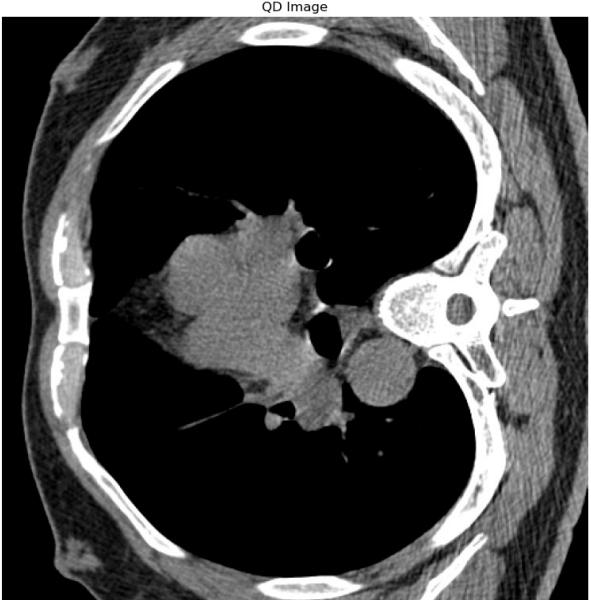
w model Predicted Image : PSNR=33.3203  
SSIM=0.8582  
RMSE=74.4886



QD Image



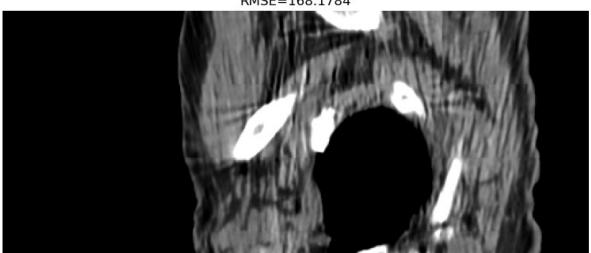
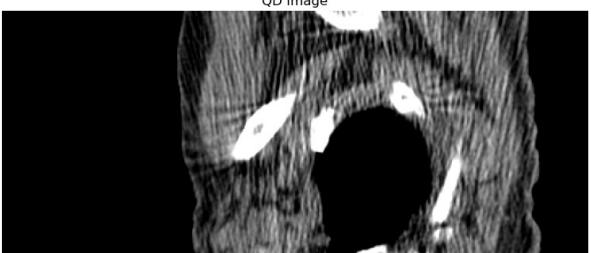
w model Predicted Image : PSNR=33.0199  
SSIM=0.9401  
RMSE=79.8241

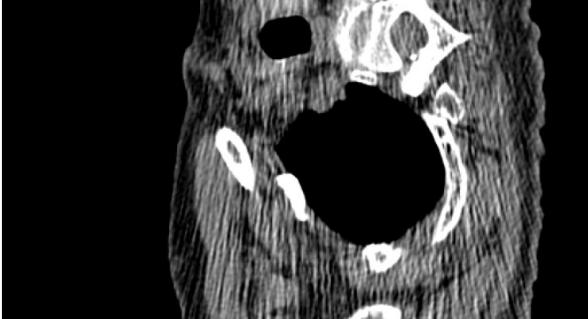


QD Image

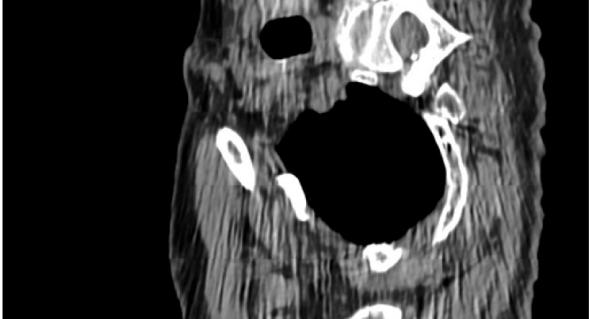


w model Predicted Image : PSNR=29.7835  
SSIM=0.9513  
RMSE=168.1784

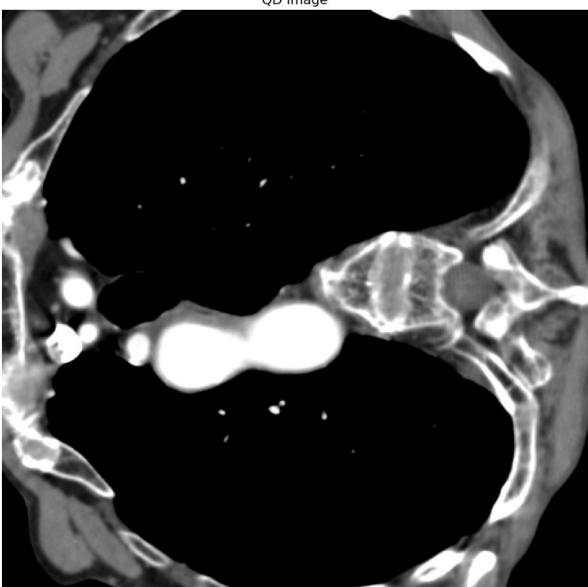




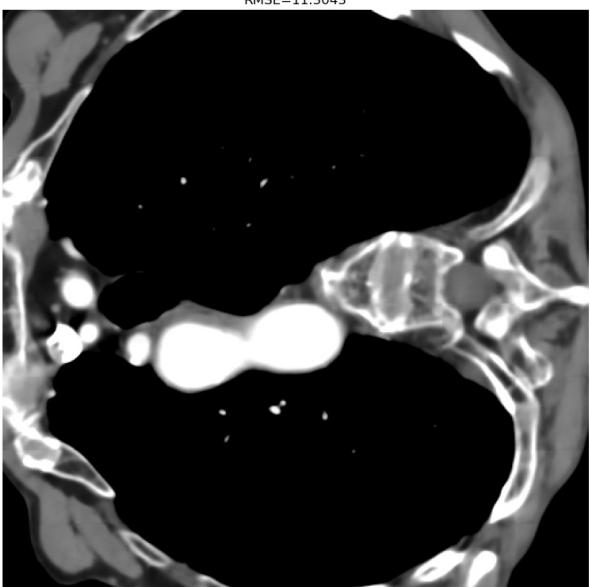
QD Image



w model Predicted Image : PSNR=41.5088  
SSIM=0.9928  
RMSE=11.3043



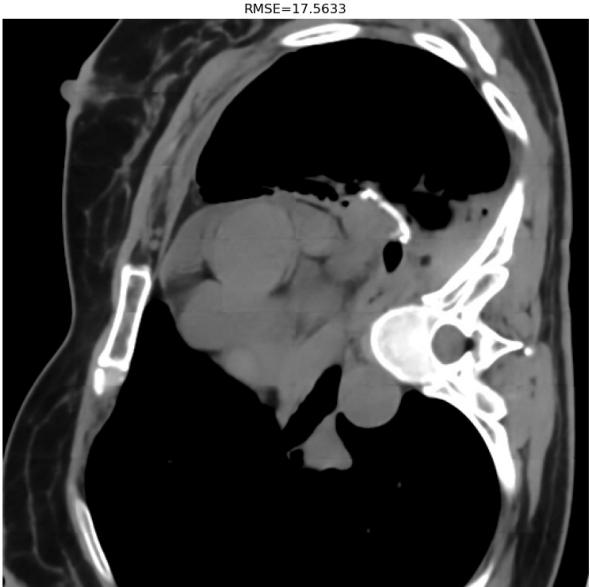
QD Image



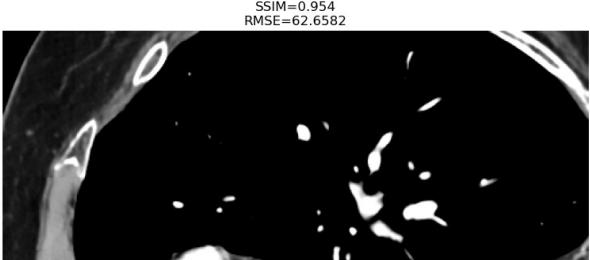
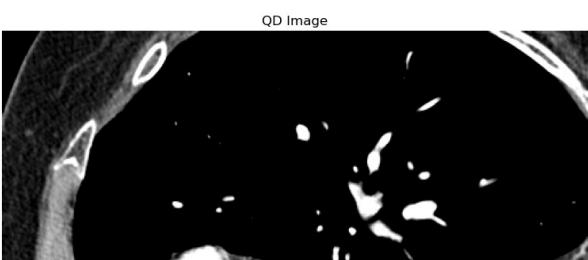
w model Predicted Image : PSNR=39.5951  
SSIM=0.975  
RMSE=17.5633

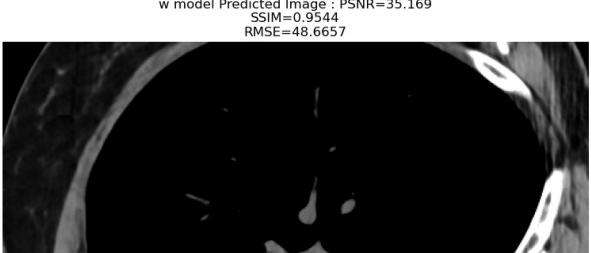
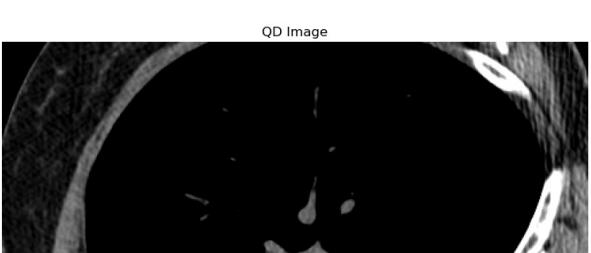
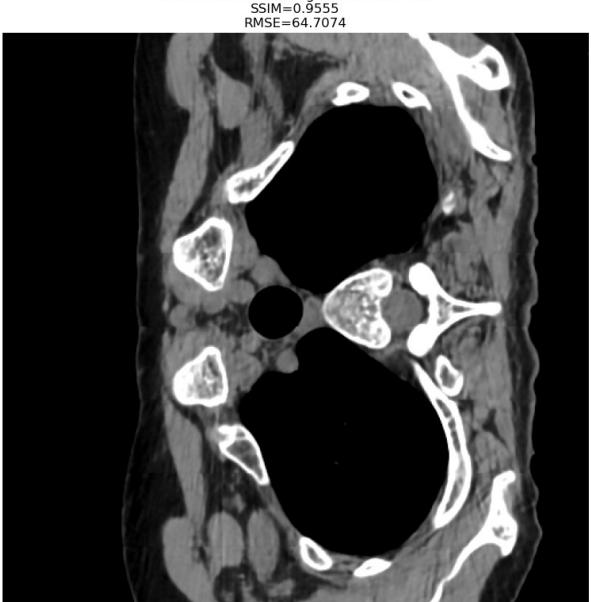
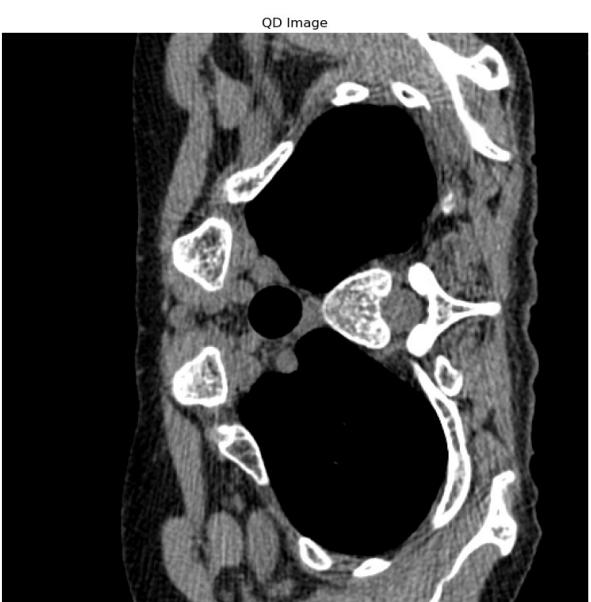
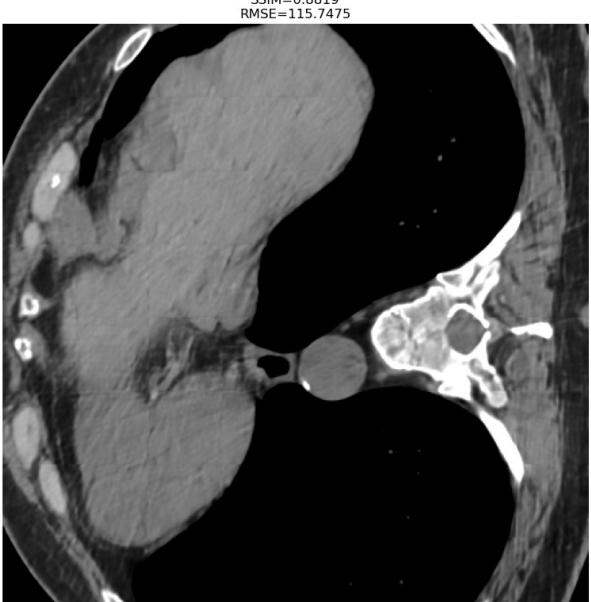
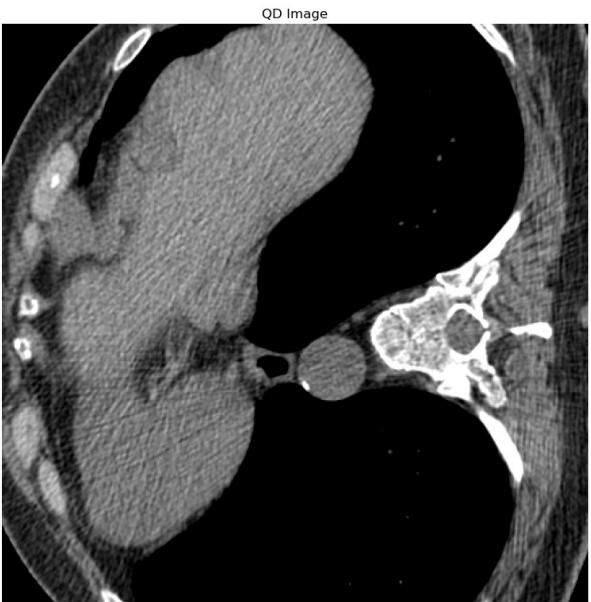
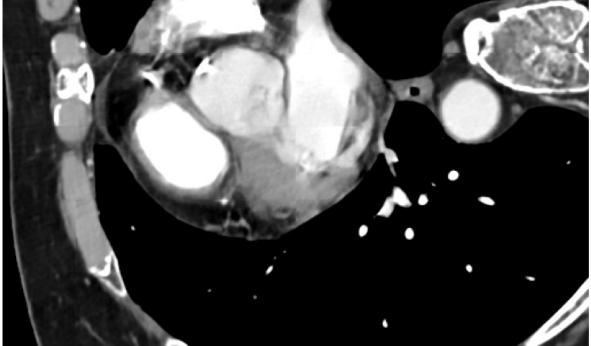
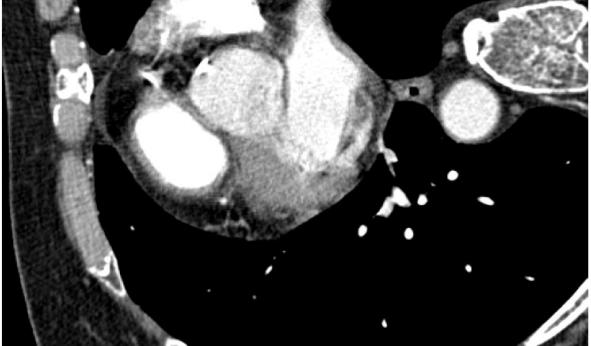


QD Image



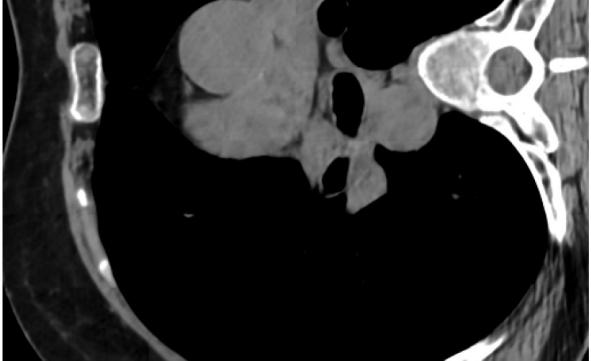
w model Predicted Image : PSNR=34.0714  
SSIM=0.954  
RMSE=62.6582



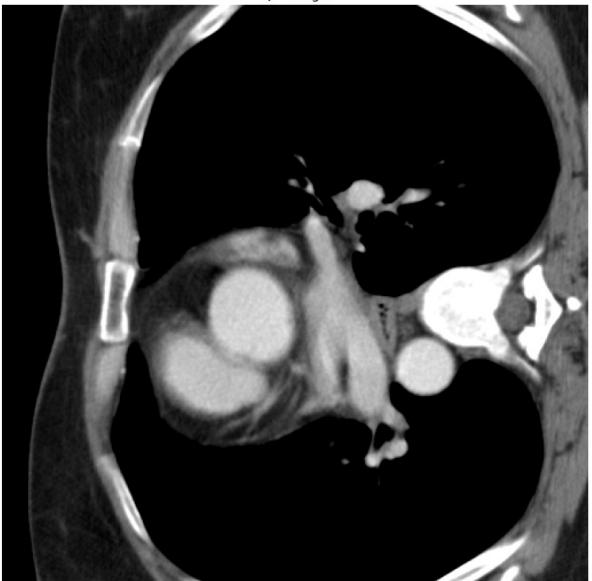




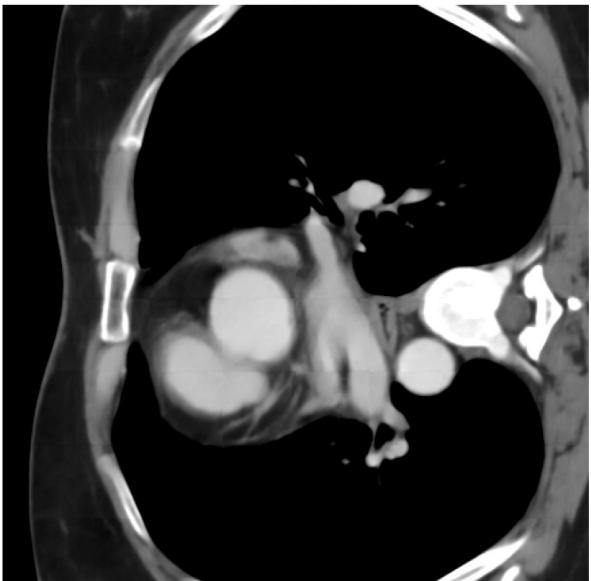
QD Image



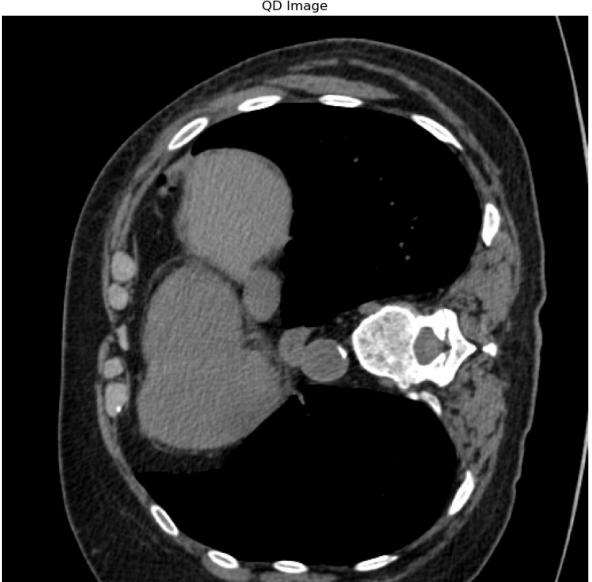
w model Predicted Image : PSNR=39.8881  
SSIM=0.9756  
RMSE=16.4175



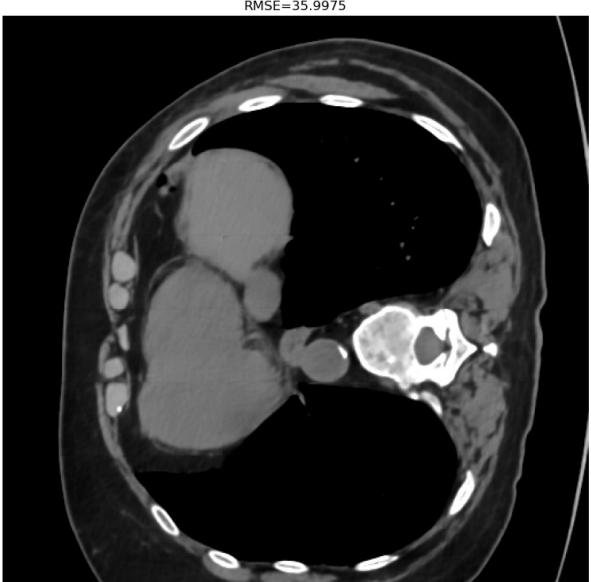
QD Image



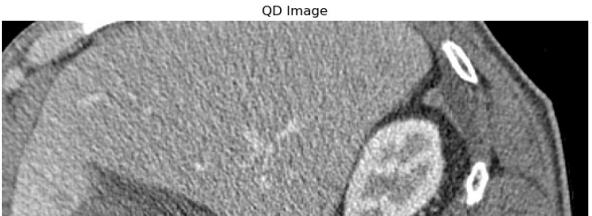
w model Predicted Image : PSNR=36.4785  
SSIM=0.9499  
RMSE=35.9975

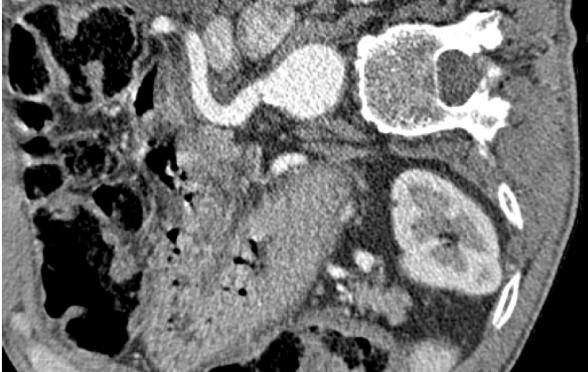


QD Image

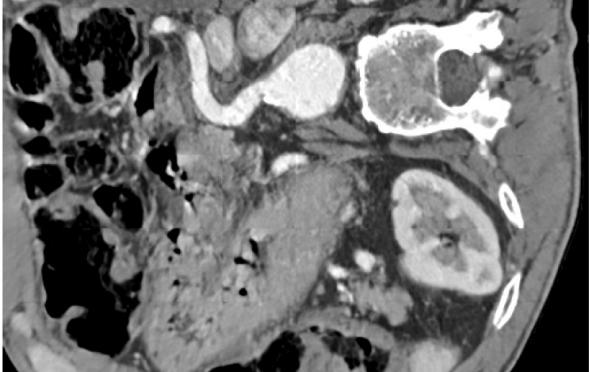


w model Predicted Image : PSNR=27.106  
SSIM=0.8185  
RMSE=311.5448





QD Image



w model Predicted Image : PSNR=28.7437  
SSIM=0.8207  
RMSE=213.6733























```
# Wavelet with wmodel

# W model with wavelet
from pytorch_wavelets import DTCWTForward, DTCWTInverse
dwt = DTCWTForward(J=3).cuda()
idwt = DTCWTInverse().cuda()
wavelet_w_model_prediction_patches = []

with torch.no_grad():
    for i in range(noisy_image_patches_array.shape[0] // 64):
        noisy = noisy_image_patches_array[i * 64 : i * 64 + 64]
        noisy = torch.from_numpy(noisy.cpu().numpy()).to('cuda')
        prediction = w_model(noisy)

        prediction_img = torch.transpose(prediction, 1, 3)
        transposed_noisy_image = torch.transpose(noisy, 1, 3)
```

```

prediction_approx, prediction_high_freq = dwt(prediction_img.cuda())
prediction_high_freq_low, prediction_high_freq_mid, prediction_high_freq_coarse = prediction_hi

noisy_approx, noisy_high_freq = dwt(transposed_noisy_image.cuda())

noisy_high_freq_fine, noisy_high_freq_mid, noisy_high_freq_coarse = noisy_high_freq[0], noisy_hi

reconstructed_prediction_image_with_high_freq_swap = idwt((noisy_approx, noisy_high_freq))
reconstructed_prediction_image_with_high_freq_swap = torch.transpose(reconstructed_prediction_im

#denoise_high_freq = torch.transpose(torch.from_numpy(denoised_high_freq), 1, 3)
#denoised_noisy_approx, denoised_high_freq = dwt(denoise_high_freq.cuda())
#noisy_high_freq_fine, noisy_high_freq_mid, noisy_high_freq_coarse = denoised_high_freq[0], denc

wavelet_high_freq_swapped = [None] * 3
wavelet_high_freq_swapped[0] = noisy_high_freq_fine
wavelet_high_freq_swapped[1] = prediction_high_freq_mid
wavelet_high_freq_swapped[2] = prediction_high_freq_coarse

reconstructed_prediction_image = idwt((prediction_approx, wavelet_high_freq_swapped))
reconstructed_prediction_image = torch.transpose(reconstructed_prediction_image, 1, 3)
wavelet_w_model_prediction_patches.append(reconstructed_prediction_image.detach().cpu().numpy())

wavelet_w_predictions = [None] * noisy_array.shape[0]
def reconstruct(patches, num_images):
    num_patches_per_image = patches.shape[0] // num_images
    for i in range(num_images):

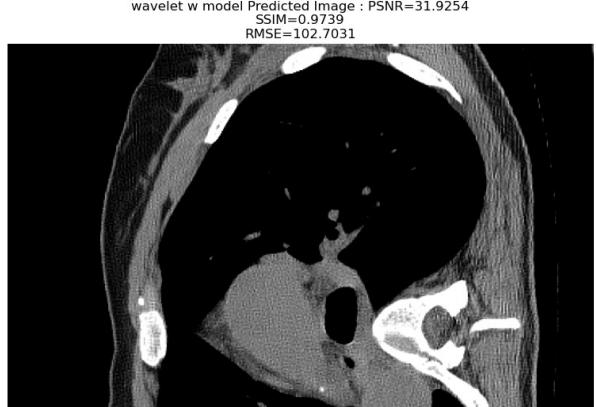
        image_patches = patches[i * num_patches_per_image:i * num_patches_per_image + num_patches_per_im

        reconstruct_image = (reconstruct_image_from_patches(image_patches, 8))
        wavelet_w_predictions[i] = reconstruct_image
wavelet_w_model_prediction_patches = np.concatenate(wavelet_w_model_prediction_patches)
reconstruct(wavelet_w_model_prediction_patches, 28)

WARNING:tensorflow:From C:\Users\Tarun\AppData\Local\Temp\ipykernel_11168\1428046907.
Instructions for updating:
Use tf.identity instead.

```

```
visualize_predictions(w_model, noisy_array, len(noisy_array), wavelet_w_predictions, "wavelet w model")
```

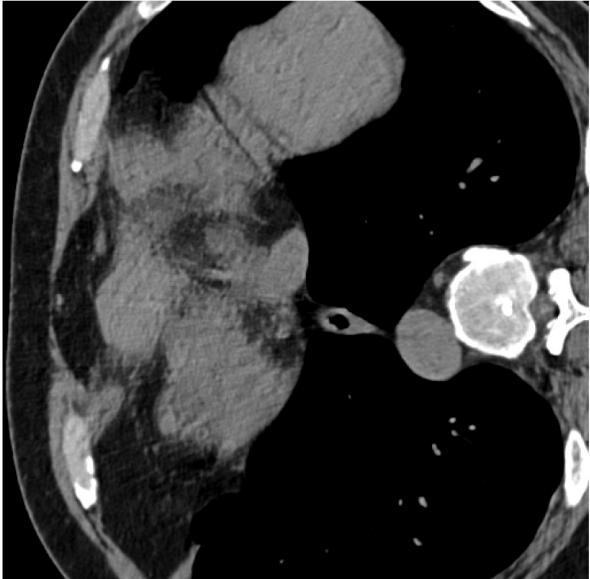




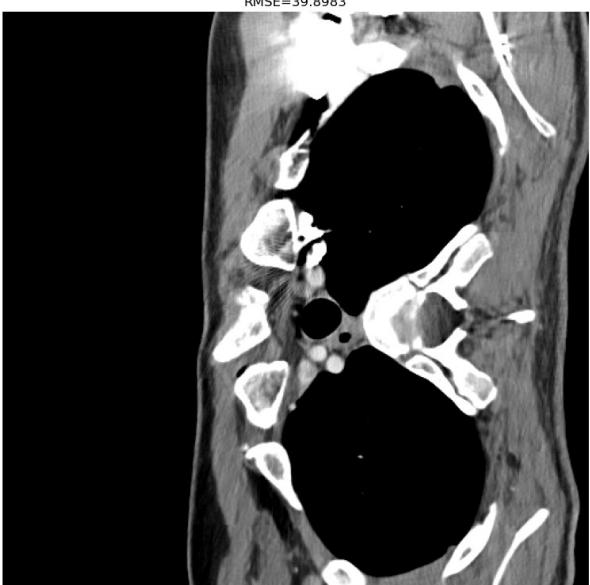
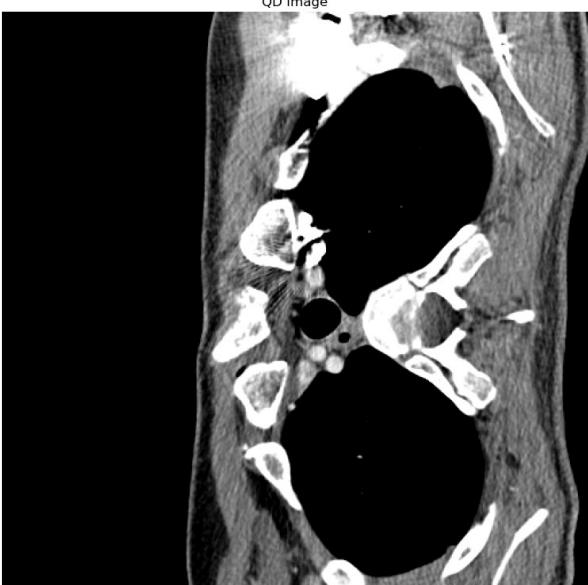
QD Image



wavelet w model Predicted Image : PSNR=33.8063  
SSIM=0.9421  
RMSE=66.6025



wavelet w model Predicted Image : PSNR=36.0317  
SSIM=0.9785  
RMSE=39.8983



wavelet w model Predicted Image : PSNR=30.9012  
SSIM=0.967  
RMSE=130.0179





QD Image



wavelet w model Predicted Image : PSNR=37.9865  
SSIM=0.9336  
RMSE=25.4373



QD Image



wavelet w model Predicted Image : PSNR=34.1222  
SSIM=0.9626  
RMSE=61.9297

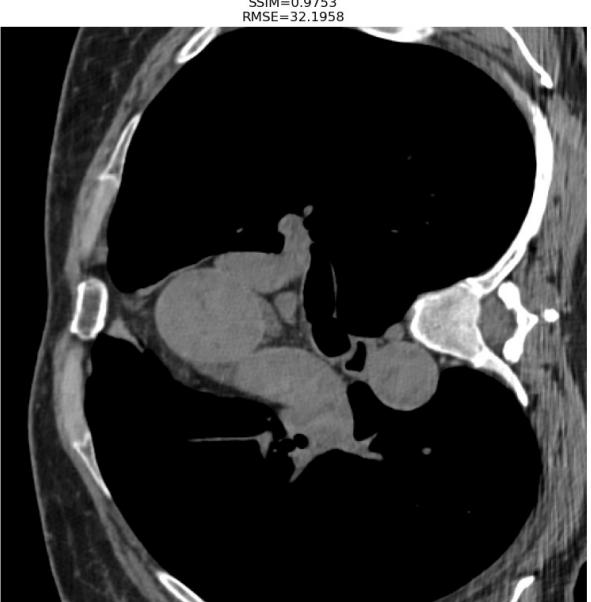
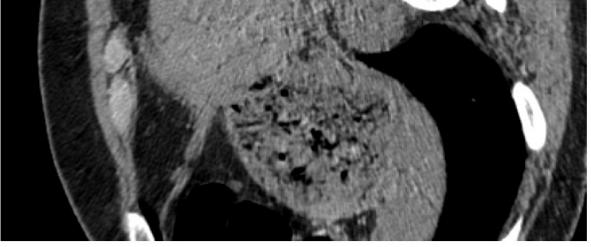
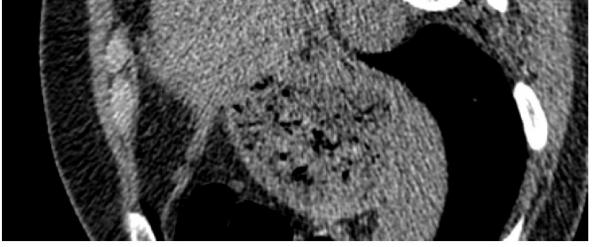


QD Image



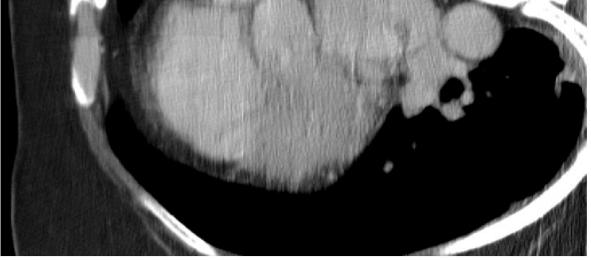
wavelet w model Predicted Image : PSNR=29.0101  
SSIM=0.9011  
RMSE=200.9599







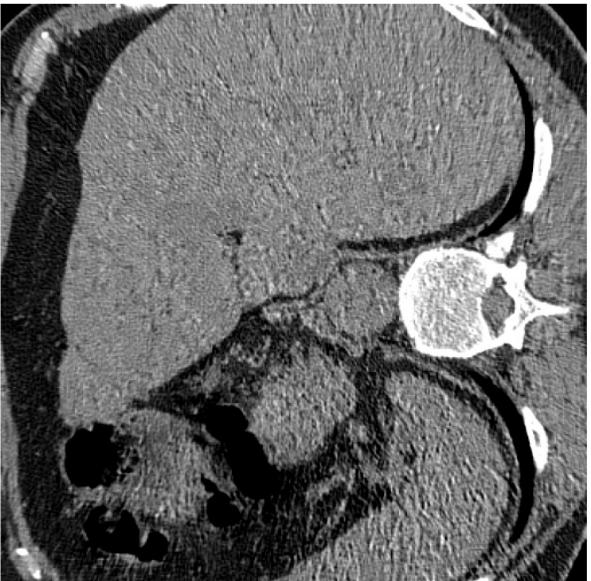
QD Image



wavelet w model Predicted Image : PSNR=23.1912  
SSIM=0.8522  
RMSE=767.359



QD Image



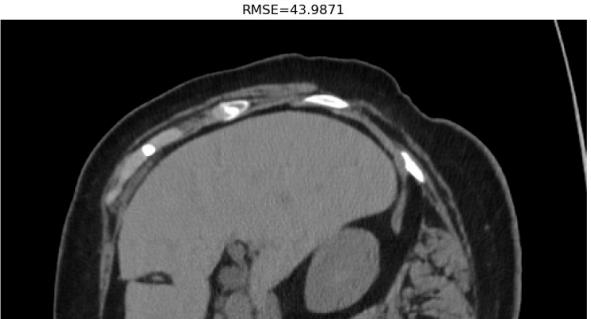
wavelet w model Predicted Image : PSNR=36.6535  
SSIM=0.9619  
RMSE=34.5753



QD Image



wavelet w model Predicted Image : PSNR=35.6079  
SSIM=0.9456  
RMSE=43.9871

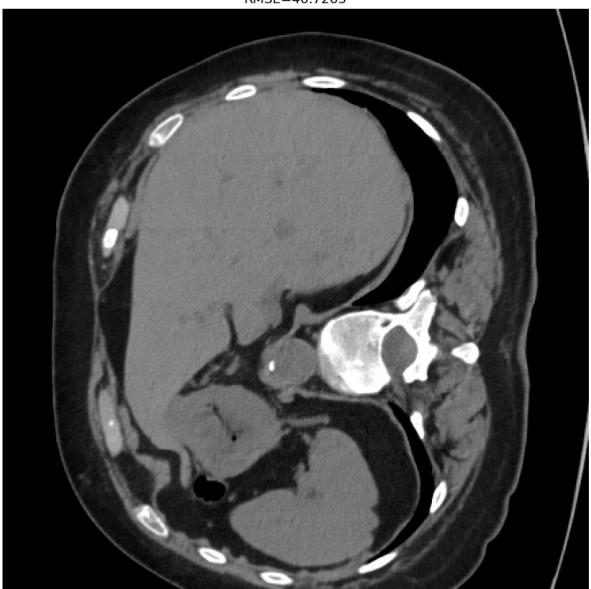




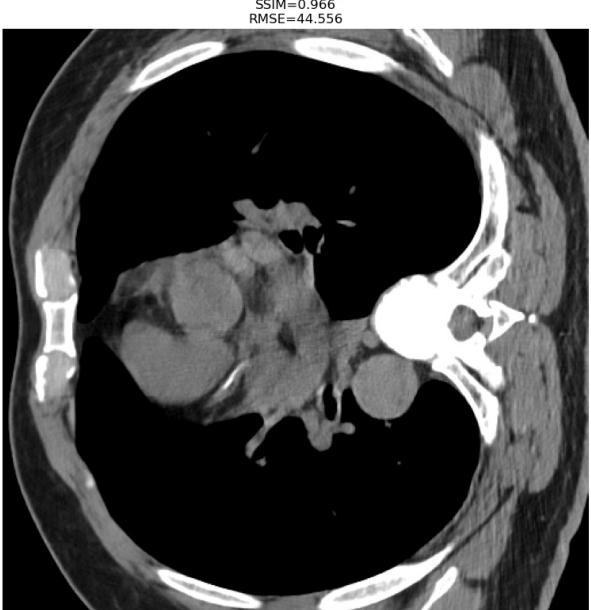
QD Image



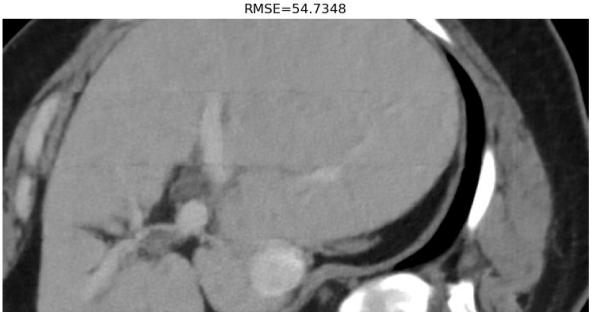
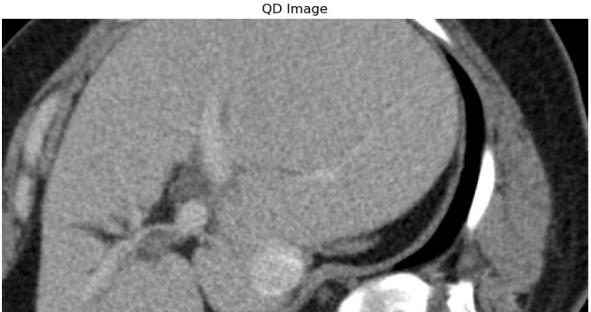
wavelet w model Predicted Image : PSNR=35.3456  
SSIM=0.9355  
RMSE=46.7263

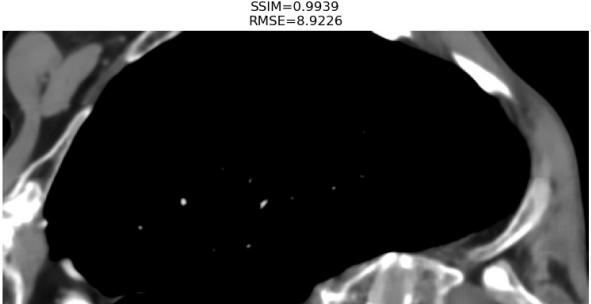
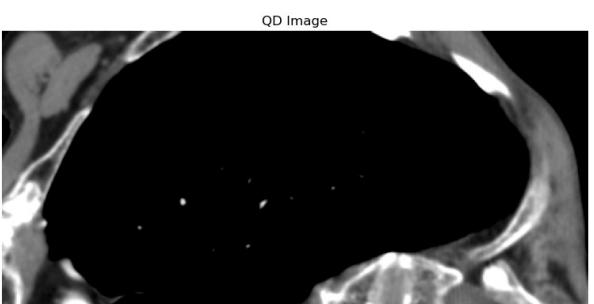
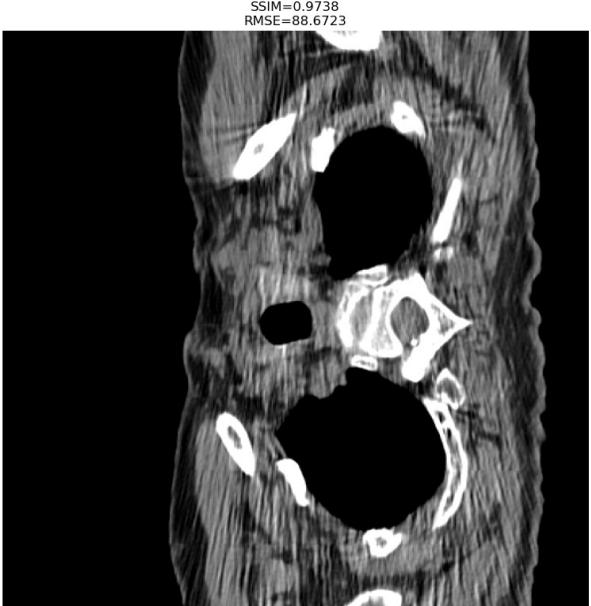
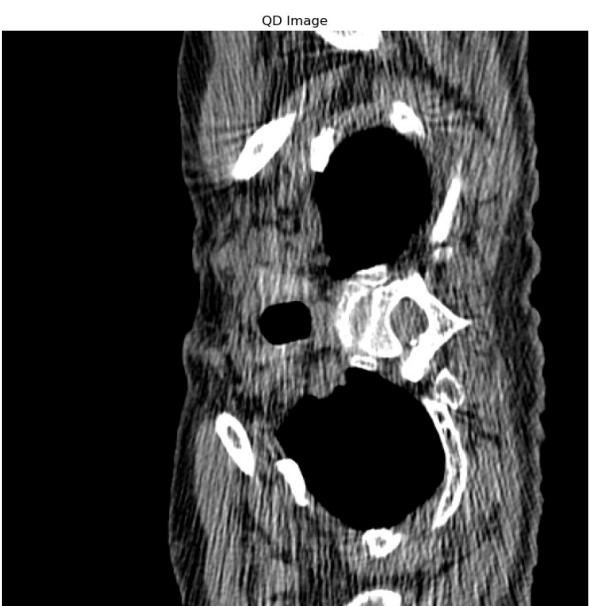
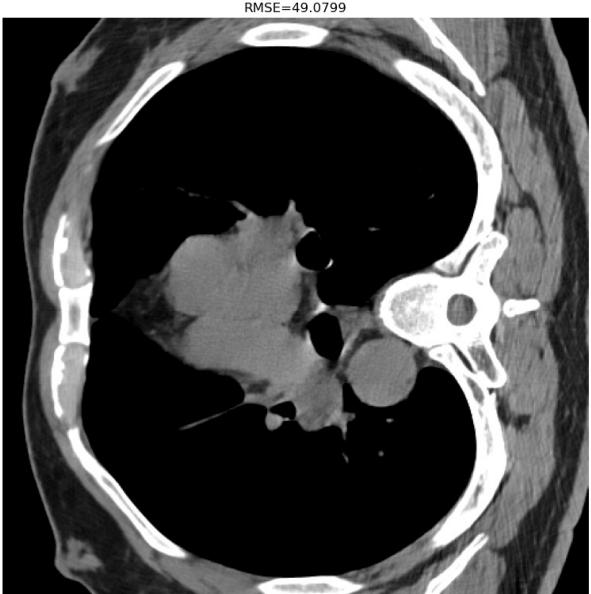
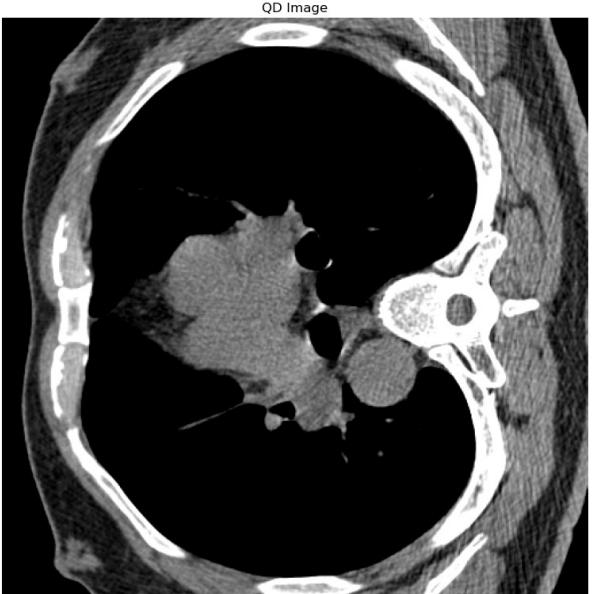
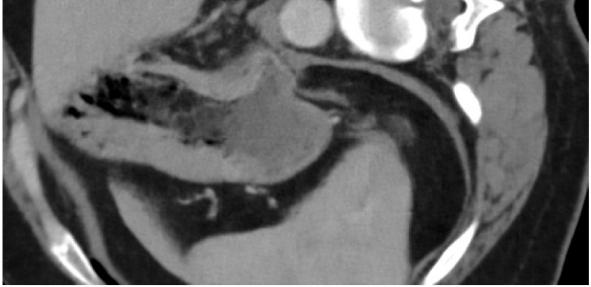
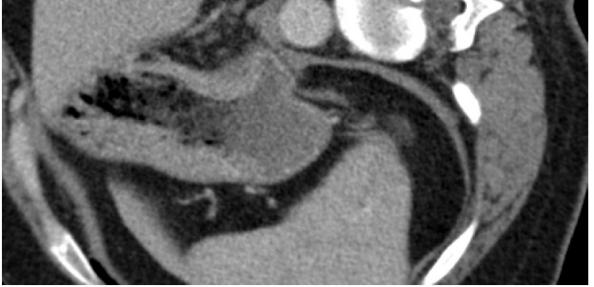


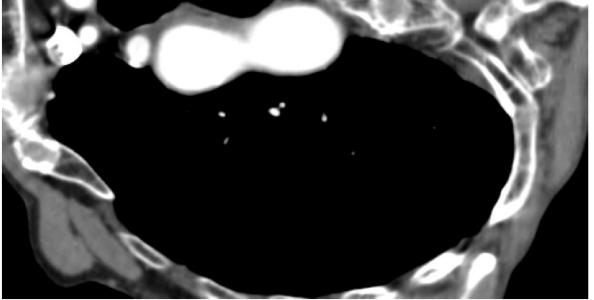
wavelet w model Predicted Image : PSNR=35.5521  
SSIM=0.966  
RMSE=44.556



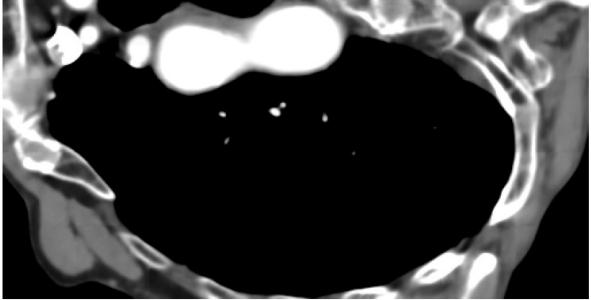
wavelet w model Predicted Image : PSNR=34.6586  
SSIM=0.8981  
RMSE=54.7348



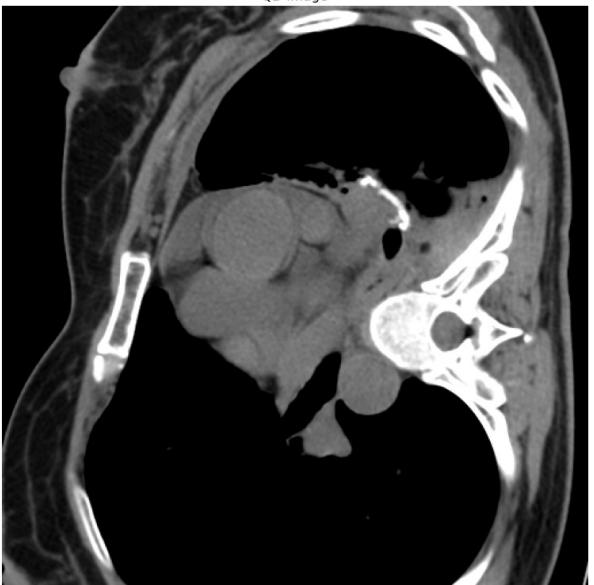




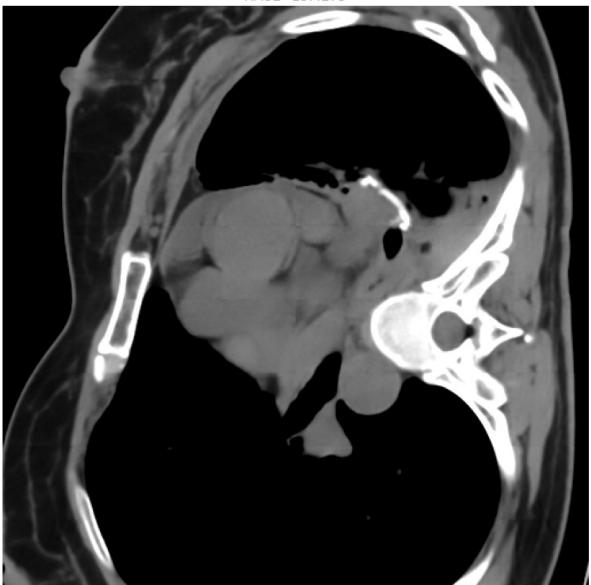
QD Image



wavelet w model Predicted Image : PSNR=40.7613  
SSIM=0.9798  
RMSE=13.4275



QD Image



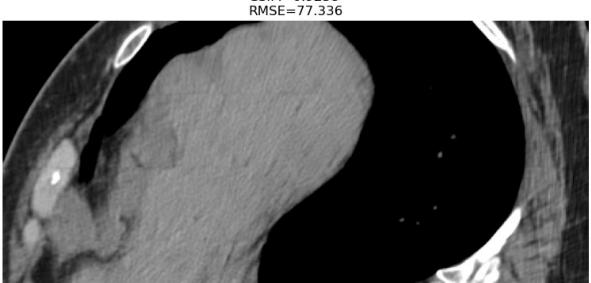
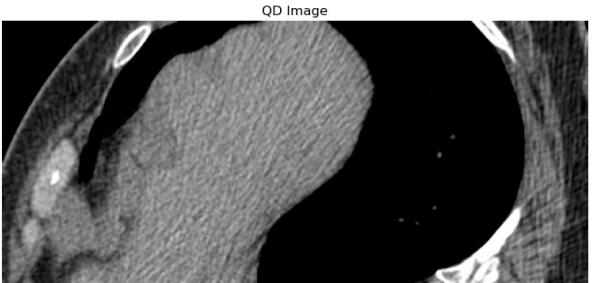
wavelet w model Predicted Image : PSNR=36.2547  
SSIM=0.9727  
RMSE=37.9012

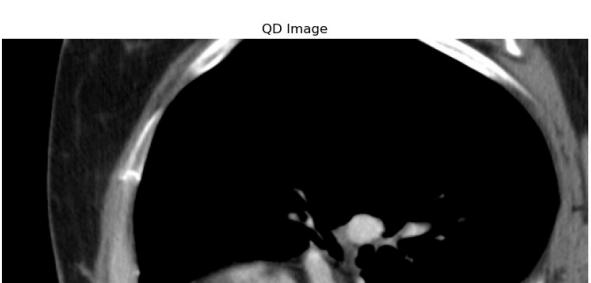
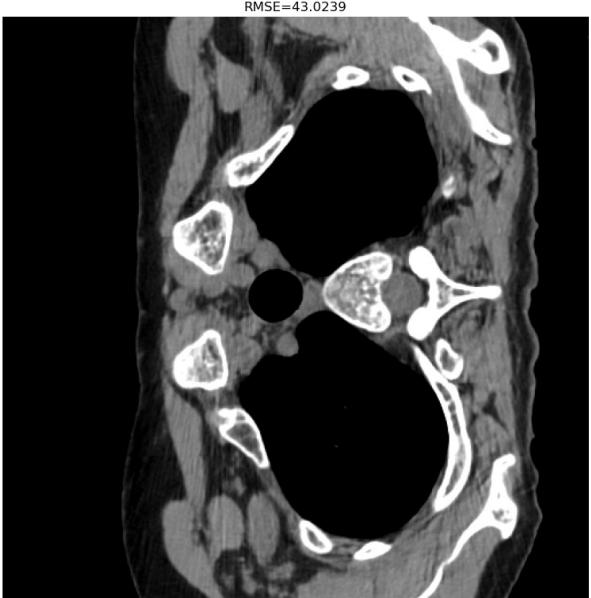
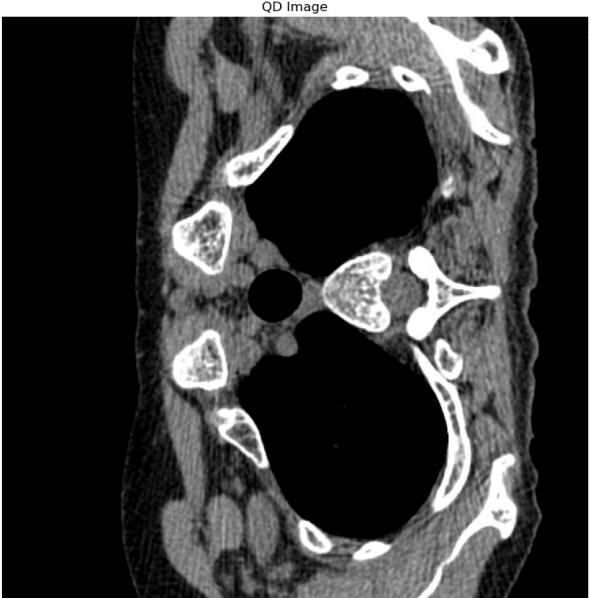
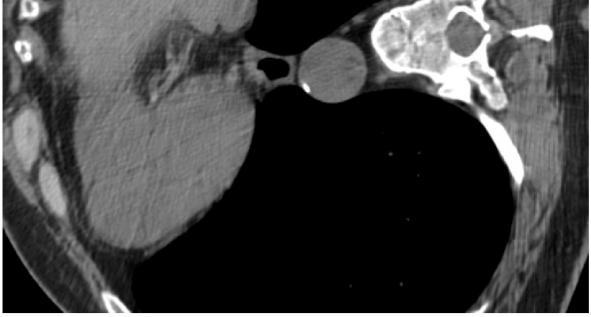
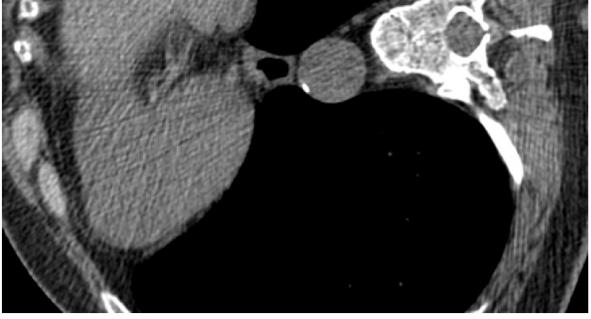


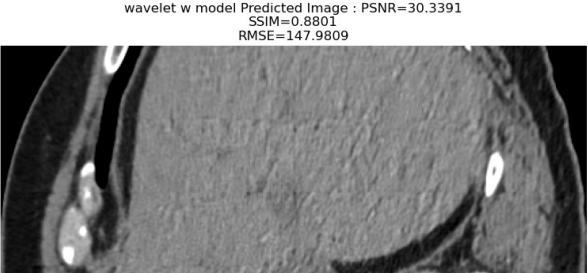
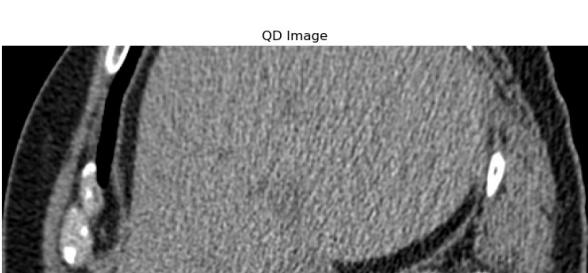
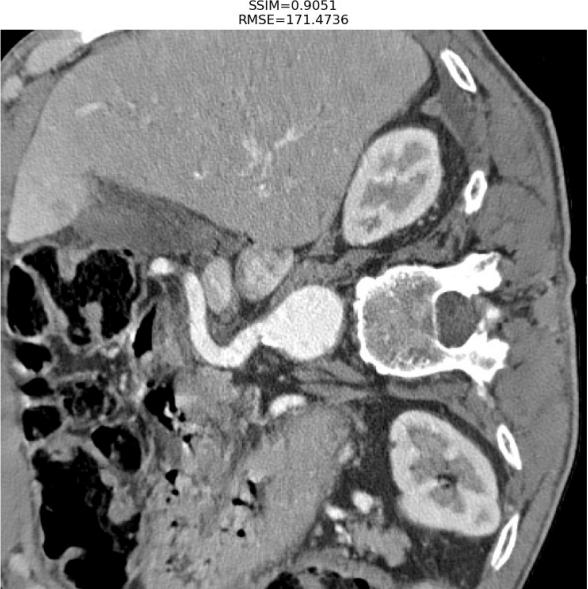
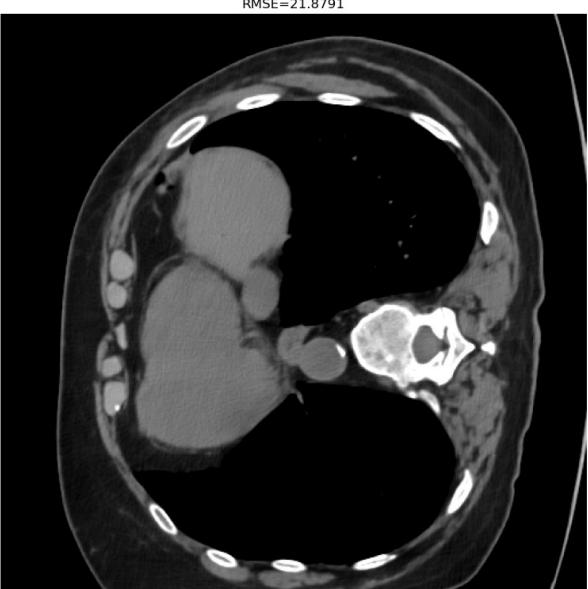
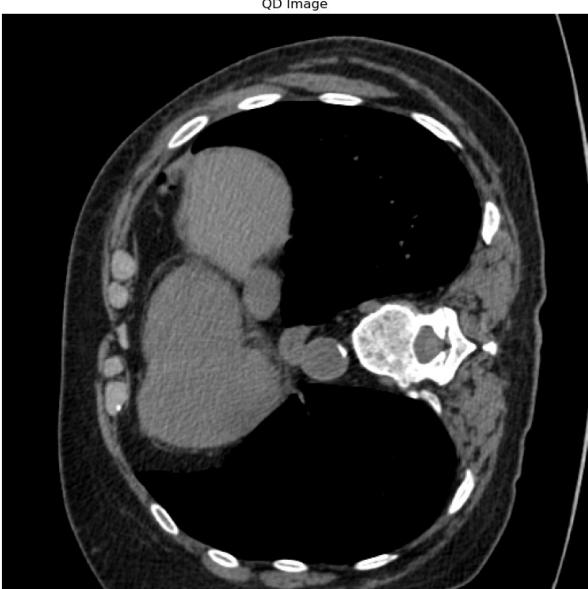
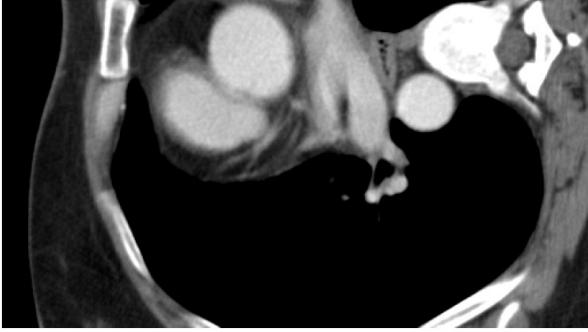
QD Image

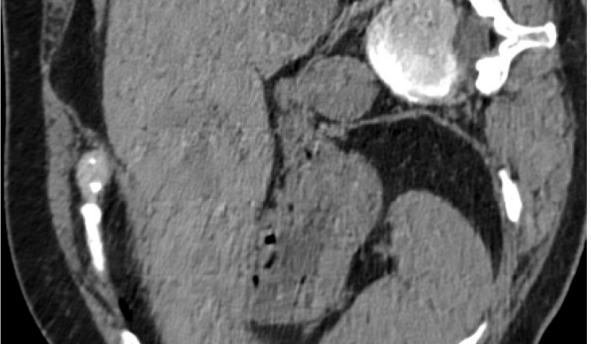
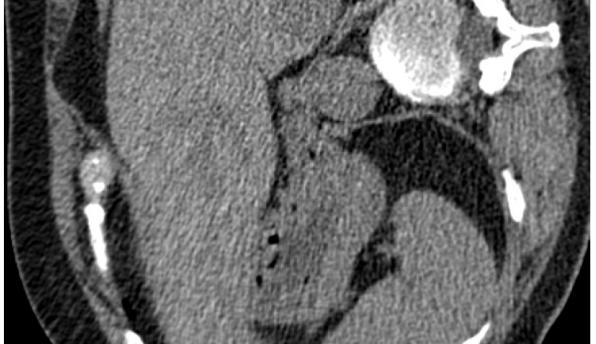


wavelet w model Predicted Image : PSNR=33.1574  
SSIM=0.9238  
RMSE=77.336





























## ❖ Part V : Side by side comparison of all models

```
from prettytable import PrettyTable

pt = PrettyTable()
pt.field_names = ["Model", "PSNR", "SSIM", "MSE"]

hformer_metrics = get_average_metrics(hformer_predictions, noisy_array)
x_metrics= get_average_metrics(np.concatenate(x_model_prediction_patches, axis=0), noisy_array)
y_metrics= get_average_metrics(np.concatenate(y_model_prediction_patches, axis=0), noisy_array)
sa_metrics = get_average_metrics(sa_predictions, noisy_array)
z_metrics = get_average_metrics(z_predictions, noisy_array)
w_metrics = get_average_metrics(w_predictions, noisy_array)
wavelet_w_metrics = get_average_metrics(wavelet_w_predictions, noisy_array)
```

```
Predicted average gt-predicted PSNR -> 32.99037891246741
Predicted average gt-predicted SSIM -> 0.916629681798598
Predicted average gt-predicted MSE-> 185.87344353967302
Predicted average gt-predicted PSNR -> 8.728047851372207
Predicted average gt-predicted SSIM -> 0.332313969732048
Predicted average gt-predicted MSE-> 22949.383764070844
Predicted average gt-predicted PSNR -> 4.6170096551702695
Predicted average gt-predicted SSIM -> 0.0001014814520292278
Predicted average gt-predicted MSE-> 55501.08398167486
Predicted average gt-predicted PSNR -> 26.12000539421971
Predicted average gt-predicted SSIM -> 0.7756457741615084
Predicted average gt-predicted MSE-> 413.72842258304127
Predicted average gt-predicted PSNR -> 32.877298514970875
Predicted average gt-predicted SSIM -> 0.9232649297146059
Predicted average gt-predicted MSE-> 145.19556484099186
```

```

Predicted average gt-predicted PSNR -> 32.471087393232125
Predicted average gt-predicted SSIM -> 0.9093229017453929
Predicted average gt-predicted MSE-> 165.67989832823488
Predicted average gt-predicted PSNR -> 34.73367343929717
Predicted average gt-predicted SSIM -> 0.9482066967907415
Predicted average gt-predicted MSE-> 88.90021109734904

```

```

pt.add_row(["Original X-y pairs (No Model)", '-', '-', "-"])
pt.add_row(["Hformer", str(hformer_metrics[0]), str(hformer_metrics[1]), str(round(hformer_metrics[2], 4))]
pt.add_row(["X Model", str(x_metrics[0]), str(x_metrics[1]), str(round(x_metrics[2], 4))])
pt.add_row(["Y Model", str(y_metrics[0]), str(y_metrics[1]), str(round(y_metrics[2], 4))])
pt.add_row(["SA Model", str(sa_metrics[0]), str(sa_metrics[1]), str(round(sa_metrics[2], 4))])
pt.add_row(["Z Model", str(z_metrics[0]), str(z_metrics[1]), str(round(z_metrics[2], 4))])
pt.add_row(["W Model", str(w_metrics[0]), str(w_metrics[1]), str(round(w_metrics[2], 4))])
pt.add_row(["W Wavelet Model", str(wavelet_w_metrics[0]), str(wavelet_w_metrics[1]), str(round(wavelet_w_]

print(pt)

+-----+-----+-----+-----+
|      Model | PSNR | SSIM | MSE  |
+-----+-----+-----+-----+
| Original X-y pairs (No Model) | -    | -    | -    |
| Hformer          | 32.9904 | 0.9166 | 185.8734 |
| X Model          | 8.728   | 0.3323 | 22949.3838 |
| Y Model          | 4.617   | 0.0001 | 55501.084  |
| SA Model         | 26.12   | 0.7756 | 413.7284  |
| Z Model          | 32.8773 | 0.9233 | 145.1956  |
| W Model          | 32.4711 | 0.9093 | 165.6799  |
| W Wavelet Model | 34.7337 | 0.9482 | 88.9002  |
+-----+-----+-----+-----+

```

## Part 6 : Output of predictions of all 4 models side by side for direct visualize comparison

```

def visualize_predictions_all_models(X_test, n, hformer_predictions, x_predictions, y_predictions, sa_predictions):
    random_numbers = list(range(n)) # not very random
    for i in random_numbers:
        gt_image= X_test[i]

        hformer_pred = hformer_predictions[i]
        x_pred = x_predictions[i]
        y_pred = y_predictions[i]
        sa_pred = sa_predictions[i]
        z_pred = z_predictions[i]
        w_pred = w_predictions[i]
        wavelet_w_pred = wavelet_w_predictions[i]

        models = ["HFORMER", "X", "Y", "SA", "Z", "W", "Wavelet W"]
        predictions = [hformer_pred, x_pred, y_pred, sa_pred, z_pred, w_pred, wavelet_w_pred]

```

```
# Display QD and FD images
f, axarr = plt.subplots(1, 1 + len(models), figsize=(41,41))

axarr[0].imshow(trunc(denormalize(gt_image)), cmap='gray', vmin=-160.0, vmax=240.0)
axarr[0].set_title("FD Image")
axarr[0].set_axis_off()

for j, (model_name, predicted_image) in enumerate(zip(models, predictions), start=1):
    if predicted_image.shape[-1] == 3:
        predicted_image = rgb2gray(predicted_image)

    psnr_recon = calculate_psnr(trunc(denormalize(gt_image)), trunc(denormalize(predicted_image)))
    ssim_recon = calculate_ssim(trunc(denormalize(gt_image)), trunc(denormalize(predicted_image)))
    rmse_recon = calculate_rmse(trunc(denormalize(gt_image)), trunc(denormalize(predicted_image)))

    psnr_recon = round(psnr_recon, 4)
    ssim_recon = round(ssim_recon, 4)
    mse_recon = round(rmse_recon, 4)

    axarr[j].imshow(trunc(denormalize(predicted_image)), cmap='gray', vmin=-160.0, vmax=240.0)
    axarr[j].set_title("{}\nPSNR={}\nSSIM={}\nMSE={}".format(model_name, psnr_recon, ssim_recon, mse_recon))
    axarr[j].set_axis_off()

plt.savefig('..../output/novel_comparison_lodopabct/combined_outputs_image_index_{}.png'.format(index))
plt.show()
```

visualize\_predictions\_all\_models(noisy\_array, len(noisy\_array), hformer\_predictions, np.concatenate(x\_r)

